

CHAPTER 1

OVERVIEW OF CONTAINER TECHNOLOGY

Are there other competing container formats?

- lxd, used by lxc runtime <https://linuxcontainers.org/>
- aci from appc
<https://github.com/appc/spec/blob/master/spec/aci.md#app-container-image>
 - used by rkt from coreos
- docker (v1) deprecated 2/28/2017
<https://github.com/moby/moby/blob/master/image/spec/v1.md>
- docker (2v1) overly complicated b/c of backwards compat with docker v1
<https://github.com/docker/distribution/blob/master/docs/spec/manifest-v2-1.md>
- docker (2v2)
<https://github.com/docker/distribution/blob/master/docs/spec/manifest-v2-2.md>
- oci (originally based on docker 2v2)
<https://github.com/opencontainers/image-spec/blob/master/spec.md>

Consider the format used in one of the containers used in kubernetes:

```
$ skopeo inspect docker://k8s.gcr.io/pause
$ podman pull k8s.gcr.io/pause:latest
$ podman inspect k8s.gcr.io/pause:latest | grep Manifest
  "ManifestType": "application/vnd.docker.distribution.manifest.v2+json",
```

Consider what the registries used by cloud vendors support:

Google cloud: <https://cloud.google.com/container-registry/docs/image-formats>

Azure:

<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-image-formats>

aws: <https://docs.aws.amazon.com/AmazonECR/latest/userguide/image-manifest-formats.html>

OVERVIEW OF CONTAINER (RUNTIME) ARCHITECTURE

What are namespaces ?

- Innovations that followed include:
 - **Namespaces** - responsible for resource isolation
 - Think of the view-master !



- man 7 namespaces, nsenter
 - “A namespace wraps a global system resource in an abstraction that makes it appear to the processes within the namespace that they have their own isolated instance of the global resource”
- /proc/\$\$/ns/
 - Uts: hostname (uname -n)
 - Ipc: interprocess communications, shared memory (ipcs)
 - Net: ipv4/ipv6 stacks, routing, firewall, sockets
 - Pid: process id namespace (allows for migrating containers to new host, suspend/resume)
 - User: distinct uids and gids
 - Cgroup: cgroup root directory, prevents process X from escaping the limits imposed by ancestor cgroups

<https://www.redhat.com/sysadmin/pid-namespace>

[student@workstation ~]\$ unshare -Urp --mount-proc

```
[root@workstation ~]# sleep 9000 &
```

```
[1] 32
```

```
[root@workstation ~]# ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	11:48	pts/1	00:00:00	-bash
root	32	1	0	11:49	pts/1	00:00:00	sleep 9000
root	33	1	0	11:49	pts/1	00:00:00	ps -ef

(from a different terminal)

```
[student@workstation ~]$ ps -ef | grep 9000
```

student	4497	4466	0	11:49	pts/1	00:00:00	sleep 9000
student	4583	4414	0	11:51	pts/2	00:00:00	grep --color=auto 9000

(notice the different pid for this same sleep process)

```
[student@workstation ~]$ sudo nsenter -t 4497 -a
```

```
-bash: /root/.bash_profile: Permission denied
```

```
[root@workstation /]# ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	11:48	pts/1	00:00:00	-bash
root	32	1	0	11:49	pts/1	00:00:00	sleep 9000
root	34	0	3	11:51	pts/2	00:00:00	-bash
root	61	34	0	11:51	pts/2	00:00:00	ps -ef

What are control groups ?

- **Control groups** (cgroups) - limits what resources a process group can consume. Allows processes to be organized into hierarchical groups so that usage of particular resources can be limited and monitored.

Recently, distros like Fedora31 have made cgroups V2 the default. See

<https://www.redhat.com/sysadmin/fedora-31-control-group-v2>

```
[root@workstation ~]# cd /sys/fs/cgroup/pids/system.slice/sss.service
```

```
[root@workstation sssd.service]# ls
```

```
cgroup.clone_children cgroup.procs notify_on_release pids.current pids.events pids.max tasks
```

```
[root@workstation sssd.service]# cat pids.max
```

```
36445
```

```
[root@workstation sssd.service]# cat pids.current
```

3

- NOTE: memory accounting from /proc/meminfo is NOT namespaced. So, a container's view of memory from tools like free/top will show the system accounting vs /sys/fs/cgroup/memory/memory.usage_in_bytes
 - <https://ops.tips/blog/why-top-inside-container-wrong-memory/>

What is seccomp?

- **Seccomp** - limits what system calls a process can make... even if running as root ! "secure computing mode"
 - /proc/sys/kernel/seccomp/actions_avail
 - /proc/sys/kernel/seccomp/actions_logged
 - man 2 seccomp {fedora has this one}
 - Limits what system calls by either read(), write(), _exit(), and sigreturn() or by a list of allowed calls given as "filters"
 - The Seccomp field of the /proc/[pid]/status file provides a method of viewing the seccomp mode of a process
 - [student@workstation ~]\$ cat /proc/3421/status | grep -i seccomp
 - Seccomp: 0
 - 0 means SECCOMP_MODE_DISABLED; 1 means SECCOMP_MODE_STRICT; 2 means SECCOMP_MODE_FILTER
 - In SECCOMP_MODE_STRICT, it cannot use any system calls except exit(), sigreturn(), read() and write().
 - In SECCOMP_MODE_FILTER, since linux 3.5, it is possible to define advanced custom filters based on the BPF (Berkley Packet Filters) to limit what system calls and their arguments can be used by the process.
 - podman run --security-opt=seccomp=unconfined
 - Will use /usr/share/containers/seccomp.json as the default profile

What are capabilities?

- **Capabilities** -A related and additionally important feature are Capabilities.
 - **man 7 capabilities** “Starting with kernel 2.2, Linux divides the privileges traditionally associated with superuser into distinct units, known as capabilities, which can be independently enabled and disabled. Capabilities are a per-thread attribute.”

```
[student@workstation ~]$ ps -ef | grep sshd
```

```
root    1049  1 0 08:11 ?    00:00:00 /usr/sbin/sshd -D
```

```
[student@workstation ~]$ grep Cap /proc/1049/status
```

```
CapInh: 0000000000000000
```

```
CapPrm: 000001ffffffff
```

```
CapEff: 000001ffffffff
```

```
CapBnd: 000001ffffffff
```

```
CapAmb: 0000000000000000
```

```
[student@workstation ~]$ capsh --decode=000001ffffffff
```

```
0x000001ffffffff=cap_chown,cap_dac_override,cap_dac_read_search,cap_fowner,cap_fsetid,cap_kill,cap_setgid,cap_setuid,cap_setpcap,cap_linux_immutable,cap_net_bind_service,cap_net_broadcast,cap_net_admin,cap_net_raw,cap_ipc_lock,cap_ipc_owner,cap_sys_module,cap_sys_rawio,cap_sys_chroot,cap_sys_ptrace,cap_sys_pacct,cap_sys_admin,cap_sys_boot,cap_sys_nice,cap_sys_resource,cap_sys_time,cap_sys_tty_config,cap_mknod,cap_lease,cap_audit_write,cap_audit_control,cap_setfcap,cap_mac_override,cap_mac_admin,cap_syslog,cap_wake_alarm,cap_block_suspend,cap_audit_read,cap_perfmon,cap_bpf,cap_checkpoint_restore
```

CapPrm = Permitted Capabilities

CapBnd = Bounding Capabilities

CapEff = Effective Capabilities

Capabilities can also be assigned to a file binary (getcap, setcap):

```
[student@workstation sbin]$ getcap /sbin/arping  
/sbin/arping cap_net_raw=p
```

- e: Effective
This means the capability is “activated”.
- p: Permitted
This means the capability can be used/is allowed.
- i: Inherited
The capability is kept by child/subprocesses upon execve() for example

What container specific selinux labels are used ?

- **SELinux** - protects processes from each other since they will all be running on the host system
 - ps -Z

```
[student@workstation ~]$ cat /usr/share/containers/selinux/contexts  
process = "system_u:system_r:container_t:s0"  
file = "system_u:object_r:container_file_t:s0"  
ro_file="system_u:object_r:container_ro_file_t:s0"  
kvm_process = "system_u:system_r:container_kvm_t:s0"  
init_process = "system_u:system_r:container_init_t:s0"  
engine_process = "system_u:system_r:container_engine_t:s0"
```

What happened to svirt_sandbox_file_t ?

```
[root@workstation ~]# mkdir /testdir  
[root@workstation ~]# chcon -t svirt_sandbox_file_t /testdir  
[root@workstation ~]# podman run -it -v /testdir:/data rhel7 /bin/bash  
[root@a8b47cb39617 /]# touch /data/test1
```

```
# rpm -q selinux-policy
```

selinux-policy-3.13.1-229.el7.noarch

yum install setools-console

seinfo -tcontainer_file_t -x

```
container_file_t
    device_node
    file_type
    filesystem_type
    mountpoint
    non_auth_file_type
    non_security_file_type
    noxattrfs
    ptynode
Aliases
    svirt_sandbox_file_t
    svirt_lxc_file_t
```

"""

Aliases are alternate names used to refer to a type. We can use an alias anywhere that we would use a type name, including TE rules, security contexts, and labeling statements. Aliases are typically used for compatibility when making policy changes. For example, an older policy might refer to the type `netscape_t`. An updated policy might switch to the type name to `mozilla_t`, but provide `netscape_t` as an alias to allow older modules to correctly compile.

"""

<https://flylib.com/books/en/2.803.1.40/1/>

<https://danwalsh.livejournal.com/81756.html>

What other selinux target contexts are allowed by selinux policy ?

[student@workstation ~]\$ **sudo yum install selinux-policy-doc**


```
[student@workstation ~]$ man container_selinux
```

MANAGED FILES

The SELinux process type `container_t` can manage files labeled with the following file types:

```
cephfs_t
cifs_t
container_file_t
fusefs_t
hugetlbfs_t
nfs_t
onload_fs_t
```

```
[student@workstation ~]$ podman run -d -p 8080:8080
```

```
registry.redhat.io/rhscv/httpd-24-rhel7
```

```
[student@workstation ~]$ podman exec -it 8b /bin/bash
```

```
bash-4.2$ cd /var/www/html/
```

```
bash-4.2$ echo "helloworld" > index.html
```

```
bash-4.2$ ls -lZ
```

```
-rw-r--r--. default root system_u:object_r:fusefs_t:s0      index.html
```

```
bash-4.2$ ps -efZ
```

LABEL	UID	PID	PPID	C	STIME	TTY	TIME	CMD
system_u:system_r:container_t:s0:c468,c532	default	1	0	0	16:38	?	00:00:00	httpd -D
BACKGROUND								
system_u:system_r:container_t:s0:c468,c532	default	40	1	0	16:38	?	00:00:00	/usr/bin/cat
system_u:system_r:container_t:s0:c468,c532	default	41	1	0	16:38	?	00:00:00	/usr/bin/cat
system_u:system_r:container_t:s0:c468,c532	default	42	1	0	16:38	?	00:00:00	/usr/bin/cat
system_u:system_r:container_t:s0:c468,c532	default	43	1	0	16:38	?	00:00:00	/usr/bin/cat
system_u:system_r:container_t:s0:c468,c532	default	44	1	0	16:38	?	00:00:00	httpd -D
BACKGROUND								
system_u:system_r:container_t:s0:c468,c532	default	45	1	0	16:38	?	00:00:00	httpd -D
BACKGROUND								
system_u:system_r:container_t:s0:c468,c532	default	52	1	0	16:38	?	00:00:00	httpd -D
BACKGROUND								
system_u:system_r:container_t:s0:c468,c532	default	67	1	0	16:38	?	00:00:00	httpd -D
BACKGROUND								
system_u:system_r:container_t:s0:c468,c532	default	69	1	0	16:38	?	00:00:00	httpd -D
BACKGROUND								
system_u:system_r:container_t:s0:c468,c532	default	102	0	0	16:48	pts/0	00:00:00	/bin/bash
system_u:system_r:container_t:s0:c468,c532	default	111	102	0	16:49	pts/0	00:00:00	ps -efZ

```
[student@workstation ~]$ curl localhost:8080  
helloworld
```

```
[student@workstation ~]$ podman rm -a -f
```

MANAGING CONTAINERS WITH PODMAN

How to install podman and container-tools in RHEL ?

```
[root@rhel7 ~]# subscription-manager repos --enable rhel-7-server-extras-rpms  
[root@rhel7 ~]# yum install podman
```

This training environment uses RHEL8.2 on workstation VM:

```
[root@workstation ~]# subscription-manager repos --list-enabled  
This system has no repositories available through subscriptions.
```

```
[root@workstation ~]# yum repolist
```

Updating Subscription Management repositories.

Unable to read consumer identity

This system is not registered to Red Hat Subscription Management. You can use subscription-manager to register.

repo id	repo name
rhel-8.2-for-x86_64-appstream-rpms	Red Hat Enterprise Linux 8.2
AppStream (dvd)	
rhel-8.2-for-x86_64-baseos-rpms	Red Hat Enterprise Linux 8.2
BaseOS (dvd)	

(DONT ACTUALLY RUN THE NEXT COMMAND)

```
[root@rhel8 ~]# yum module list container-tools
```

Updating Subscription Management repositories.

Last metadata expiration check: 0:00:48 ago on Mon 19 Aug 2019 10:51:59 AM EDT.

Red Hat Enterprise Linux 8 for x86_64 - AppStream (RPMs)

Name	Stream	Profiles	Summary
------	--------	----------	---------

container-tools	1.0	common [d]	Common tools
and dependencies for container runtimes			
container-tools	rhel8 [d][e]	common [d] [i]	Common tools
and dependencies for container runtimes			

Hint: [d]efault, [e]nabled, [x]disabled, [i]nstalled

[root@workstation ~]# **yum module info container-tools**

```

: buildah-0:1.11.6-7.module+el8.2.0+5856+b8046c6d.x86_64
: common-2:2.0.6-1.module+el8.2.0+5182+3136e5d4.x86_64
: podman-0:1.6.4-10.module+el8.2.0+6063+e761893a.src
: runc-0:1.0.0-65.rc10.module+el8.2.0+5762+aaee29fb.src
: skopeo-1:0.1.40-10.module+el8.2.0+5955+6cd70ceb.src
: slirp4netns-0:0.4.2-3.git21fdece.module+el8.2.0+5658+9a15711d.src

```

Is there a module or group that can assist installation in Fedora ?

[root@badger ~]# **cat /etc/redhat-release**

Fedora release 34 (Thirty Four)

[root@badger ~]# **dnf repolist**

repo id	repo name
fedora	Fedora 34 - x86_64
fedora-cisco-openh264	Fedora 34 openh264 (From Cisco) - x86_64
fedora-modular	Fedora Modular 34 - x86_64
google-chrome	google-chrome
rpmfusion-free	RPM Fusion for Fedora 34 - Free
rpmfusion-free-updates	RPM Fusion for Fedora 34 - Free - Updates
updates	Fedora 34 - x86_64 - Updates
updates-modular	Fedora Modular 34 - x86_64 - Updates

[root@badger ~]# **dnf install -y @container-tools**

Last metadata expiration check: 1:17:32 ago on Tue 21 Sep 2021 05:29:15 AM CDT.

Module or Group 'container-tools' is not available.

Error: Nothing to do.

Although there is no container-tools group or module in Fedora, there is a container-management group:

```
[root@badger ~]# dnf groupinfo "Container Management"
```

Last metadata expiration check: 1:16:42 ago on Tue 21 Sep 2021 05:29:15 AM CDT.

Group: Container Management

Description: Tools for managing Linux containers

Default Packages:

podman

Optional Packages:

buildah

flatpak

flatpak-builder

origin-clients

```
[root@badger ~]# dnf install -y @container-management
```

CHAPTER 2 CREATING CONTAINERIZED SERVICES

- *Search for and fetch container images with Podman.*
- *Run and configure containers locally.*
- *Use the Red Hat Container Catalog.*

FETCHING CONTAINER IMAGES WITH PODMAN

How can you determine the available tags within a repository ?

```
$ podman search registry.access.redhat.com/rhel7 --list-tags
```

(but this only includes a limited list of tags)

Use skopeo:

```
[student@workstation ~]$ skopeo inspect docker://registry.access.redhat.com/rhel7 | grep 7.5
```

```
"7.5-424",
"7.5-245.1527091554",
"7.5-409.1533127727",
"7.5",
"7.5-231",
"7.5-404",
"7.5-433",
"7.5-245",
"7.5-409",
"7.7-529",
```

Alternatively,

```
[student@workstation ~]$ curl -L https://registry.access.redhat.com/v2/rhel7/tags/list
{"name": "rhel7", "tags": ["7.3-74", "7.4-120", "7.2-56", "7.3-89", "7.3-66", "7.5-424",
"7.5-245.1527091554", "7.4-129", "7.1-12", "7.6-122", "7.3-82", "7.7-384.1575996163",
"7.5-409.1533127727", "7.2-75", "7.2-38", "7.6", "7.7-348", "7.4", "7.5", "7.6-301.1561066494",
"7.4-164", "7.7", "7.8", "7.2-35", "7.7-269", "7.1-6", "7.6-122.1547747894", "7.5-231", "7.5-404",
"7.1-9", "7.1-24", "7.4-81", "7.6-362", "7.6-252", "7.6-202.1553789841", "7.2-104", "7.3-97",
"7.4-113", "7.7-384.1580117710", "7.6-119", "7.1-16", "7.1-4", "7.6-301", "7.5-433", "7.0-27",
"7.3-95", "7.6-115", "7.3-79", "7.0-21", "7.8-265", "7.4-152", "7.7-310", "7.5-245",
"7.6-252.1561619826", "7.3-53", "7.0-23", "7.7-481", "7.6-151.1550575774", "7.3", "7.2",
"7.4-105", "7.6-202.1554729462", "7.3-45", "7.5-409", "7.2-46", "7.6-151", "7.2-44", "7.2-43",
"7.1-11", "7.7-529", "7.2-61", "latest", "7.7-384", "7.2-84", "7.6-202"]}]}
```

We can make that prettier with json_reformat:

```
[[student@workstation ~]$ curl -L https://registry.access.redhat.com/v2/rhel7/tags/list |
json_reformat
```

```
[student@workstation ~]$ curl -L https://registry.access.redhat.com/v2/rhel7/tags/list |
json_reformat | grep 7.7
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload Total Spent Left	Speed			
100	467	100	467	0	0	1040	0 --:--:-- --:--:-- --:--:-- 1042
100	901	100	901	0	0	1852	0 --:--:-- --:--:-- --:--:-- 1852

```
"7.5-424",
"7.5-245.1527091554",
```

```
"7.5-409.1533127727",  
"7.5",  
"7.5-231",  
"7.5-404",  
"7.5-433",  
"7.5-245",  
"7.5-409",  
"7.7-529",
```

Now that we can identify an image within a repository that we want, let's pull a local copy using the "latest" tag:

```
[student@workstation ~]$ podman pull registry.access.redhat.com/rhel7:latest  
9a3387c8f6bc9b63b119dc61ddbaed6bb20795a7b187908ca1b5ecabc5c19aac
```

How does that compare to:

```
[student@workstation ~]$ podman pull registry.access.redhat.com/rhel7  
9a3387c8f6bc9b63b119dc61ddbaed6bb20795a7b187908ca1b5ecabc5c19aac
```

(same image id). What about a different version of rhel7?

```
[student@workstation ~]$ podman pull registry.access.redhat.com/rhel7:7.7  
6682529ce3faf028687cef4fc6ffb30f51a1eb805b3709d31cb92a54caeb3daf
```

What about trying to pull using a short-name alias ?

Consider:

```
[student@workstation ~]$ head /etc/containers/registries.conf.d/001-rhel-shortnames.conf  
[aliases]  
"3scale-amp2/3scale-rhel7-operator-metadata" =  
"registry.redhat.io/3scale-amp2/3scale-rhel7-operator-metadata"  
"3scale-amp2/3scale-rhel7-operator" = "registry.redhat.io/3scale-amp2/3scale-rhel7-operator"  
"3scale-amp24/wildcard-router" = "registry.redhat.io/3scale-amp24/wildcard-router"
```

```
[student@workstation ~]$ grep rhel7  
/etc/containers/registries.conf.d/001-rhel-shortnames.conf  
"rhel7/open-vm-tools" = "registry.access.redhat.com/rhel7/open-vm-tools"  
"rhel7" = "registry.access.redhat.com/rhel7"  
"rhel7/rhel-atomic" = "registry.access.redhat.com/rhel7/rhel-atomic"
```

"rhel7/rhel" = "registry.access.redhat.com/rhel7/rhel"

```
[student@workstation ~]$ podman pull rhel7:latest
Trying to pull registry.access.redhat.com/rhel7:latest...
Getting image source signatures
```

```
[student@workstation ~]$ podman images
REPOSITORY          TAG   IMAGE ID   CREATED   SIZE
registry.access.redhat.com/rhel7 latest 2664aa19856f 2 weeks ago 216 MB
registry.access.redhat.com/rhel7 7.7   6682529ce3fa 22 months ago 215 MB
```

RUNNING CONTAINERS

How to run my first container

```
[student@workstation ~]$ podman run 9a3387c8f6bc cat /etc/redhat-release
Red Hat Enterprise Linux Server release 7.9 (Maipo)

[student@workstation ~]$ podman run e64297b706b7 cat /etc/redhat-release
Red Hat Enterprise Linux Server release 7.5 (Maipo)
```

What about running a container we haven't "pulled" yet ?

```
[student@workstation ~]$ podman run rhel7:7.7 cat /etc/redhat-release
Red Hat Enterprise Linux Server release 7.7 (Maipo)
```

Podman run is also capable of pulling images that aren't available in the local image storage:

```
[student@workstation ~]$ podman run rhel7:7.8 cat /etc/redhat-release
Resolved "rhel7" as an alias (/etc/containers/registries.conf.d/001-rhel-shortnames.conf)
Trying to pull registry.access.redhat.com/rhel7:7.8...
Getting image source signatures
Checking if image destination supports signatures
Copying blob b13ffc206103 done
Copying blob 872582724f33 done
Copying config 9da37a6819 done
Writing manifest to image destination
Storing signatures
```

Red Hat Enterprise Linux Server release 7.8 (Maipo)

Notice we now have several images listed here:

```
[student@workstation ~]$ podman images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
registry.access.redhat.com/rhel7	latest	c7344c9fb18c	3 weeks ago	216 MB
registry.access.redhat.com/rhel7	7.8	9da37a681956	2 years ago	215 MB
registry.access.redhat.com/rhel	7.7	6682529ce3fa	2 years ago	215 MB
registry.access.redhat.com/rhel7	7.7	6682529ce3fa	2 years ago	215 MB

Let's take a closer look at the cat commands that were executed using **podman run**:

```
[student@workstation ~]$ podman run --help
```

NAME:

podman run - Run a command in a new container

USAGE:

podman run [command options] IMAGE [COMMAND [ARG...]]

- Containers will exit after command execution. So let's see what info we can find on those
 - Not in ps output

```
[student@workstation ~]$ ps -ef | grep cat
```

```
gdm    4068  3897  0 16:15 ?        00:00:00 /usr/libexec/gsd-print-notifications
root   7313  6645  0 18:30 pts/0    00:00:00 grep --color=auto ca
```

- Use **podman ps** instead:

```
[student@workstation ~]$ podman ps
```

```
[student@workstation ~]$ podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED
27a3d872ff92	registry.access.redhat.com/rhel:7.4	cat /etc/redhat-rel...	About a minute ago
Exited (0)	About a minute ago	confident_dubinsky	false


```
6905b9e93f0d registry.access.redhat.com/rhel7:7.5-404 cat /etc/redhat-rel... 6 minutes ago
Exited (0) 6 minutes ago dreamy_panini false
```

```
de3fa91850ee registry.access.redhat.com/rhel7:latest cat /etc/redhat-rel... 6 minutes ago
Exited (0) 6 minutes ago competent_bell false
```

- Some images have prebuilt commands that will be executed if the COMMAND is missing from the podman run:

```
[student@workstation ~]$ podman inspect registry.access.redhat.com/rhel7:latest | less
```

- (notice /bin/bash)
- So what happens when we run this without giving any command:

```
[student@workstation ~]$ podman run registry.access.redhat.com/rhel7:latest
```

```
[student@workstation ~]$
```

- Did it fail ? ps -a (no, just non-interactive)

```
[student@workstation ~]$ podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	IS INFRA
04e62d0683ff	registry.access.redhat.com/rhel7:latest	/bin/bash	5 seconds ago
Exited (0) 4 seconds ago	youthful_benz	false	

- So, to make it interactive run:

```
[student@workstation ~]$ podman run -i registry.access.redhat.com/rhel7:latest
```

```
asdljflkasdjf
```

```
/bin/bash: line 3: asdljflkasdjf: command not found
```

```
echo hello world
```

```
hello world
```

```
whoami
```

```
root
```

```
tty
```

```
not a tty
```

```
exit
```

- What about a terminal ? Let's add the -t flag

```
[student@workstation ~]$ podman run -it registry.access.redhat.com/rhel7:latest
```

```
[root@66f4d93191b7 /]# tty
```

```
/dev/pts/0
```

```
[root@66f4d93191b7 /]# exit
```

- Another useful option when running a container is the -d or detach, for example:

```
[student@workstation ~]$ podman run -d registry.access.redhat.com/rhel7:latest sleep 5000
```

```
829e8264c3f722e047002ebf9bf55b38fcc9b2be3b6f0a2afdfb4088d01a3a7f
```

```
[student@workstation ~]$
```

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
829e8264c3f7	registry.access.redhat.com/rhel7:7.5	sleep 5000	4 seconds ago	Up 3 seconds ago
	loving_montalcini			

- Sometimes we will want to pass environment variables for the processes in the container. Do this with -e:

```
[student@workstation ~]$ podman run registry.access.redhat.com/rhel7:latest env
```

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

```
TERM=xterm
```

```
HOSTNAME=497c3c41f75f
```

```
container=oci
```

```
HOME=/root
```

```
[student@workstation ~]$ podman run -e FOO="hello world"
```

```
registry.access.redhat.com/rhel7:latest env
```

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

```
TERM=xterm
```

```
HOSTNAME=c730b9772e9a
```

```
container=oci
```

```
FOO=hello world
```

```
HOME=/root
```

- If you'd like to name a container use --name like:

```
[student@workstation ~]$ podman run -d --name mycontainer
```

```
registry.access.redhat.com/rhel7:latest sleep 5000
```

```
09b5adcfbcc0f894ef2d1782ebe5a28ba78e2bc67901814726b8d3b63f9545b5
```

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
09b5adcfbcc0	registry.access.redhat.com/rhel7:7.5	sleep 5000	3 seconds ago
	Up 2 seconds ago	mycontainer	

- To run a new command inside a running container use podman-exec:

```
[student@workstation ~]$ podman exec -it mycontainer /bin/bash
```

```
[root@db380c01c168 /]# ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
-----	-----	------	---	-------	-----	------	-----

```

root    1      0 0 18:58 ?    00:00:00 sleep 5000
root    5      0 3 19:00 pts/0      00:00:00 /bin/bash
root   17      5 0 19:00 pts/0      00:00:00 ps -ef

```

```
[root@db380c01c168 /]# cat /proc/1/cgroup
```

```

[root@07d1eca25e39 /]# grep Cap /proc/1/status
CapInh: 0000000000000000
CapPrm: 00000000800425fb
CapEff: 00000000800425fb
CapBnd: 00000000800425fb
CapAmb: 0000000000000000
[root@07d1eca25e39 /]# grep -i seccomp /proc/1/status
Seccomp: 2

```

```

[root@db380c01c168 /]# ps -efZ
LABEL                UID    PID  PPID  C STIME TTY    TIME CMD
system_u:system_r:container_t:s0:c478,c651 root 1  0  0 18:58 ?      00:00:00 sleep 5000

```

```

[root@db380c01c168 /]# ipcs -a
[root@db380c01c168 /]# exit
[student@workstation ~]$

```

How to change the hostname or otherwise identify which container you are running in ?

The hostname will be set to the container id by default:

```

[student@workstation ~]$ podman run -it rhel:latest /bin/bash
[root@560cf16fe847 /]# uname -n
560cf16fe847

```

```

[student@workstation ~]$ podman ps -a
CONTAINER ID  IMAGE                                COMMAND                  CREATED        STATUS
PORTS        NAMES
560cf16fe847 registry.access.redhat.com/rhel:latest /bin/bash              2 minutes ago Up 2 minutes
ago          kind_elion

```

/run/.containerenv tells us that we are in a container:

```
[root@560cf16fe847 /]# ls -l /run/.containerenv
-rw-r--r--. 1 root root 0 Nov 12 13:44 /run/.containerenv
```

Env tells us what kind:

```
[root@560cf16fe847 /]# env
container=oci
```

Hostname could be set to something however:

```
[student@workstation ~]$ podman run --hostname foo --name foo -it rhel:latest /bin/bash
```

Actually running podman commands in a container wouldn't be an easy thing to allow. Why?

```
[student@workstation ~]$ podman start foo
[student@workstation ~]$ podman exec -it foo /bin/bash
[root@foo /]# podman ps
bash: podman: command not found
```

Consider Nested containers:

https://developers.redhat.com/blog/2019/08/14/best-practices-for-running-buildah-in-a-container#running_buildah_inside_a_container

<https://github.com/containers/podman/issues/5188>

- To redirect host traffic to a container on a specific port use -p host:container

```
[student@workstation ~]$ man podman-run
```

-p, --publish=[]

Publish a container's port, or range of ports, to the host

Format: ip:hostPort:containerPort | ip::containerPort | **hostPort:containerPort** | containerPort

```
[student@workstation ~]$ podman run -d -p 8080:8080 --name httpd-basic
registry.access.redhat.com/rhsc1/httpd-24-rhel7
```

...

77f803ed8546027f65f9c4422f9d81d0696ce1de708b1ca9cdebafbb552890e1

```
[student@workstation ~]$ podman run -d -p 8080:8080 --name httpd-basic registry.redhat.io/rhsc1/httpd-24-rhel7
```

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
77f803ed8546	docker.io/library/httpd:2.4	httpd-foreground	30 seconds ago Up 29 seconds ago
0.0.0.0:8080->80/tcp	httpd-basic		

Check netstat:

```
[student@workstation ~]$ netstat -tunap | grep 8080
```

(Not all processes could be identified, non-owned process info will not be shown, you would have to be root to see it all.)

tcp6	0	0 :::8080	:::*	LISTEN	4827/rootlessport
------	---	-----------	------	--------	-------------------

```
[student@workstation ~]$ curl localhost:8080
```

<body>

<h1>Red Hat Enterprise Linux Test Page</h1>

What's that common process ?

Conmon is a monitoring program and communication tool between a container manager (like [podman](#) or [CRI-O](#)) and an OCI runtime (like [runc](#) or [crun](#)) for a single container.

```
[student@workstation ~]$ ps -ef | grep httpd
```

student	3761	3751	0 07:36 ?	00:00:00 httpd -DFOREGROUND
100000	3775	3761	0 07:36 ?	00:00:00 httpd -DFOREGROUND
100000	3776	3761	0 07:36 ?	00:00:00 httpd -DFOREGROUND
100000	3777	3761	0 07:36 ?	00:00:00 httpd -DFOREGROUND

```
[student@workstation ~]$ ps -ef | grep 3751
```

```
student      3751    1 0 07:36 ?    00:00:00 /usr/bin/common --api-version 1 -c
815d8f11fd3c3dc67bf2e6913fcdf6d2a517c3387eed801d89588ca2e7b1e2e1 -u
815d8f11fd3c3dc67bf2e6913fcdf6d2a517c3387eed801d89588ca2e7b1e2e1 -r /usr/bin/runc -b
/home/student/.local/share/containers/storage/overlay-containers/815d8f11fd3c3dc67bf2e6913f
cdf6d2a517c3387eed801d89588ca2e7b1e2e1/userdata
```

Another way to look at this is with pstree:

```
[student@workstation ~]$ pstree
```

```

|---common---|---httpd---3*[httpd---26*[{httpd}]]
|             |
|             |---{common}
```

```
[student@workstation ~]$ /usr/libexec/podman/common -h
```

<https://github.com/containers/common>

Can you run an image built from architectures different than the container host ?

No, but you can pull them. See:

<https://www.redhat.com/sysadmin/specify-architecture-pulling-podman-images>

```
[student@workstation ~]$ podman pull --arch=arm64 registry.access.redhat.com/ubi8
```

```
[student@workstation ~]$ podman pull registry.access.redhat.com/ubi8
```

```
[student@workstation ~]$ podman images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
registry.access.redhat.com/rhel7/rhel	latest	e2c37c467077	2 weeks ago	216 MB
registry.access.redhat.com/rhel7	7.9	e2c37c467077	2 weeks ago	216 MB
registry.access.redhat.com/rhel7	latest	e2c37c467077	2 weeks ago	216 MB
registry.access.redhat.com/ubi8	latest	d5c70d09f361	3 weeks ago	246 MB
<none>	<none>	2fd9e1478809	3 weeks ago	225 MB

```
[student@workstation ~]$ podman run 2fd9e1478809 uname -a
```

```
Linux f1b40336fff3 4.18.0-348.2.1.el8_5.x86_64 #1 SMP Mon Nov 8 13:30:15 EST 2021
x86_64 x86_64 x86_64 GNU/Linux
[student@workstation ~]$ podman run d5c70d09f361 uname -a
standard_init_linux.go:228: exec user process caused: exec format error
```

It is possible to use a qemu process to emulate different architectures. See <https://github.com/multiarch/qemu-user-static>

USING ROOTLESS CONTAINERS

What are the disadvantages of running rootless containers ?

SEE <https://github.com/containers/podman/blob/master/rootless.md>

Some subcommands will not work, esp ones that depend on features like cgroups:

```
[student@workstation ~]$ podman pause 382
Error: pause is not supported for rootless containers
```

```
[student@workstation ~]$ podman stats 382
Error: stats is not supported in rootless mode without cgroups v2
```

<https://www.redhat.com/sysadmin/behind-scenes-podman>

<https://opensource.com/article/19/2/how-does-rootless-podman-work>

https://indico.cern.ch/event/757415/contributions/3421994/attachments/1855302/3047064/Podman_Rootless_Containers.pdf

Also, networking is handled differently for rootless as a non-root user has limitations on what it can do to the host's network:

```
[student@workstation ~]$ podman run -d -p 808:8080 --name myhttpd  
registry.access.redhat.com/rhsc1/httpd-24-rhel7
```

```
Error: error from slirp4netns while setting up port redirection: map[desc:bad request:  
add_hostfwd: slirp_add_hostfwd failed]
```

How to better understand user namespaces?

Root = not different from the host

Rootless = maps user and group IDs to appear to be running under a different ID. Uses a “pause” process

```
[student@workstation ~]$ podman run -it rhel7  
[root@5367563cc886 /]# whoami  
root  
[root@5367563cc886 /]# id  
uid=0(root) gid=0(root) groups=0(root)
```

man 7 user_namespaces:

In particular, a process can have a normal unprivileged user ID outside a user namespace while at the same time having a user ID of 0 inside the namespace; in other words, the process has full privileges for operations inside the user namespace, but is unprivileged for operations outside the namespace.

User and group ID mappings: uid_map and gid_map

When a user namespace is created, it starts out without a mapping of user IDs (group IDs) to the parent user namespace. The /proc/[pid]/uid_map and /proc/[pid]/gid_map files (available since Linux 3.5) expose the mappings for user and group IDs inside the user namespace for the process pid.

Each line in the uid_map file specifies a 1-to-1 mapping of a range of contiguous user IDs between two user namespaces. The first two numbers specify the starting user ID in each of the two user namespaces. The third number specifies the length of the mapped range.


```
[root@5367563cc886 ~]$ cat /proc/self/uid_map
```

(start of range)	(parent ns)	(range)
0	1000	1
1	100000	65536

So, the “root” user inside this namespace maps to the user with uid=1000 in the parent namespace (ie the “student” user).

A user with uid=1 in the child namespace would have a uid of 100000 in the parent namespace and increment up from there in the respective namespaces:

```
[root@c0ec71b5ed9e /]# cat /etc/passwd
```

```
[root@c0ec71b5ed9e /]# id 1
```

```
uid=1(bin) gid=1(bin) groups=1(bin)
```

```
[root@c0ec71b5ed9e /]# id 2
```

```
uid=2(daemon) gid=2(daemon) groups=2(daemon)
```

```
[root@c0ec71b5ed9e /]# id 3
```

```
uid=3(adm) gid=4(adm) groups=4(adm)
```

These users would map to 100000, 100001, and 100002 respectively:

uid=2 would be 100001

uid=3 -> 100002

uid=x -> 100000+(x-1)

What happens when we create a new user inside this container ?

```
[root@c0ec71b5ed9e /]# useradd foo
```

```
[root@c0ec71b5ed9e /]# id foo
```

```
uid=1000(foo) gid=1000(foo) groups=1000(foo)
```

Within this container user_namespace, the “foo” user has a uid=1000. What would be that user’s id outside the container ?

Use our mapping algorithm: uid=x -> 100000+(x-1)

The foo user with uid=1000 inside the container would thus have a uid of 100000+(1000-1) or 100999.

We can inspect the ownership of the files in the home directory for the upperdir (ephemeral storage) to prove it:

```
[student@workstation ~]$ podman inspect 99 | less
(look for UpperDir)
[student@workstation ~]$ cd
/home/student/.local/share/containers/storage/overlay/7088d79675cdfe1be4cb8be64428f3
0a510108688758b78e33ea9990e8c8edd5/diff
[student@workstation diff]$ ls
etc home root run var
[student@workstation diff]$ cd home/
[student@workstation home]$ ls -l
total 0
drwx-----. 2 100999 100999 62 Sep 21 16:37 foo
```

Understanding rootless networking

Root = virtual ethernet device

Rootless = Slirp, tap device

Container networking normally uses CNI plugins to configure a bridge, but that would require root. For rootless, podman will execute /usr/bin/slirp4netns to setup networking. This command will create a tap device that is injected inside the new networking namespace.

Also, ping might not work depending on the RHEL version:

```
[student@workstation ~]$ podman run -it ubi8 /bin/bash
[root@840855c79201 /]# yum install iputils
[root@ff226094dfd3 /]# ping google.com
PING google.com (172.217.1.238) 56(84) bytes of data.
^C
--- google.com ping statistics ---
57 packets transmitted, 0 received, 100% packet loss, time 57365ms
```

Fixed per https://bugzilla.redhat.com/show_bug.cgi?id=2037807

```
[student@workstation ~]$ rpm -q systemd
systemd-239-58.el8.x86_64
```

```
[student@workstation ~]$ rpm -q --changelog systemd
```

```
* Mon Feb 07 2022 systemd maintenance team <systemd-maint@redhat.com> - 239-57
- hash-funcs: introduce macro to create typesafe hash_ops (#2037807)
- hash-func: add destructors for key and value (#2037807)
- util: define free_func_t (#2037807)
- hash-funcs: make basic hash_ops typesafe (#2037807)
- test: add tests for destructors of hashmap or set (#2037807)
- man: document the new sysctl.d/ - prefix (#2037807)
- sysctl: if options are prefixed with "-" ignore write errors (#2037807)
- sysctl: fix segfault (#2037807)
```

<https://github.com/containers/podman/blob/main/troubleshooting.md#5-rootless-containers-cannot-ping-hosts>

```
[student@workstation ~]$ sysctl -a | grep ping
net.ipv4.ping_group_range = 0 2147483647
```

```
[student@workstation ~]$ podman run -it ubi8 /bin/bash
[root@34cb445d6819 /]# yum install iputils -y
```

```
[root@34cb445d6819 /]# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=255 time=4.21 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=255 time=1.70 ms
```

```
[root@34cb445d6819 /]# exit
```

Understanding rootless storage

Root = overlay2

Rootless = FUSE-overlayFS is legacy, now (RHEL8.5+) native overlay just inside a mount namespace

<https://www.redhat.com/sysadmin/podman-rootless-overlay>

“”

The fuse-overlay has been great. However, it is a user-space file system, which means it needs to do almost twice as much work as the kernel. Every read/write has to be interpreted by the fuse-overlay before being passed onto the host kernel. For heavy workloads that hammer the file system, the performance of fuse-overlay suffers. You could see the fuse-overlays pegging out the CPU. Bottom line, we should see better performance with native overlays, especially for heavy read/write containers in rootless mode. For example, podman build . performance should improve significantly. Note that when writing to volumes, the fuse-overlays is seldom used, so performance will not be affected

“”

<https://www.redhat.com/en/blog/whats-new-red-hat-enterprise-linux-85-container-tools> See “Better Performance with Native OverlayFS”

To see the mount you’ll have to look inside the mount namespace for a running container:

```
[student@workstation ~]$ lsns -t mnt
      NS TYPE NPROCS  PID USER      COMMAND
4026531840 mnt      69  2315 student /usr/lib/systemd/systemd --user
4026532257 mnt       2  2464 student podman
4026532464 mnt      1 24819 student sleep 5000
4026532599 mnt       1 24805 student /usr/bin/slipr4netns --disable-host-loopback --mtu=65520
--enable-sandbox --enable-seccomp -c -e 3 -r 4 --net
[student@workstation ~]$ cat /proc/24819/mounts | grep overlay
overlay / overlay
rw,context="system_u:object_r:container_file_t:s0:c111,c779",relatime,lowerdir=/home/student/.local/share/containers/storage/overlay//QP4SN4QNWT5H3ILUUG6SKBUFNI:/home/student/.local/share/containers/storage/overlay//GS7XTMTUHQDG36F4YKAI6PQYDZ,upperdir=/home/student/.local/share/containers/storage/overlay/d460fa813b863b8d70195872e9abe50811cc87dfb10663df41bbe362566892d4/diff,workdir=/home/student/.local/share/containers/storage/overlay/d460fa813b863b8d70195872e9abe50811cc87dfb10663df41bbe362566892d4/work 0 0
```

CHAPTER 3: MANAGING CONTAINERS

CONTAINER LIFE CYCLE MANAGEMENT WITH PODMAN

Objective: manage the life cycle of a container from creation to deletion

- Highlight Figure 3.1: Podman managing subcommands
- Getting syntax help on any of those commands - man pages !
- Podman has subcommands to: create a new container (run), delete a container (rm), list containers (ps), stop a container (stop), and start a process in a container (exec).

There is a rhel7 container image but not a rhel8 image. Where is the rhel8 container image?

Check out ubi on <https://catalog.redhat.com/software/containers/explore/>



ubi8/ubi

Red Hat Universal Base Image 8

by Red Hat, Inc.

Provides the latest release of the Red Hat Universal Base Image 8.

Updated 28 days ago

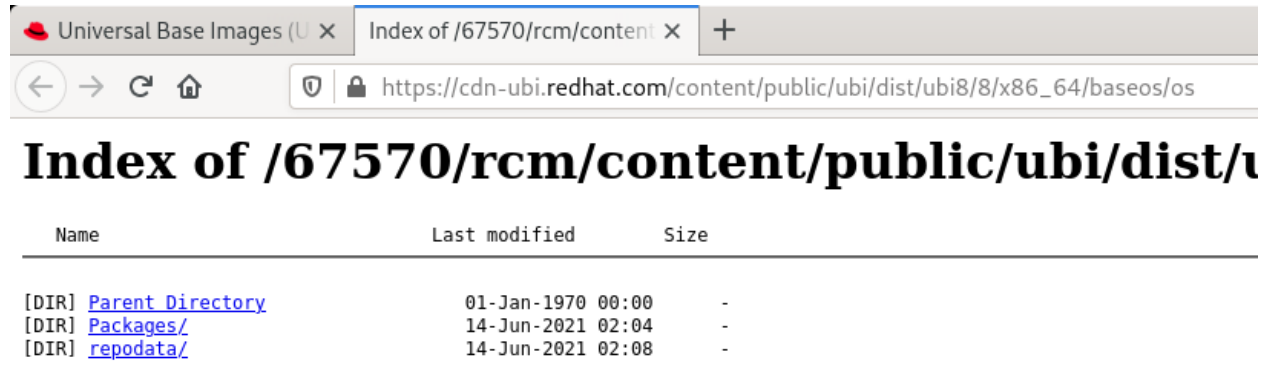
```
[student@workstation ~]$ podman pull registry.access.redhat.com/ubi8
Trying to pull registry.access.redhat.com/ubi8...
```

```
[student@workstation ~]$ podman run 2722 cat /etc/redhat-release
Red Hat Enterprise Linux release 8.4 (Ootpa)
```

```
[student@workstation ~]$ podman run 2722 cat /etc/yum.repos.d/ubi.repo
[ubi-8-baseos]
name = Red Hat Universal Base Image 8 (RPMs) - BaseOS
baseurl = https://cdn-ubi.redhat.com/content/public/ubi/dist/ubi8/8/$basearch/baseos/os
enabled = 1
gpgkey = file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
gpgcheck = 1
```

<https://access.redhat.com/articles/4238681>

https://cdn-ubi.redhat.com/content/public/ubi/dist/ubi8/8/x86_64/baseos/os



Name	Last modified	Size
[DIR] Parent Directory	01-Jan-1970 00:00	-
[DIR] Packages/	14-Jun-2021 02:04	-
[DIR] repodata/	14-Jun-2021 02:08	-

Universal base images are great containers to use as building blocks for any application. They make great “parent” images.

How are the names autogenerated by podman determined ?

Names will be autogenerated for containers with the form ***adjective_famousperson***

SEE the sourcecode:

<https://github.com/containers/podman/blob/main/vendor/github.com/docker/docker/pkg/namesgenerator/names-generator.go>

```
[ablum@badger ~]$ cd /home/ablum/go/src/namesgenerator
```

```
[ablum@badger namesgenerator]$ go mod init
[ablum@badger namesgenerator]$ go build
[ablum@badger namesgenerator]$ ./namesgenerator
suspicious_colden
[ablum@badger namesgenerator]$ ./namesgenerator
priceless_davinci
[ablum@badger namesgenerator]$ ./namesgenerator
nice_varahamihira
```

```
    left = [...]string{
        "admiring",
        "adoring",
        "affectionate",
        "agitated",
        "amazing",
        "angry",

        right = [...]string{
            // Muhammad ibn Jābir al-Ḥarrānī al-Battānī was a founding father of astronomy.
https://en.wikipedia.org/wiki/Mu%E1%B8%A5ammad\_ibn\_J%C4%81bir\_al-%E1%B8%A4arr%C4%81n%C4%AB\_al-Batt%C4%81n%C4%AB
            "albattani",

            // Frances E. Allen, became the first female IBM Fellow in 1989. In 2006, she
            became the first female recipient of the ACM's Turing Award.
https://en.wikipedia.org/wiki/Frances\_E.\_Allen
            "allen",

func main() {
begin:
    rand.Seed(time.Now().UnixNano())
    name := fmt.Sprintf("%s_%s", left[rand.Intn(len(left))], right[rand.Intn(len(right))])
    if name == "boring_wozniak" /* Steve Wozniak is not boring */ {
        goto begin
    }
    fmt.Println(name)
}
```

What is actually running inside the httpd image from rhsc1 ?

```
[student@workstation ~]$ podman ps --no-trunc
```

CONTAINER ID	IMAGE
a4b6429a3108095e1fdf1c509e105132f96a2a1da2eb2ce38614283f1151fb59	
registry.access.redhat.com/rhsc1/httpd-24-rhel7:latest	/usr/bin/run-httpd 12 minutes ago Up 12 minutes ago myhttpd
1d4db13a4995d493eb4176a3e88b05255e15cf9899dc336927b77882944734dc	
registry.access.redhat.com/rhsc1/httpd-24-rhel7:latest	/usr/bin/run-httpd 13 minutes ago Up 13 minutes ago objective_khorana

We could run another program inside the same namespaces of our myhttpd container using `podman exec`:

```
[student@workstation ~]$ podman exec -it myhttpd /bin/bash
```

```
bash-4.2$
```

```
bash-4.2$ ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
default	1	0	0	14:55	?	00:00:00	httpd -D FOREGROUND
default	40	1	0	14:55	?	00:00:00	/usr/bin/cat
default	41	1	0	14:55	?	00:00:00	/usr/bin/cat
default	42	1	0	14:55	?	00:00:00	/usr/bin/cat
default	43	1	0	14:55	?	00:00:00	/usr/bin/cat
default	44	1	0	14:55	?	00:00:00	httpd -D FOREGROUND
default	45	1	0	14:55	?	00:00:00	httpd -D FOREGROUND
default	54	1	0	14:55	?	00:00:00	httpd -D FOREGROUND
default	66	1	0	14:55	?	00:00:00	httpd -D FOREGROUND
default	69	1	0	14:55	?	00:00:00	httpd -D FOREGROUND
default	90	0	0	15:09	pts/0	00:00:00	/bin/bash
default	99	90	0	15:10	pts/0	00:00:00	ps -ef

What about the run-httpd ?

```
bash-4.2$ cat /usr/bin/run-httpd
```

```
...SNIP...
```

```
process_extending_files ${HTTPD_APP_ROOT}/src/httpd-pre-init/  
${HTTPD_CONTAINER_SCRIPTS_PATH}/pre-init/
```

```
exec httpd -D FOREGROUND $@
```


So, this script (a wrapper) executed the `httpd -D FOREGROUND` we see running within this namespace.

What about creating a `systemd.unit` file so that this container is started on system boot ?

```
[student@workstation ~]$ podman generate systemd -n myhttpd
```

```
# container-myhttpd.service
# autogenerated by Podman 1.6.4
# Mon Sep 13 15:42:02 EDT 2021
```

```
[Unit]
```

```
Description=Podman container-myhttpd.service
Documentation=man:podman-generate-systemd(1)
```

```
[Service]
```

```
Restart=on-failure
ExecStart=/usr/bin/podman start myhttpd
ExecStop=/usr/bin/podman stop -t 10 myhttpd
KillMode=none
Type=forking
PIDFile=/run/user/1000/overlay-containers/a4b6429a3108095e1fdf1c509e105132f96a2a1da2e
b2ce38614283f1151fb59/userdata/common.pid
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
[student@workstation ~]$ podman stop myhttpd
```

```
[student@workstation ~]$ mkdir -p ~/.config/systemd/user
```

```
[student@workstation ~]$ podman generate systemd -n myhttpd >
~/.config/systemd/user/myhttpd.service
```

```
[student@workstation ~]$ systemctl --user daemon-reload
```

```
[student@workstation ~]$ systemctl --user enable myhttpd.service
```

```
[student@workstation ~]$ systemctl --user start myhttpd
```

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
a4b6429a3108	registry.access.redhat.com/rhsc	/usr/bin/run-http...	5
hours ago	Up 8 seconds ago	myhttpd	

1d4db13a4995 registry.access.redhat.com/rhsc/htp-24-rhel7:latest /usr/bin/run-htp... 5
hours ago Up 5 hours ago objective_khorana

```
[student@workstation ~]$ systemctl --user status myhttpd
```

- myhttpd.service - Podman container-myhttpd.service

Loaded: loaded (/home/student/.config/systemd/user/myhttpd.service; enabled; vendor preset: enabled)

Active: active (running) since Mon 2021-09-13 16:03:49 EDT; 37s ago

```
[student@workstation ~]$ systemctl --user stop myhttpd
```

SEE also

https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html-single/building_running_and_managing_containers/index#proc_enabling-systemd-services_assembly_porting-containers-to-systemd-using-podman

Now, it's time for us to remove the container...we are completely done with it. Consider, we have some running, but then some exited containers:

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
a4b6429a3108	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp... 5	
hours ago	Up 3 minutes ago	myhttpd	
1d4db13a4995	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp... 5	
hours ago	Up 5 hours ago	objective_khorana	

```
[student@workstation ~]$ podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
a4b6429a3108	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp... 5	
hours ago	Up 3 minutes ago	myhttpd	
1d4db13a4995	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp... 5	
hours ago	Up 5 hours ago	objective_khorana	
0a395937282a	registry.access.redhat.com/ubi8:latest	cat /etc/yum.repo... 5 hours	
ago	Exited (0) 5 hours ago	awesome_maxwell	
e4d157d70fad	registry.access.redhat.com/ubi8:latest	cat /etc/redhat-r... 5 hours	
ago	Exited (0) 5 hours ago	quirky_sinoussi	

To remove, we should use `podman rm`:

```
[student@workstation ~]$ podman rm myhttpd
```

Error: cannot remove container
a4b6429a3108095e1fdf1c509e105132f96a2a1da2eb2ce38614283f1151fb59 as it is running -
running or paused containers cannot be removed without force: container state improper

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
a4b6429a3108	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp...	5 hours ago
Up 6 seconds ago		myhttpd	
1d4db13a4995	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp...	5 hours ago
Up 5 hours ago		objective_khorana	

```
[student@workstation ~]$ podman stop myhttpd
```

```
A4b6429a3108095e1fdf1c509e105132f96a2a1da2eb2ce38614283f1151fb59
```

Ok, good, now we check to make sure its stopped:

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
a4b6429a3108	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp...	5 hours ago
Up Less than a second ago		myhttpd	
1d4db13a4995	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp...	5 hours ago
Up 5 hours ago		objective_khorana	

Grr, systemd you are too good...its still running:

```
[student@workstation ~]$ systemctl --user stop myhttpd
```

```
[student@workstation ~]$ systemctl --user disable myhttpd
```

```
Removed /home/student/.config/systemd/user/multi-user.target.wants/myhttpd.service.
```

```
[student@workstation ~]$ podman rm myhttpd
```

```
A4b6429a3108095e1fdf1c509e105132f96a2a1da2eb2ce38614283f1151fb59
```

```
[student@workstation ~]$ podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
1d4db13a4995	registry.access.redhat.com/rhsc/htp-24-rhel7:latest	/usr/bin/run-htp...	5 hours ago
Up 5 hours ago		objective_khorana	
0a395937282a	registry.access.redhat.com/ubi8:latest	cat /etc/yum.repo...	5 hours ago
Exited (0) 5 hours ago		awesome_maxwell	
e4d157d70fad	registry.access.redhat.com/ubi8:latest	cat /etc/redhat-r...	5 hours ago
Exited (0) 5 hours ago		quirky_sinoussi	

```
[student@workstation ~]$
```

Let's remove all these:

```
[student@workstation ~]$
```

```
[student@workstation ~]$ podman rm -a
```

```
0a395937282a3a57e0f5fddcd19563f68d29ada3e0bd895f95a520164cd3edfb
```

```
e4d157d70fad0e2252b7fa516ca0442847123c3a6e5109dbf6cdbff9f30dd74a
```

```
Error: cannot remove container
```

```
1d4db13a4995d493eb4176a3e88b05255e15cf9899dc336927b77882944734dc as it is running -  
running or paused containers cannot be removed without force: container state improper
```

```
[student@workstation ~]$ podman rm -a -f
```

```
1d4db13a4995d493eb4176a3e88b05255e15cf9899dc336927b77882944734dc
```

```
[student@workstation ~]$ podman ps -a
```

```
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
```

How to extract metadata from `podman inspect` ?

```
[student@workstation ~]$ podman inspect distracted_grothendieck
```

```
[student@workstation ~]$ man podman-inspect
```

This displays the low-level information on containers and images identified by name or ID. By default, this will render all results in a JSON array.

First, Let's understand JSON a bit:

(From: <https://developers.squarespace.com/what-is-json>)

JSON, or JavaScript Object Notation, is a minimal, readable format for structuring data.

Primarily this is done using "key" and "value" pairs...or

"Key" : "value"

Values can be any of the following:

- **String:** Several plain text characters which usually form a word
- **Boolean:** True or false.

- **Number:** An integer.
- **Object:** An associative array of key/value pairs....a “dictionary” { }
- **Array:** An associative array of value...a “list” []

Let’s go back to the podman-inspect man page:

`--format, -f="FORMAT"`

Format the output using the given Go template. The keys of the returned JSON can be used as the values for the `--format` flag (see examples below).

- Uses Go syntax: “A template variable can be a boolean, string, character, integer, floating-point, imaginary, or complex constant in Go syntax. Data passed to the template can be accessed using dot `{{ . }}`”.
- “Actions” — data evaluations or control structures — is delimited by `{{and}}`. The data evaluated inside it is called a *Pipeline*. Anything outside them is sent to the output unchanged.
- If the data is a complex type then it’s fields can be accessed using the dot `{{ .FieldName }}`. Dots can be chained together if the data contains multiple complex structures. `{{ .Struct.StructTwo.Field }}`
 - <https://curtisvermeeren.github.io/2017/09/14/Golang-Templates-Cheatsheet>
 - <https://www.openshift.com/blog/customizing-oc-output-with-go-templates>

```
[student@workstation ~]# podman inspect 166196236b59 --format '{{.Created}}'
2019-07-29 19:48:31.858078856 +0000 UTC
```

```
[student@workstation ~]# podman inspect 5b6 --format '{{.State.Pid}}'
3404
```

What about processing a complex object that contains a list of other objects. Consider the `Ulimits` list here:

```
[student@workstation ~]# podman inspect 5b6
```

```
"HostConfig": {
...SNIP...
  "Ulimits": [
    {
      "Name": "RLIMIT_NOFILE",
      "Soft": 1048576,
      "Hard": 1048576
```

```

    },
    {
      "Name": "RLIMIT_NPROC",
      "Soft": 1048576,
      "Hard": 1048576
    }
  ],

```

```
[student@workstation ~]# podman inspect 5b6 --format '{{.HostConfig.Ulimits}}'
[{{RLIMIT_NOFILE 1048576 1048576}} {{RLIMIT_NPROC 1048576 1048576}}]
```

Lets loop through the Ulimits list printing only the names for each item:

```
[student@workstation ~]# podman inspect 5b6 --format '{{range .HostConfig.Ulimits}}{{.Name}}{{end}}'
RLIMIT_NOFILERLIMIT_NPROC
```

```
[student@workstation ~]# podman inspect 5b6 --format '{{range .HostConfig.Ulimits}}name: {{.Name}} hard: {{.Hard}}{{end}}'
name: RLIMIT_NOFILE hard: 1048576name: RLIMIT_NPROC hard: 1048576
```

```
[student@workstation ~]# podman inspect 5b6 --format '{{range .HostConfig.Ulimits}}name: {{.Name}} hard: {{.Hard}}{{"\n"}}{{end}}'
name: RLIMIT_NOFILE hard: 1048576
name: RLIMIT_NPROC hard: 1048576
```

```
[student@workstation ~]# podman inspect a21 --format '{{index .HostConfig.Ulimits 0}}'
{{RLIMIT_NOFILE 1048576 1048576}}
[student@workstation ~]# podman inspect a21 --format '{{index .HostConfig.Ulimits 1}}'
{{RLIMIT_NPROC 1048576 1048576}}
```

If ulimits are not available, try the BoundingCaps:

```
[student@workstation ~]$ podman inspect myhttpd --format '{{range .BoundingCaps}}capability:{{.}}{{"\n"}}{{end}}'
capability:CAP_CHOWN
capability:CAP_DAC_OVERRIDE
capability:CAP_FOWNER
capability:CAP_FSETID
capability:CAP_KILL
```

```
capability:CAP_NET_BIND_SERVICE
capability:CAP_NET_RAW
capability:CAP_SETFCAP
capability:CAP_SETGID
capability:CAP_SETPCAP
capability:CAP_SETUID
capability:CAP_SYS_CHROOT
```

```
[student@workstation sbin]$ podman inspect 07d1eca25e39 --format '{{range
.BoundingCaps}}{{if eq . "CAP_KILL"}}eek this can kill{{end}}{{end}}'
eek this can kill
```

```
[student@workstation sbin]$ podman inspect 07d1eca25e39 --format '{{range
.BoundingCaps}}{{if eq . "CAP_FOO"}}eek this can kill{{end}}{{end}}'
```

```
[student@workstation ~]# podman inspect 5b6 --format '{{.State.Pid}}'
3404
```

```
[student@workstation ~]# ps -fp 3404
UID          PID    PPID  C STIME TTY          TIME CMD
root   3404   3392  0 06:56 ?        00:00:00 sleep 5000
```

Ok - but where does this fit in with openshift ? Consider
<https://www.openshift.com/blog/customizing-oc-output-with-go-templates>

```
[student@workstation ~]# podman exec -it 166196236b59 /bin/bash
bash-4.2$ exit
exit
```

You can also use this with podman ps to help create tables that are useful to inspect information.

For example,

```
[student@workstation ~]$ podman ps -a --format json
[student@workstation ~]$ podman ps -a --format '{{.Names}} {{.State}} {{.Image}}'
```

What is the use case for `podman pause` ?

`podman pause` uses the cgroup “freezer” to freeze (halt) a task without stopping it or without the task knowing.

NOTE: This is NOT supported with rootless due to a limit in the freezer cgroup.

The container and its processes are paused while the image is committed. This minimizes the likelihood of data corruption when creating the new image.

(man podman-commit)

<https://www.kernel.org/doc/Documentation/cgroup-v1/freezer-subsystem.txt>

The cgroup freezer will also be useful for checkpointing running groups of tasks. The cgroup freezer is hierarchical. Freezing a cgroup freezes all tasks belonging to the cgroup and all its descendant cgroups

```
[root@badger ~]# podman run -d docker.io/library/httpd
```

```
[root@badger ~]# podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
b0f5ce994715	docker.io/library/httpd:latest	httpd-foreground	4 seconds ago	Up 3 seconds

ago elated_montalcini

```
[root@badger ~]# ps -ef | grep httpd
```

```
root  27369 27357  2 15:31 ?        00:00:00 httpd -DFOREGROUND
bin    27389 27369  0 15:31 ?        00:00:00 httpd -DFOREGROUND
bin    27390 27369  0 15:31 ?        00:00:00 httpd -DFOREGROUND
bin    27391 27369  0 15:31 ?        00:00:00 httpd -DFOREGROUND
root   27509 19190  0 15:31 pts/1    00:00:00 grep --color=auto httpd
```

```
[root@badger ~]# cat /sys/fs/cgroup/freezer/machine.slice/libpod-b0f5ce*/freezer.state
THAWED
```

```
[root@badger ~]# podman pause b0f5ce994715
```

```
[root@badger ~]# cat /sys/fs/cgroup/freezer/machine.slice/libpod-b0f5ce*/freezer.state
```


FROZEN

```
[root@badger ~]# curl 10.88.0.30
(hangs)
```

```
[root@badger ~]# podman commit b0f5ce994715 docker.io/library/httpd:mypause
```

```
[root@badger ~]# podman images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
docker.io/library/httpd	mypause	1ab4200ebb2a	9 seconds ago	170 MB

```
[root@badger ~]# podman unpause b0f5ce994
```

```
b0f5ce9947155595b0d2d2c2c31dded774e8107c77e9658df0b2e23fc5c7306c
```

```
[root@badger ~]# curl 10.88.0.30
```

```
<html><body><h1>It works!</h1></body></html>
```

```
[root@badger ~]# cat /sys/fs/cgroup/freezer/machine.slice/libpod-b0f5ce*/freezer.state
THAWED
```

ATTACHING PERSISTENT STORAGE TO CONTAINERS

How does overlay work?

```
[root@workstation overlay-images]# cat /etc/containers/storage.conf | grep -v ^#
```

```
[storage]
```

```
driver = "overlay"
```

```
runroot = "/var/run/containers/storage"
```

```
graphroot = "/var/lib/containers/storage"
```

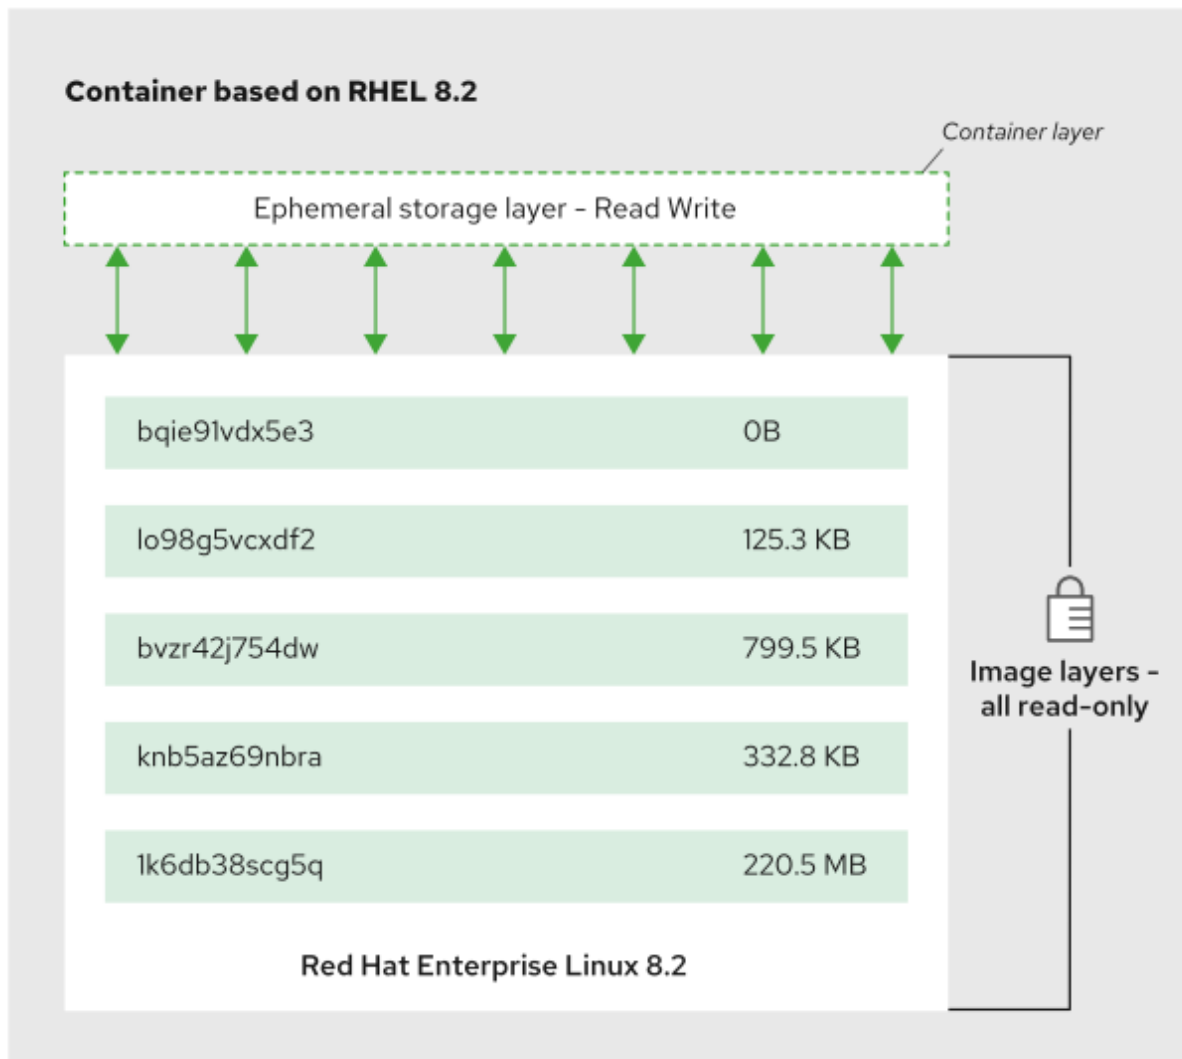


Figure 3.3: Container layers

<https://docs.docker.com/storage/storagedriver/overlayfs-driver/#how-the-overlay2-driver-works>

The overlay2 driver natively supports up to 128 lower OverlayFS layers - the original overlay driver only worked with 2 layers, extra layers in overlay relied on hard-linked directories. This created excessive use of inodes (known limitation)

NOTE1: rootless uses `/home/student/.local/share/containers/` for it's ephemeral storage

NOTE2: ~~rootless uses fuse overlays:~~ <https://github.com/containers/fuse-overlayfs>

```
[root@workstation storage]# mkdir -p /data/lower{1..3} /data/upper /data/work /data/merged
```

```
[root@workstation storage]# touch /data/lower1/file{1..3} /data/lower2/file{4..6}  
/data/lower3/file{7..9} /data/upper/file10
```

```
[root@workstation storage]# tree /data
```

```
[root@workstation storage]# mount -t overlay overlay -o  
lowerdir=/data/lower1:/data/lower2:/data/lower3,upperdir=/data/upper,workdir=/data/work  
/data/merged
```

```
[root@workstation storage]# tree /data
```

```
[root@workstation storage]# ls -l /data/merged
```

```
total 0
```

```
-rw-r--r--. 1 root root 0 Jul 29 20:29 file1  
-rw-r--r--. 1 root root 0 Jul 29 20:29 file2  
-rw-r--r--. 1 root root 0 Jul 29 20:29 file3  
-rw-r--r--. 1 root root 0 Jul 29 20:30 file4  
-rw-r--r--. 1 root root 0 Jul 29 20:30 file5  
-rw-r--r--. 1 root root 0 Jul 29 20:30 file6  
-rw-r--r--. 1 root root 0 Jul 29 20:30 file7
```

```
[root@workstation storage]# echo hello > /data/merged/file10
```

Which directory(ies) should show a change ?

```
[root@workstation storage]# ls -lR /data/
```

Notice its 6 bytes in both the /data/merged AND /data/upper layer. Ok now let's modify a different file, one in the "lower" layers.

```
[root@workstation storage]# echo hello > /data/merged/file9
```

Which directory(ies) should show a change ?

```
[root@workstation storage]# ls -lR /data/
```

Notice "lower3" is where the original file9 was created in, but it didn't change:

```
/data/lower3:
```

```
total 0
-rw-r--r--. 1 root root 0 Aug 13 11:37 file7
-rw-r--r--. 1 root root 0 Aug 13 11:37 file8
-rw-r--r--. 1 root root 0 Aug 13 11:37 file9
```

The change shows up in the “upper” layer (and merged):

```
/data/upper:
total 8
-rw-r--r--. 1 root root 6 Aug 13 11:39 file10
-rw-r--r--. 1 root root 6 Aug 13 11:40 file9
```

Containers use overlay to share underlying images in “lower” layers. The upper is used to track changes only.

- Notice the layers for this container:

```
[root@workstation storage]# podman run -d rhsc1/httpd-24-rhel7
[root@workstation storage]# podman inspect 37 | less
```

```
    "GraphDriver": {
      "Name": "overlay",
      "Data": {
        "LowerDir":
"/var/lib/containers/storage/overlay/77b158205d2cc791c9dc6f1ca04db6dca5e0737fc8a44f6c60
df42626d4838be/diff:/var/lib/containers/storage/overlay/4d3aa8111d6a2805f45d11997708bb5a1
79add74d61da9543b10818caa72e1fd/diff:/var/lib/containers/storage/overlay/5201771aee3980a
1208cf5111c23763492e29cdacbae70d85f0dfbbeb3fa069c/diff:/var/lib/containers/storage/overlay
/da289ed398e809e9e58320f71e3f32a0cfc881fc7db55e1a3a7bb1125e5b8c1e/diff",
        "MergedDir":
"/var/lib/containers/storage/overlay/49701e617da01a61fa60a9ced94efaaa1697828c621a37f98d
6ab3650de3bc1d/merged",
        "UpperDir":
"/var/lib/containers/storage/overlay/49701e617da01a61fa60a9ced94efaaa1697828c621a37f98d
6ab3650de3bc1d/diff",
        "WorkDir":
"/var/lib/containers/storage/overlay/49701e617da01a61fa60a9ced94efaaa1697828c621a37f98d
6ab3650de3bc1d/work"
      }
    }
```

- To reclaim storage on the host system run **podman rm**
 - Consider both **podman ps** and **podman ps -a**

Has overlay2 always been the storage driver used on container hosts ?

No. See https://bugzilla.redhat.com/show_bug.cgi?id=1475625

- Do we support devicemapper ?

[1/30/2020]

We currently support both devicemapper and overlayfs (overlay2) with docker. Before RHEL 7.5, the default docker storage configuration is devicemapper in direct-lvm[1] mode. Starting with 7.5, the default docker storage configuration is overlay2[2]. I do not know how the various versions/releases of OCP 3.x change/overwrite the default docker storage configuration.

Podman, buildah, and CRI-O only support overlayfs[3][4], and this is configured by default.

CRI-O and docker can be installed and running on the same RHEL 7 host, but they use completely different, incompatible storage configurations.

So, the original question is do we currently support the docker devicemapper storage driver, correct? If so, yes, we still support devicemapper in direct-lvm mode. Hope that helps!

-Derrick

[1]

<https://docs.docker.com/storage/storagedriver/device-mapper-driver/#configure-direct-lvm-mode-for-production>

[2] <https://docs.docker.com/storage/storagedriver/overlayfs-driver/>

[3]

<https://www.redhat.com/en/blog/working-container-storage-library-and-tools-red-hat-enterprise-linux>

[4] https://bugzilla.redhat.com/show_bug.cgi?id=1774789#c3

What if we want persistent, performant storage ?

- Add “volumes” which our directories on the host system exposed to the container.
- Permissions and selinux matter !
 - Rootless containers will need to set permissions using `podman unshare`

```
[student@workstation ~]$ mkdir mydata
```

```
[student@workstation ~]$ podman run -d -v /home/student/mydata:/var/www/html
registry.redhat.io/rhsc1/httpd-24-rhel7
```

```
Trying to pull registry.redhat.io/rhsc1/httpd-24-rhel7...
```

```
8971a0ce72fb6fb35615f6d4dcb28b77747c16dec095befb08a2d4d558e8a3e
```

```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
8971a0ce72fb	registry.redhat.io/rhsc1/httpd-24-rhel7:latest	/usr/bin/run-http...	10 seconds ago
Up 9 seconds ago		musing_shirley	

```
[student@workstation ~]$
```

```
[student@workstation ~]$
```

```
[student@workstation ~]$ podman inspect 89 --format='{{.Mounts}}'
```

```
[{"bind /home/student/mydata /var/www/html [rbind] true rprivate}]
```

```
[student@workstation ~]$ podman exec -it 89 /bin/bash
```

```
bash-4.2$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
fuse-overlays	9.9G	6.6G	3.3G	67%	/
tmpfs	64M	0	64M	0%	/dev
tmpfs	580M	128K	580M	1%	/etc/hosts
shm	63M	0	63M	0%	/dev/shm
/dev/vda3	9.9G	6.6G	3.3G	67%	/var/www/html

```
bash-4.2$ touch /var/www/html/index.html
```

```
touch: cannot touch '/var/www/html/index.html': Permission denied
```

```
bash-4.2$ id
```

```
uid=1001(default) gid=0(root) groups=0(root)
```

```
bash-4.2$ ls -ld /var/www/html
```

```
drwxrwxr-x. 2 root root 6 Sep 14 14:20 /var/www/html
```

```
bash-4.2$ chown 1001 /var/www/html
```

```
chown: changing ownership of '/var/www/html': Permission denied
```

```
bash-4.2$
```

```
bash-4.2$ exit
```

exit

Error: non zero exit code: 1: OCI runtime error

Ok, let's try to fix these permissions outside the container

```
[student@workstation ~]$ ls -ldn mydata
drwxrwxr-x. 2 1000 1000 6 Sep 14 10:20 mydata
[student@workstation ~]$ chown 1001 mydata
chown: changing ownership of 'mydata': Operation not permitted
```

Failed. You'll need to use `podman unshare` to run the chown as if student's UID were 0 (ie root). SEE man podman-unshare

```
[student@workstation ~]$ podman unshare chown -R 1001 mydata
[student@workstation ~]$ ls -ldn mydata
drwxrwxr-x. 2 101000 1000 6 Sep 14 10:20 mydata
[student@workstation ~]$ cat /etc/subuid
student:100000:65536
devops:165536:65536
[student@workstation ~]$ bc
bc 1.07.1
Copyright 1991-1994, 1997, 1998, 2000, 2004, 2006, 2008, 2012-2017 Free Software
Foundation, Inc.
This is free software with ABSOLUTELY NO WARRANTY.
For details type `warranty'.
101000-100000
1000
```

NOTE: 1000 is 1 less than the mapped uid (1001) used within the user namespace of the containerized process. Consider making this a different uid (one typically used with mysql):

```
[student@workstation ~]$ podman unshare chown 27:27 mydata
[student@workstation ~]$ ls -ldn mydata
drwxrwxr-x. 2 100026 100026 6 Sep 14 10:20 mydata
```

Think it will be "27" inside the user namespace, but it maps to "100026" outside. This uid falls within the subordinate range (aka subuid) allocated to the student user:

```
[student@workstation ~]$ man subuid
```

Each line in /etc/subuid contains a user name and a range of subordinate user ids that user is allowed to use. This is specified with three fields delimited by colons (":"). These fields are:

- login name or UID
- numerical subordinate user ID
- numerical subordinate user ID count

Ok, let's return the ownership back to a subordinate ID that will map to permissions of the working user within the httpd container's namespace (ie the "default" user):

```
[student@workstation ~]$ podman unshare chown 1001:0 mydata
[student@workstation ~]$ ls -ldn mydata
drwxrwxr-x. 2 101000 1000 6 Sep 14 10:20 mydata
```

Verify with:

```
[student@workstation ~]$ podman exec -it 89 /bin/bash
bash-4.2$ ls -ldn /var/www/html
drwxrwxr-x. 2 1001 0 6 Sep 14 14:20 /var/www/html
bash-4.2$ id
uid=1001(default) gid=0(root) groups=0(root)
bash-4.2$ touch /var/www/html/index.html
touch: cannot touch '/var/www/html/index.html': Permission denied
```

Grrr... still not working, any thoughts ? SELinux ?

```
[student@workstation ~]$ sudo grep mydata /var/log/audit/audit.log
type=AVC msg=audit(1631631220.011:137): avc: denied { write } for pid=3436 comm="touch"
name="mydata" dev="vda3" ino=26100164
scontext=system_u:system_r:container_t:s0:c408,c911
tcontext=unconfined_u:object_r:user_home_t:s0 tclass=dir permissive=0
```

```
[student@workstation ~]$ ls -ldZ mydata
drwxrwxr-x. 2 101000 student unconfined_u:object_r:user_home_t:s0 6 Sep 14 10:20 mydata
```

```
[student@workstation ~]$ ps -eZ | grep httpd
system_u:system_r:container_t:s0:c408,c911 2896 ? 00:00:00 httpd
```

If we try to fix the permissions as the student user:

```
[student@workstation ~]$ chcon -t container_file_t mydata
chcon: failed to change context of 'mydata' to 'unconfined_u:object_r:container_file_t:s0':
Operation not permitted
```


We must do this with `podman unshare` like before:

```
[student@workstation ~]$ podman unshare chcon -t container_file_t mydata
[student@workstation ~]$ ls -ldZ mydata
drwxrwxr-x. 2 101000 student unconfined_u:object_r:container_file_t:s0 6 Sep 14 10:20 mydata
```

```
[student@workstation ~]$ podman exec -it 89 /bin/bash
bash-4.2$ echo hello > /var/www/html/index.html
bash-4.2$
bash-4.2$
bash-4.2$ exit
exit
[student@workstation ~]$ podman exec 89 curl -s "http://localhost:8080"
hello
```

works!

How to apply selinux labels ?

The selinux label on a volume can be made persistent on a host with:

```
[student@workstation ~]$ sudo semanage fcontext -a -t container_file_t
'/home/student/mydata(/.*)?'
[student@workstation ~]$ sudo restorecon -Rv /home/student/mydata/
/home/student/mydata not reset as customized by admin to
unconfined_u:object_r:container_file_t:s0
```

Alternatively use the :Z option like --

```
[student@workstation ~]$ mkdir mydata1
[student@workstation ~]$ podman unshare chown 1001 mydata1
[student@workstation ~]$ podman run -d -v /home/student/mydata1:/var/www/html:Z
registry.redhat.io/rhsc1/httpd-24-rhel7
9c9d8752085e22ad32407b1f655a1a49e2427ceeae46ca6455665c6b0412db96
[student@workstation ~]$ podman exec 9c ls -ldZ /var/www/html
drwxrwxr-x. default root system_u:object_r:container_file_t:s0:c1002,c1014 /var/www/html
[student@workstation ~]$ podman exec 9c ls -ldZ /var/www/html
drwxrwxr-x. default root system_u:object_r:container_file_t:s0:c1002,c1014 /var/www/html
```

```
[student@workstation ~]$ podman exec 9c touch /var/www/html/index.html
[student@workstation ~]$ ls -lZ mydata1
total 0
-rw-r--r--. 1 101000 student system_u:object_r:container_file_t:s0:c1002,c1014 0 Sep 14 11:03
index.html
```

ACCESSING CONTAINERS

Objective: Describe the basics of networking with containers.

Objective: Remotely connect to services within a container.

NOTE: Rootless container networking is different and uses slirp4netns which provides user-mode networking ("slirp") for network namespaces.

What is used to configure networking for containers run as the root user?

- the Container Networking Interface (CNI) open source project.
<https://github.com/containernetworking/cni>
- The CNI project aims to standardize the network interface for containers in cloud native environments, such as Kubernetes and Red Hat OpenShift Container Platform
- consists of a specification and libraries for writing plugins to configure network interfaces in Linux containers

What other plugins are available ?

<https://github.com/containernetworking/plugins>

Some CNI network plugins, maintained by the containernetworking team. There are others maintained by different teams. Consider

There is an ovs cni plugin among others:

Ovs-cni plugin: <https://github.com/kubevirt/ovs-cni/blob/master/docs/cni-plugin.md>

```
[student@workstation ~]$ cat /etc/cni/net.d/87-podman-bridge.conflist
```

```
[root@workstation ~]# bridge link show
```

```
4: virbr0-nic: <BROADCAST,MULTICAST> mtu 1500 master virbr0 state disabled priority 32
cost 100
```

```
[root@workstation ~]# ip addr show
```

(note missing interface for cni-podman0)

Now, let's run a container and check the networking on the host again:

```
[root@workstation ~]# podman run -d rhsc1/httpd-24-rhel7
```

```
44d328fb7ca9f951e5e165a4eff4860789c446fc35c249844259b9f3320e67fa
```

```
[root@workstation ~]# bridge link show
```

```
4: virbr0-nic: <BROADCAST,MULTICAST> mtu 1500 master virbr0 state disabled priority 32
cost 100
```

```
6: vethf643eab8@virbr0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 master
cni-podman0 state forwarding priority 32 cost 2
```

```
[root@workstation ~]# ip addr show cni-podman0
```

```
5: cni-podman0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state
UP group default qlen 1000
```

```
    link/ether 32:f4:04:21:cd:64 brd ff:ff:ff:ff:ff:ff
```

```
    inet 10.88.0.1/16 brd 10.88.255.255 scope global cni-podman0
```

```
    valid_lft forever preferred_lft forever
```

```
    inet6 fe80::30f4:4ff:fe21:cd64/64 scope link
```

```
    valid_lft forever preferred_lft forever
```

How does ip allocation work with cni ipam “host-local” ?

IPAM=IP Address Management

Check in **/var/lib/cni/networks/podman/**

The IP address assigned to a particular container is listed as a flat text file while its contents matches its networking namespace (or containerID).

The **last_reserved_ip.0** is a “helper” file indicating the last assigned IP address.

```
[root@workstation ~]# cd /var/lib/cni/networks/podman/
[root@workstation podman]# ls
10.88.0.31 10.88.0.33 last_reserved_ip.0 lock
[root@workstation podman]# cat 10.88.0.33
e186fa7ddb77eb96c1017e5b34193c432e10d7227b0d7b060f63bdc9379e7dcb
```

```
[root@workstation ~]# podman ps
CONTAINER ID IMAGE COMMAND CREATED
STATUS PORTS NAMES
e186fa7ddb77 registry.access.redhat.com/rhscsl/httpd-24-rhel7:latest /usr/bin/run-http... 2
minutes ago Up 2 minutes ago pedantic_stonebraker
```

```
[root@workstation ~]# curl 10.88.0.33:8080 | grep title
% Total % Received % Xferd Average Speed TimeTime Time Current
Dload Upload Total Spent Left Speed
100 3985 100 3985 0 0 486k 0 --:--:-- --:--:-- --:--:-- 555k
<title>Test Page for the Apache HTTP Server on Red Hat Enterprise Linux</title>
```

But, does this work from a different machine in our network?

```
[root@workstation ~]# ssh student@bastion
```

```
[student@bastion ~]$
[student@bastion ~]$
[student@bastion ~]$ curl 10.88.0.33:8080 | grep title
% Total % Received % Xferd Average Speed TimeTime Time Current
Dload Upload Total Spent Left Speed
0 0 0 0 0 0 0 0 0 --:--:-- 0:00:14 --:--:-- 0
^C
[student@bastion ~]$
[student@bastion ~]$ nc -v 10.88.0.33 8080
Ncat: Version 7.70 ( https://nmap.org/ncat )
^C
[student@bastion ~]$ nc -v workstation 22
Ncat: Version 7.70 ( https://nmap.org/ncat )
Ncat: Connected to 172.25.250.9:22.
SSH-2.0-OpenSSH_8.0
```

So, the apache process running in the container on workstation is not available outside of workstation.

- Containers running on other hosts aren't connected by podman either.
- Mapping a networking port on the workstation (host) to the container's networking namespace (apache) would be required even for a container run as the root user

SEE Figure 3.4: Basic Linux container networking

MAPPING NETWORK PORTS

How can we connect to a containerized application from the host's network ?

Ports !

```
[student@workstation ~]# podman run -d -p 8888:8080 ubi8 sleep 5000
845072a99bbdca4d0581e0a75326fa72d44c75d1e4f52f63420c1dfc810d1110
[student@workstation ~]# netstat -tunap | grep 8888
tcp        0      0 0.0.0.0:8888          0.0.0.0:*            LISTEN     3439/slirp4netns
```

```
[student@workstation ~]$ ps -ef | grep slirp
student      3439   1 0 06:54 pts/1        00:00:00 /usr/bin/slirp4netns --api-socket
/run/user/1000/libpod/tmp/930f971c0e8c14bdba1605f09f7fd3ea14c9aa0d1384f3a7a348e4181f
a126f1.net --disable-host-loopback --mtu 65520 --enable-sandbox -c -e 3 -r 4 --netns-type=path
/run/user/1000/netns/cni-200e5dfd-dabc-4c02-b2a5-392a7dc4ce87 tap0
```

```
[student@workstation ~]$ podman exec -it 93 /bin/bash
[root@930f971c0e8c /]# yum install iproute -y
[root@930f971c0e8c /]# ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```

    inet 127.0.0.1/8 scope host lo
    valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
    valid_lft forever preferred_lft forever
2: tap0: <BROADCAST,UP,LOWER_UP> mtu 65520 qdisc fq_codel state UNKNOWN group
default qlen 1000
    link/ether 36:c3:4a:58:15:04 brd ff:ff:ff:ff:ff:ff
    inet 10.0.2.100/24 brd 10.0.2.255 scope global tap0
    valid_lft forever preferred_lft forever
    inet6 fe80::34c3:4aff:fe58:1504/64 scope link
    valid_lft forever preferred_lft forever

```

This network exists within a separate networking namespace from the global one. Consider findmnt. Look for nsfs (namespace filesystem):

```

[student@workstation ~]$ podman unshare findmnt | grep nsfs
|      └─/run/user/1000/netns/cni-e74fcbb7-b359-3084-eab8-599b903d2647
      nsfs[net:[4026532260]]                                nsfs                rw,seclabel
[student@workstation ~]$ findmnt | grep nsfs
[student@workstation ~]$

```

Nsfs is a pseudo-filesystem. See the commit in the kernel here:

<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/commit/?id=e149ed2b805fefdccf7ccdfc19eca22fdd4514ac>

NOTE: if run as root, common is the listening process:

```

[root@workstation ~]# podman run -d -p 8081:8080 rhsc1/httpd-24-rhel7
eef01bf84d01da35e289ec27b8dd4ea4aec188114c18daee567fef3d5f22a51d
[root@workstation ~]# netstat -tunap | grep 8081
tcp        0      0 0.0.0.0:8081          0.0.0.0:*             LISTEN      8956/common

```

```

[student@workstation ~]$ curl localhost:8081 | grep title
% Total    % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload  Total   Spent    Left     Speed
100 3985 100 3985 0    0 1945k    0 --:--:-- --:--:-- --:--:-- 1945k

```

<title>Test Page for the Apache HTTP Server on Red Hat Enterprise Linux</title>

Since we are using host's network, the host's firewall will need to be modified:

```
[student@workstation ~]# sudo firewall-cmd --add-port=8081/tcp
```

```
[student@workstation ~]$ ssh bastion
```

Activate the web console with: `systemctl enable --now cockpit.socket`

This system is not registered to Red Hat Insights. See <https://cloud.redhat.com/>

To register this system, run: `insights-client --register`

Last login: Mon Jun 14 12:22:35 2021 from 172.25.250.9

```
[student@bastion ~]$ curl workstation:8081 | grep title
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
		Dload	Upload	Total	Spent	Left	Speed
100	3985	100	3985	0	0	353k	0 --:--:-- --:--:-- --:--:-- 353k

<title>Test Page for the Apache HTTP Server on Red Hat Enterprise Linux</title>

Is Podman 4.0 is moving away from CNI ?

<https://www.redhat.com/sysadmin/podman-new-network-stack>

Of the new features in [Podman v4.0](#), one of the most important is a new network stack, written from scratch in [Rust](#) to support Podman. The new stack is composed of two tools, the [Netavark](#) network setup tool and the [Aardvark DNS](#) server.

Existing containers in nondefault networks cannot be converted to Netavark, and Netavark doesn't support advanced CNI plugins (for example, connecting to Kubernetes networks created using Flannel). To ensure a smooth transition, we will continue to support CNI with Podman, and existing Podman installations will continue to use CNI for networking.

New installations can opt to use CNI by explicitly specifying it via the `containers.conf` configuration file, using the `network_backend` field. CNI and Netavark cannot be used simultaneously in order to avoid conflicts in the configurations the two create.

What is used currently in DO180 classroom ? CNI see:

```
[root@workstation ~]# cat /usr/share/containers/containers.conf
```

[network]

```
# Network backend determines what network driver will be used to set up and tear down
container networks.
# Valid values are "cni" and "netavark".
# The default value is empty which means that it will automatically choose CNI or netavark. If
there are
# already containers/images or CNI networks preset it will choose CNI.
#
# Before changing this value all containers must be stopped otherwise it is likely that
# iptables rules and network interfaces might leak on the host. A reboot will fix this.
#
#network_backend = ""
network_backend = "cni"
```

Are there python modules that can work directly with libpod ?

No. Not directly, but there is a RESTful API available with podman started by running `podman system service tcp:localhost:8080 --log-level=debug --time=0` like:

This could be adapted to python code using the requests module ie
<https://www.redhat.com/sysadmin/podman-python-bash>

There are some efforts to use this api in python here: <https://github.com/containers/podman-py>

A new API is coming with version 2.0 of podman (we are using 1.6.4 in the training env):
<https://www.redhat.com/sysadmin/podmans-new-rest-api>

https://docs.podman.io/en/latest/_static/api.html

CHAPTER 4: MANAGING CONTAINER IMAGES

Where can I find the actual OCI specification ?

OCI spec: <https://opencontainers.org/>

Runtime-spec <https://github.com/opencontainers/runtime-spec>

Image-spec <https://github.com/opencontainers/image-spec>

CONFIGURING REGISTRIES IN PODMAN

How can you block access to certain registries system-wide ?

This should now be possible because of https://bugzilla.redhat.com/show_bug.cgi?id=1787667

..but you can't really - https://bugzilla.redhat.com/show_bug.cgi?id=1811098

(older podman configuration format)

```
[student@workstation ~]$ sudo vi /etc/containers/registries.conf
```

```
[registries.block]
```

```
registries = ['docker.io']
```

(newer podman v2 configuration format)

```
[student@workstation ~]$ sudo vi /etc/containers/registries.conf
```

```
unqualified-search-registries = ["registry.access.redhat.com", "registry.redhat.io", "quay.io"]
```

```
[[registry]]
```

```
prefix = "quay.io"
```

```
location = "quay.io"
```

```
insecure = false
```

blocked = true

(end new format)

```
[student@workstation ~]$ podman pull quay.io/redhattraining/httpd-parent:latest
Trying to pull quay.io/redhattraining/httpd-parent:latest...
Error: initializing source docker://quay.io/redhattraining/httpd-parent:latest: registry quay.io is blocked in /etc/containers/registries.conf or /home/student/.config/containers/registries.conf.d
```

```
[student@workstation ~]$ podman pull quay.io/ajblum/mytest:latest
Trying to pull quay.io/ajblum/mytest:latest...
Error: initializing source docker://quay.io/ajblum/mytest:latest: registry quay.io is blocked in /etc/containers/registries.conf or /home/student/.config/containers/registries.conf.d
```

It's also possible to block registries from a particular namespace. For this, use the location for matching instead of the prefix:

```
unqualified-search-registries = ["registry.access.redhat.com", "registry.redhat.io", "quay.io"]
```

```
[[registry]]
location = "quay.io/ajblum"
insecure = false
blocked = true
```

```
[student@workstation ~]$ podman pull quay.io/ajblum/mytest:latest
Trying to pull quay.io/ajblum/mytest:latest...
Error: initializing source docker://quay.io/ajblum/mytest:latest: registry quay.io is blocked in /etc/containers/registries.conf or /home/student/.config/containers/registries.conf.d
```

Fails as expected, let's try a different

```
[student@workstation ~]$ podman pull quay.io/redhattraining/httpd-parent:latest
Trying to pull quay.io/redhattraining/httpd-parent:latest...
Getting image source signatures
```

Works. So, only a specific namespace can be blocked.

Ok - Looks good, right ? Now, we create a local registries.conf as a non-root user (one that allows access to docker.io):

```
[student@workstation ~]$ mkdir -p ~/.config/containers/  
[student@workstation ~]$ touch ~/.config/containers/registries.conf
```

```
[student@workstation ~]$ podman pull quay.io/ajblum/mytest:latest  
Trying to pull quay.io/ajblum/mytest:latest...  
Getting image source signatures
```

Worked ... eek ! But, really nothing can stop a motivated user to work around this global config. From BZ 1811098 “a local user could still pull an image via curl or by pointing the tools to another path.” https://bugzilla.redhat.com/show_bug.cgi?id=1811098

How to use the registry http api directly ?

```
[root@workstation ~]# podman search quay.io/mytest
```

INDEX	NAME	DESCRIPTION	STARS	OFFICIAL
AUTOMATED				
quay.io	quay.io/ihoukai/mytest		0	
quay.io	quay.io/little_arhat/mytest	test of homu/quay integration	0	
quay.io	quay.io/guenael/mytest		0	
quay.io	quay.io/ajblum/mytest		0	

```
[root@workstation ~]# curl https://quay.io/v2/ajblum/mytest/tags/list  
{ "name": "ajblum/mytest", "tags": [ "1.0", "latest", "2.0", "3.0", "4.0", "5.0" ] }
```

Additional curl troubleshooting: <https://access.redhat.com/articles/3560571>

Docker Registry API docs: <https://docs.docker.com/registry/spec/api/>

For Quay: <https://docs.quay.io/api/swagger/>

repository : List, create and manage repositories.

GET	/api/v1/repository
POST	/api/v1/repository
POST	/api/v1/repository/{repository}/changevisibility
POST	/api/v1/repository/{repository}/changetrust
DELETE	/api/v1/repository/{repository}
GET	/api/v1/repository/{repository}
PUT	/api/v1/repository/{repository}

```
    "error_message": "string",
    "error_type": "string"
}
```

Try it out!

Hide Response

Request URL

```
https://quay.io/api/v1/repository?public=true&namespace=ajblum
```

```
[student@workstation ~]$ curl -L
```

```
"https://quay.io/api/v1/repository?public=true&namespace=ajblum"
```

```
{
  "repositories": [
    {
      "namespace": "ajblum",
      "name": "mytest",
      "description": "",
      "is_public": true,
      "kind": "image",
      "state": "NORMAL",
      "quota": null
    },
    {
      "namespace": "ajblum",
      "name": "myapp",
      "description": null,
      "is_public": true,
      "kind": "image",
      "state": "NORMAL",
      "quota": null
    },
    {
      "namespace": "ajblum",
      "name": "httpd-systemd",
      "description": null,
      "is_public": true,
      "kind": "image",
      "state": "NORMAL",
      "quota": null
    }
  ]
}
```

```
[student@workstation ~]$ curl -Ls
```

```
"https://quay.io/api/v1/repository?public=true&namespace=ajblum" | jq
```

```
 '.repositories[].name'
```

```
"mytest"
```

```
"myapp"
```

```
"httpd-systemd"
"versioned-hello"
"myubi"
"foo"
"helloworld"
"debezium-connector-postgres"
"rhel7-attr"
"hello-openshift"
"myubitest"
"mysigtest"
"do180"
```

A better tool, skopeo:

```
[root@workstation ~]# skopeo inspect docker://quay.io/ajblum/mytest
{
  "Name": "quay.io/ajblum/mytest",
  "Tag": "latest",
  "Digest":
"sha256:6cd0217844a2d778786dcc8c9c948aecc6ca1a36f8f16e5e4bbd4151f7ba5a61",
  "RepoTags": [
    "1.0",
    "latest",
    "2.0",
    "3.0",
    "4.0",
    "5.0"
  ]
}
```

There is also an RFE against podman for searching:

https://bugzilla.redhat.com/show_bug.cgi?id=1757531

Other ways to copy images locally other than podman pull ?

```
[student@workstation ~]$ mkdir /tmp/mytest
[student@workstation ~]$ skopeo copy docker://quay.io/ajblum/mytest:1.0 dir:/tmp/mytest
[student@workstation mytest]$ cd /tmp/mytest/
[student@workstation mytest]$ ls
```

```
[student@workstation mytest]$ cat manifest.json
[student@workstation mytest]$ cat manifest.json | json_reformat
[student@workstation mytest]$ file
a38d7adc1eb9f56b95435dfb6a51d26e225ef0181c0c71f9f8434c79e98aa59f
[student@workstation mytest]$ tar xvf
a38d7adc1eb9f56b95435dfb6a51d26e225ef0181c0c71f9f8434c79e98aa59f
```

```
[student@workstation ~]$ skopeo copy docker://quay.io/ajblum/mytest:1.0
containers-storage:quay.io/ajblum/mytest:1.0
```

```
[student@workstation ~]$ skopeo copy docker://quay.io/ajblum/mytest:1.0
oci-archive:/tmp/mytest/mytest.tar
```

```
[student@workstation ~]$ podman load -i /tmp/mytest/mytest.tar
Getting image source signatures
```

```
[student@workstation ~]$ podman images
REPOSITORY TAG IMAGE ID CREATED SIZE
<none> <none> a6a3e178a6bc 4 months ago 215MB
```

What about registry.redhat.io ?

This works:

```
[student@workstation ~]# skopeo inspect docker://registry.access.redhat.com/rhel
{
  "Name": "registry.access.redhat.com/rhel",
  "Digest":
"sha256:2d215868e282e68998adece762d374ea49d66266d9dee67776eddc80a3d8e168",
  "RepoTags": [
    "7.3-74",
```

But, not this:

```
[student@workstation ~]# skopeo inspect docker://registry.redhat.io/rhel
FATA[0000] unable to retrieve auth token: invalid username/password
```

Skopeo will use the same authentication used by podman.

How can you pull an image using its digest ?

Suppose you are interested in specific images from
registry.access.redhat.com/rhsc1/httpd-24-rhel7

```
[student@workstation storage]$ skopeo inspect
docker://registry.access.redhat.com/rhsc1/httpd-24-rhel7:latest | head -10
{
  "Name": "registry.access.redhat.com/rhsc1/httpd-24-rhel7",
  "Digest":
"sha256:02152fd99c0bcfae06af21301ad92ffa122a46e537465d2b6f064f56e5c0685f",
  "RepoTags": [
    "2.4-170.1638430400-source",
    "2.4-170",
    "2.4-172",
    "2.4-146-source",
    "2.4-136.1614612498",
    "2.4-170.1638430400",
```

Compare to

```
[student@workstation ~]$ skopeo inspect
docker://registry.access.redhat.com/rhsc1/httpd-24-rhel7:2.4-172 | head -10
{
  "Name": "registry.access.redhat.com/rhsc1/httpd-24-rhel7",
  "Digest":
"sha256:ed835f1a45efb7dfd62894274692f494ddb83d1072019ecafc040574cce5886",
  "RepoTags": [
    "2.4-170.1638430400-source",
    "2.4-170",
    "2.4-172",
    "2.4-146-source",
    "2.4-136.1614612498",
    "2.4-170.1638430400",
```

We could make a local copy using the tag “2.4-172” but lets try using this digest:

```
[student@workstation storage]$ podman pull
registry.access.redhat.com/rhsc1/httpd-24-rhel7@sha256:ed835f1a45efb7dfd62894274692f494ddb83d1072019ecafc040574cce5886
```

...SNIP...

fcea1b0658e6a351aec4119d8c9ee2adb725e151536b98aa8c13d4c6b8e8647b

```
[student@workstation storage]$ podman images
```

```
registry.access.redhat.com/rhsc1/httpd-24-rhel7 <none> fcea1b0658e6 2 months ago 329 MB
```

We see later how we can assign a local tag to this image if we want.

REGISTRY AUTHENTICATION

See <https://access.redhat.com/RegistryAuthentication>

- Some container registries require authentication.

```
[student@workstation ~]$ podman pull registry.redhat.io/rhel7
```

```
Trying to pull registry.redhat.io/rhel7...Failed
```

```
error pulling image "registry.redhat.io/rhel7": unable to pull registry.redhat.io/rhel7: unable to pull image: Error determining manifest MIME type for docker://registry.redhat.io/rhel7:latest: unable to retrieve auth token: invalid username/password
```

```
[student@workstation ~]$ podman login -u rhn-support-ablum registry.redhat.io
```

```
Password:
```

```
Login Succeeded!
```

```
[student@workstation ~]$ podman pull registry.redhat.io/rhel7
```

```
Trying to pull registry.redhat.io/rhel7...Getting image source signatures
```

```
Copying blob
```

```
sha256:c9281c141a1bfec06e291d2ad29bfdedfd10a99d583fc0f48d3c26723ebe0761
```

```
...
```

Although this may be deprecated at some point, this doesn't require authentication:

```
[student@workstation ~]$ podman pull registry.access.redhat.com/rhel
```

```
Trying to pull registry.access.redhat.com/rhel...Getting image source signatures
```

```
Skipping fetch of repeat blob
```

```
sha256:c9281c141a1bfec06e291d2ad29bfdedfd10a99d583fc0f48d3c26723ebe0761
```


- Mention service accounts and why they would be a good idea

- <https://access.redhat.com/terms-based-registry/>

-

<https://access.redhat.com/terms-based-registry/#/token/ablum-rhel8-training>

<https://status.redhat.com/incidents/bjqjyxcknf86>

Container Registry Login Outage

Incident Report for Red Hat

Resolved

This incident has been resolved.

Posted about 3 hours ago. Aug 21, 2019 - 09:44 EDT

Update

We are continuing to monitor for any further issues.

Posted about 5 hours ago. Aug 21, 2019 - 07:10 EDT

Monitoring

Web user logins should be working again. We'll continue to monitor the issue.

Posted about 5 hours ago. Aug 21, 2019 - 07:07 EDT

Investigating

We're investigating issues with logins to the container registry.

Posted about 6 hours ago. Aug 21, 2019 - 06:51 EDT

<https://redhat.service-now.com/surl.do?n=INC0930151>

“”

When performing user credential auth to the registry there is a dependence on a restricted party screening service external to Red Hat. In this case that service looks like it had some issues and returned errors and then false export blocks. We need to follow up with the vendor to know what happened on their end. The important point here is that access to the registry through web credentials is not considered a C1 function because of this extra dependency. Resilient, production integration with the registry should always use an auth token that can be generated from the customer portal. The relevant documentation is here:

<https://access.redhat.com/RegistryAuthentication>

Any customer that is using user credentials for access to the registry in a production set up should be referred to this section

<https://access.redhat.com/RegistryAuthentication#registry-service-accounts-for-shared-environments-4>

where it reads:

Registry Service Accounts for Shared Environments

To consume container images from registry.redhat.io in shared environments such as OpenShift, it is recommended for an administrator to use a Registry Service Account, also referred to as authentication tokens, in place of an individual's Customer Portal credentials.

Service Accounts are a mechanism provided to a Customer Portal organization, used exclusively for authenticating to and retrieving content from registry.redhat.io. The use of Service Accounts is encouraged to prevent the need to use Customer Portal credentials on shared systems, in contrast to Customer Portal accounts, Registry Service Accounts are resilient to some security controls applied to Customer Portal accounts, such as mandated password resets.

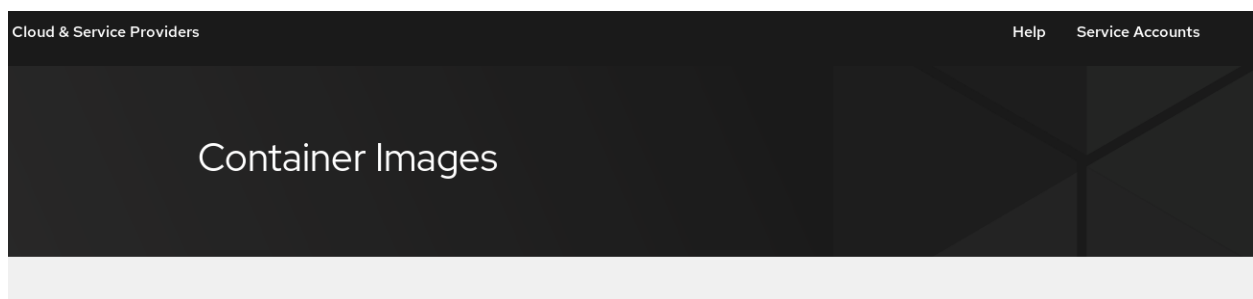
This outage is a good example of why customers should be using service accounts for any resilient registry integration. Any cases that have been open by customers against this outage are an opportunity to better **educate** the customers as to the preferred auth method for our terms based registry.

“”

To login using a service account:

Navigate to <https://catalog.redhat.com/software/containers/explore>

Click on Service Accounts in upper right:



Find your service account (or create one).

Click on Docker Login:

[Registry Service Accounts](#) > ablum-rhel8-training

Token Information

Token Information

OpenShift Secret

Docker Login

Docker Configuration

Run docker login

```
docker login -u='1979710|ablum-rhel8-training'  
-p=eyJhbGciOiJSUzUxMiJ9.eyJzdWIiOiJkZDVjYWU3ODM5MDg0OWU5ODczZDQzMWYyYjAwM2M2YiJ9.ccyliPfdji  
pKaTWVETow_NeZjFamfa7jZ8KYisTMtoF8_Bs5MerXHtYR2nmMXBLKpvIDGL6XUdRvwgABeRBCinAWTy0ePPgnqCvMjU
```

```
[root@workstation ~]# podman login -u='1979710|ablum-rhel8-training'  
-p=eyJhbGciOiJSUzUxMiJ9.eyJzdWIiOiJkZDVjYWU3ODM5MDg0OWU5ODczZDQzMWYyYjAwM2M2YiJ9.ccyliPfdji  
OFKseSYF5e0LiC-pKaTWVETow_NeZjFamfa7jZ8KYisTMtoF8_Bs  
5MerXHtYR2nmMXBLKp... registry.redhat.io
```

```
[root@workstation ~]# podman pull registry.redhat.io/rhel7
```

What is the expiration for authentication access tokens?

Normal username and password authentication with registry.redhat.io will result in the user's info being cached in **/run/user/1000/containers/auth.json**

```
{  
  "auths": {  
    "registry.redhat.io": {  
      "auth": "asdkfjkdasjfjjsadklfj"  
    }  
  }  
}
```

When podman runs commands (like pull) podman will use these credentials to obtain a token that is short-lived (300s). For example,

```
[ablum@badger ~]$ curl -Lv -u 'rhn-support-ablum:password'  
"https://sso.redhat.com/auth/realms/rhcc/protocol/redhat-docker-v2/auth?service=docker  
-registry&client_id=curl&scope=repository:rhel:pull"
```

```
bkYzs3tRYjVETakY","expires_in":300,"issued_at":"2020-05-07T11:28:46Z"}
```

When using service accounts for authentication, you can decode parts.

How to use images from a registry for disconnected customers?

<http://post-office.corp.redhat.com/archives/sbr-containers/2019-August/msg00012.html>

1.) run on a low-side machine (or one that is allowed to connect to the internet):

```
# yum module install container-tools
```

```
# podman login -u rhn-support-ablum registry.redhat.io
```

```
# skopeo copy docker://registry.redhat.io/rhel7:latest  
dir:/root/redhat_io/rhel7
```

```
# cd /root/redhat_io/rhel7 && tar -cvf /root/redhat_io_images/rhel7.tar  
* && cd ..
```

2.) loop for other images needed (or find a different way to sync the registry?)

3.) Create an ISO from tar'd images:

```
# yum install genisoimage
```

```
# mkisofs -o /root/redhat_io.iso /root/redhat_io_images
```

```
# dd if=/root/redhat_io.iso of=/dev/cdrom
```

4.) Carry ISO to disconnected network/system and copy tar'd images from ISO

5.) Follow <https://access.redhat.com/solutions/4175551> to upload image

tar files to satellite using hammer

MANIPULATING CONTAINER IMAGES

Working around container image signatures.

```
[root@workstation ~]# podman images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
registry.redhat.io/rhsc1/httpd-24-rhel7	latest	a1fdc13b5792	8 days ago	324MB
registry.access.redhat.com/rhel	latest	31cd91012c57	10 days ago	214MB
registry.redhat.io/rhel7	latest	31cd91012c57	10 days ago	214MB
registry.lab.example.com/rhel7	7.5-404	e64297b706b7	13 months ago	211MB

```
[student@workstation ~]$ podman save -o httpd.tar
```

```
registry.access.redhat.com/rhsc1/httpd-24-rhel7:latest
```

```
Getting image source signatures
```

```
Checking if image destination supports signatures
```

```
Error: Can not copy signatures to docker-archive:httpd.tar: Storing signatures for docker tar files is not supported
```

Trying to save the ubi8:latest gives the same

```
[student@workstation ~]$ podman save 1264065f6ae8 -o ubi8.tar
```

```
Getting image source signatures
```

```
Checking if image destination supports signatures
```

```
Error: Can not copy signatures to docker-archive:ubi8.tar: Storing signatures for docker tar files is not supported
```

Is this true for all images ?

```
[student@workstation ~]$ podman pull quay.io/ajblum/hello-openshift:latest
```

```
Trying to pull quay.io/ajblum/hello-openshift:latest...
```

```
Getting image source signatures
```

```
Copying blob a3ed95caeb02 done
```

```
Copying blob b30065c58b6f done
```

```
Copying config 7af3297a3f done
```

```
Writing manifest to image destination
```

```
Storing signatures
```

```
7af3297a3fb4487b740ed6798163f618e6eddea1ee5fa0ba340329fcae31c8f6
```

```
[student@workstation ~]$ podman save -o hello-openshift.tar
```

```
quay.io/ajblum/hello-openshift:latest
```

```
Getting image source signatures
```

```
Copying blob 5f70bf18a086 done
```

```
Copying blob da0e4d9121c7 done
```

```
Copying config 7af3297a3f done
```

```
Writing manifest to image destination
```

```
Storing signatures
```

Nope, works. Images from RedHat are signed using the same gpg key rpm packages are signed with. Podman uses the same rpm-gpg key to verify images from redhat.

```
[student@workstation ~]$ cat /etc/containers/policy.json
```

```
[student@workstation ~]$ podman image trust show
```

```
default
```

```
accept
```

```
registry.access.redhat.com signedBy
```

```
security@redhat.com, security@redhat.com
```

```
https://access.redhat.com/webassets/docker/content/sigstore
```

```
registry.redhat.io signedBy
```

```
security@redhat.com, security@redhat.com
```

```
https://registry.redhat.io/containers/sigstore
```

```
insecureAcceptAnything
```

Blog on gpg signatures used with RH images:

<https://developers.redhat.com/blog/2019/10/29/verifying-signatures-of-red-hat-container-images>

Instead, we could copy the image without the signature:

```
[student@workstation ~]$ skopeo copy docker://registry.access.redhat.com/ubi8:latest  
containers-storage:localhost/ubi8:latest --remove-signatures
```

```
Copying blob 028bdc977650 skipped: already exists
```

```
Copying blob 0c673eb68f88 [-----] 0.0b / 0.0b
```

```
Copying config 2fd9e14788 done
```

```
Writing manifest to image destination
```

```
Storing signatures
```

```
[student@workstation ~]$ podman images
```

```
REPOSITORY
```

```
TAG
```

```
IMAGE ID
```

```
CREATED
```

```
SIZE
```

```
registry.access.redhat.com/rhsc1/httpd-24-rhel7 latest 7b8d40facb4 2 weeks ago 330 MB
```

```
registry.access.redhat.com/rhel7 latest e2c37c467077 3 weeks ago 216 MB
```

```
localhost/ubi8 latest 2fd9e1478809 3 weeks ago 225 MB
```

```
registry.access.redhat.com/ubi8 latest 2fd9e1478809 3 weeks ago 225 MB
```

```
[student@workstation ~]$ podman save -o ubi8.tar ubi8:latest
```

```
Getting image source signatures
```

Copying blob 77bf63677b0c done
Copying blob 3edb0d97db5c done
Copying config 2fd9e14788 done
Writing manifest to image destination
Storing signatures

This should be avoided in future versions of podman
<https://github.com/containers/podman/pull/7956>

How can you remove images using `podman system prune` ?

From man 1 podman-system-prune:

podman system prune removes all unused containers (both dangling and unreferenced), pods and optionally, volumes from local storage.

```
[student@workstation ~]$ podman run -d -v myvol:/var/www/html rhsci/httpd-24-rhel7  
6e5f9dc5474c4853e0bf01e508ba2471bbc190830a1c51b06204dd436846f07e
```

```
[student@workstation ~]$ podman volume list  
DRIVER  VOLUME NAME  
local   myvol
```

```
[student@workstation ~]$ podman volume inspect myvol  
[  
  {  
    "Name": "myvol",  
    "Driver": "local",  
    "Mountpoint": "/home/student/.local/share/containers/storage/volumes/myvol/_data",  
    "CreatedAt": "2021-09-22T16:03:27.230242031-04:00",  
    "Labels": {  
  
    },  
    "Scope": "local",  
    "Options": {  
  
    }  
  }  
]
```

]

```
[student@workstation ~]$ podman inspect 6e --format '{{.Mounts}}'
[{{volume myvol /home/student/.local/share/containers/storage/volumes/myvol/_data
/var/www/html local [noexec nosuid nodev rbind] true rprivate}}
```

```
[student@workstation ~]$ podman exec -it 6e /bin/bash
```

```
bash-4.2$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
fuse-overlayfs	9.9G	6.3G	3.7G	64%	/
tmpfs	64M	0	64M	0%	/dev
tmpfs	580M	100K	580M	1%	/etc/hosts
shm	63M	0	63M	0%	/dev/shm
/dev/vda3	9.9G	6.3G	3.7G	64%	/var/www/html
tmpfs	2.9G	0	2.9G	0%	/sys/fs/cgroup
devtmpfs	2.8G	0	2.8G	0%	/dev/tty
tmpfs	2.9G	0	2.9G	0%	/proc/acpi
tmpfs	2.9G	0	2.9G	0%	/proc/scsi
tmpfs	2.9G	0	2.9G	0%	/sys/firmware
tmpfs	2.9G	0	2.9G	0%	/sys/fs/selinux

```
bash-4.2$ touch /var/www/html/index.html
```

```
bash-4.2$ exit
```

```
exit
```

```
[student@workstation ~]$ ls
```

```
/home/student/.local/share/containers/storage/volumes/myvol/_data
index.html
```

```
[student@workstation ~]$ podman system prune --volumes
```

WARNING! This will remove:

- all stopped containers
- all volumes not used by at least one container
- all stopped pods
- all dangling images
- all build cache

Are you sure you want to continue? [y/N] y

Deleted Pods

Deleted Containers

Deleted Volumes

```
[student@workstation ~]$
```

```
[student@workstation ~]$
```



```
[student@workstation ~]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
6e5f9dc5474c	registry.access.redhat.com/rhsc	/usr/bin/run-http... 2	minutes ago
Up 2 minutes ago	youthful_mahavira		

```
[student@workstation ~]$ podman stop 6e
```

```
6e5f9dc5474c4853e0bf01e508ba2471bbc190830a1c51b06204dd436846f07e
```

```
[student@workstation ~]$ podman system prune --volumes
```

WARNING! This will remove:

- all stopped containers
- all volumes not used by at least one container
- all stopped pods
- all dangling images
- all build cache

Are you sure you want to continue? [y/N] y

Deleted Pods

Deleted Containers

```
6e5f9dc5474c4853e0bf01e508ba2471bbc190830a1c51b06204dd436846f07e
```

Deleted Volumes

myvol

```
[student@workstation ~]$ podman volume list
```

```
[student@workstation ~]$
```

```
[student@workstation ~]$ mkdir mydir
```

```
[student@workstation ~]$ podman unshare chown 1001:1001 mydir
```

```
[student@workstation ~]$ podman unshare chcon -t container_file_t mydir
```

```
[student@workstation ~]$ podman run -d -v /home/student/mydir:/var/www/html  
rhsc/httpd-24-rhel7
```

```
5e0175ff9761696c20887cd450a4500552e3ed877cfd62fc4f43fb4c0d39c03f
```

```
[student@workstation ~]$ podman volume list
```

```
[student@workstation ~]$
```

```
[student@workstation ~]$ podman inspect 5e --format '{{.Mounts}}'
```

```
[{"bind": "/home/student/mydir:/var/www/html", "rbind": true, "private": false}]
```

```
[student@workstation ~]$ ls -ldZ /home/student/mydir
```

```
drwxrwxr-x. 2 101000 101000 unconfined_u:object_r:container_file_t:s0 6 Sep 22 16:07  
/home/student/mydir
```

```
[student@workstation ~]$ podman stop 5e
```

```
5e0175ff9761696c20887cd450a4500552e3ed877cfd62fc4f43fb4c0d39c03f
```

```
[student@workstation ~]$ podman system prune --volumes
```

WARNING! This will remove:

- all stopped containers
- all volumes not used by at least one container
- all stopped pods
- all dangling images
- all build cache

Are you sure you want to continue? [y/N] y

Deleted Pods

Deleted Containers

5e0175ff9761696c20887cd450a4500552e3ed877cfd62fc4f43fb4c0d39c03f

Deleted Volumes

See also https://bugzilla.redhat.com/show_bug.cgi?id=1811570#c15 :

“”

podman system prune should NOT be removing buildah containers/images.

“”

Modifying images using podman commit

- One strategy would be to run a container from an image, modify it, and then commit it to a new image
- Uses `podman commit`

```
[student@workstation ~]$ podman run -it ubi8 /bin/bash
```

```
[root@160bfbd24ae2 /]# echo "ablum was here" > testfile
```

```
[root@160bfbd24ae2 /]# vi /etc/motd
```

```
This is the message for today
```

```
[root@160bfbd24ae2 ~]# rm /root/.cshrc
```

```
rm: remove regular file '/root/.cshrc'? yes
```

```
[root@160bfbd24ae2 /]# exit
```

```
exit
```

```
[student@workstation ~]$ podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
PORTS	NAMES	IS INFRA		
160bfbd24ae2	registry.access.redhat.com/ubi8:latest	/bin/bash		3
minutes ago	Exited (0) 18 seconds ago	nice_faraday		

```
[student@workstation ~]$ podman diff 160
```

A /ablum_was_here
C /etc
C /etc/motd
C /root
A /root/.bash_history
D /root/.cshrc

```
[student@workstation ~]$ podman inspect 160 --format '{{.GraphDriver.Data.UpperDir}}'
[student@workstation ~]$ ls -lR
/home/student/.local/share/containers/storage/overlay/13ff5bb892ebf396a0a7402b93d61f1d0c0df1dccf81a3b5d29de36283723f52/diff
[student@workstation ~]$ ls -l
/home/student/.local/share/containers/storage/overlay/13ff5bb892ebf396a0a7402b93d61f1d0c0df1dccf81a3b5d29de36283723f52/diff/root/.cshrc
c-----, 2 student student 0, 0 Jun 29 08:27
/home/student/.local/share/containers/storage/overlay/13ff5bb892ebf396a0a7402b93d61f1d0c0df1dccf81a3b5d29de36283723f52/diff/root/.cshrc
```

```
[student@workstation ~]$ podman commit 584e167bd18c ubi8:modified
```

```
[student@workstation ~]$ podman images
REPOSITORY          TAG      IMAGE ID      CREATED      SIZE
localhost/rhel7      modified d48953c14c05  2 minutes ago 211MB
registry.lab.example.com/rhel7 latest    e64297b706b7 13 months ago 211MB
```

```
[student@workstation ~]$ podman run d6 ls -l /root/.cshrc
ls: cannot access '/root/.cshrc': No such file or directory
[student@workstation ~]$ podman run d6 cat /etc/motd
this is the message for today
[student@workstation ~]$ podman run d6 cat /ablum_was_here
[student@workstation ~]$ podman run d6 ls -l /ablum_was_here
-rw-r--r--. 1 root root 0 Jun 29 12:27 /ablum_was_here
```

```
[student@workstation ~]$ mkdir ~/mydata2
```

```
[student@workstation ~]$ chcon -Rv -t container_file_t ~/mydata2
changing security context of '/mydata2'
```

```
[root@workstation /]# podman run -it -v /home/student/mydata2:/opt rhel7 /bin/bash
[root@ee9d566b0357 /]# df /opt
Filesystem    1K-blocks    Used Available Use% Mounted on
/dev/vda1     104845184 5596740 99248444   6% /opt
```

```
[root@ee9d566b0357 /]# touch /opt/important
[root@ee9d566b0357 /]# exit
exit
```

```
[root@workstation /]# podman ps -a
CONTAINER ID  IMAGE                                COMMAND                                CREATED
STATUS        PORTS  NAMES                                IS INFRA
ee9d566b0357  registry.lab.example.com/rhel7:latest /bin/bash                            41 seconds ago
Exited (0) 6 seconds ago    elastic_johnson  false
6b25e5020b56  localhost/rhel7:modified            tail /var/log/yum.l... 7 minutes ago  Exited
(0) 7 minutes ago    modest_gates    false
584e167bd18c  registry.lab.example.com/rhel7:latest /bin/bash                            9 minutes ago  Exited
(0) 9 minutes ago    modest_mestorf  false
[root@workstation /]# podman diff ee9d566b0357
C /root
A /root/.bash_history
```

- NOTE /opt/important is NOT listed in diff
- Use `podman inspect` to see the mounts (volumes)

```
    "Mounts": [
{
    "destination": "/opt",
    "type": "bind",
    "source": "/home/student/mydata1",
    "options": [
    "rbind",
    "rw",
    "rprivate"
```

PUBLISHING IMAGES TO A REGISTRY

How to resolve error “image is signed or the destination specifies a digest” when pushing to quay.io ?

```
[ablum@badger ~]$ podman push 52de04277b39 quay.io/ajblum/mytest:latest
```

Getting image source signatures

Checking if image destination supports signatures

Error: Copying this image requires changing layer representation, which is not possible (image is signed or the destination specifies a digest)

This fails with new versions of podman

```
[ablum@badger ~]$ podman version
```

Version: 3.4.4

API Version: 3.4.4

Go Version: go1.16.8

Built: Wed Dec 8 15:45:07 2021

OS/Arch: linux/amd64

```
[ablum@badger ~]$ podman push 52de04277b39 quay.io/ajblum/mytest:latest
```

--remove-signatures

Copying blob c8013a2772b6 done

Copying blob 7699752e6ed6 done

Copying config 52de04277b done

Writing manifest to image destination

Storing signatures

Blog on gpg signatures used with RH images:

<https://developers.redhat.com/blog/2019/10/29/verifying-signatures-of-red-hat-container-images>

See also `podman image trust show`

Also, this might cause a different issue on quay which is enforcing strict schema standard related to signatures. It will cause a problem for older clients who try to pull:

<https://issues.redhat.com/browse/PROJQUAY-3285> The workaround is to add **--format v2s1** when pulling or copying using those older clients

How would a local image be shared back to a private registry ?

``podman push``

```
[root@workstation /]# podman push d48953c14c05  
registry.lab.example.com/rhel7:modified
```

Getting image source signatures

Copying blob

sha256:24a5c6254cd9693d64581b6f3df5e4ee551cfd5429cf25301d12afa82ac91037

200.88 MB / 200.88 MB [=====]
41s

```
[root@workstation /]# podman push d48953c14c05 registry.redhat.io/rhel7
```

Getting image source signatures

Copying blob

sha256:24a5c6254cd9693d64581b6f3df5e4ee551cfd5429cf25301d12afa82ac91037

8 B / 200.88 MB [>-----] 0s

Error copying image to the remote destination: Error writing blob: Error initiating layer upload to /v2/rhel7/blobs/uploads/ in registry.redhat.io: error parsing HTTP 403 response body: invalid character '<' looking for beginning of value: "<HTML><HEAD>\n<TITLE>Access Denied</TITLE>\n</HEAD><BODY>\n<H1>Access Denied</H1>\n\nYou don't have permission to access

"http://registry.redhat.io/v2/rhel7/blobs/uploads/" on this

server.<P>\nReference #18.ae8d4017.1564688844.35231e6\n</BODY>\n</HTML>\n"

What about pushing to public registries ?

<https://cloud.docker.com/repository/list>

```
[student@workstation mytmp]$ podman run -it 272209ff0ae5 /bin/bash
```

```
[root@11001ec73d8d /]# echo test > ablum_was_here
```

```
[root@11001ec73d8d /]# rm /etc/motd
```

rm: remove regular empty file '/etc/motd'? **yes**

```
[root@11001ec73d8d /]# echo '#ablum_was_here' >> /etc/hosts
```

```
[root@11001ec73d8d /]# cat /etc/hosts
```

```
[root@11001ec73d8d /]# exit
```

Exit

```
[student@workstation mytmp]$ podman ps -a | head -2
```

CONTAINER ID	IMAGE	STATUS	PORTS	NAMES	COMMAND	CREATED
11001ec73d8d	registry.access.redhat.com/ubi8:latest	minutes ago	Exited (0) 22 seconds ago		/bin/bash nervous_banach	3

```
[student@workstation mytmp]$ podman diff 1100
```

```
A /ablum_was_here
```

```
C /root
```

```
A /root/.bash_history
```

```
C /etc
```

```
D /etc/motd
```

```
[student@workstation mytmp]$ podman commit 1100 mytest:1.0
```

```
Getting image source signatures
```

```
Copying blob 1a6543399d61 skipped: already exists
```

```
Copying blob f0a77c369efd skipped: already exists
```

```
Copying blob 96f1a9906488 done
```

```
Copying config 085e96ca8c done
```

```
Writing manifest to image destination
```

```
Storing signatures
```

```
085e96ca8c163fdd011c2d246984b3cb3f781c90e04ff2fa381c646d664eb417
```

```
[student@workstation mytmp]$
```

```
[student@workstation mytmp]$
```

```
[student@workstation mytmp]$ podman login quay.io
```

```
Username: ajblum
```

```
Password:
```

```
Login Succeeded!
```

```
[student@workstation mytmp]$
```

```
[student@workstation mytmp]$ podman push mytest:1.0 quay.io/ajblum/mytest:13.0
```

```
Getting image source signatures
```

```
Copying blob f0a77c369efd done
```

```
Copying blob 96f1a9906488 done
```

```
Copying blob 1a6543399d61 done
```

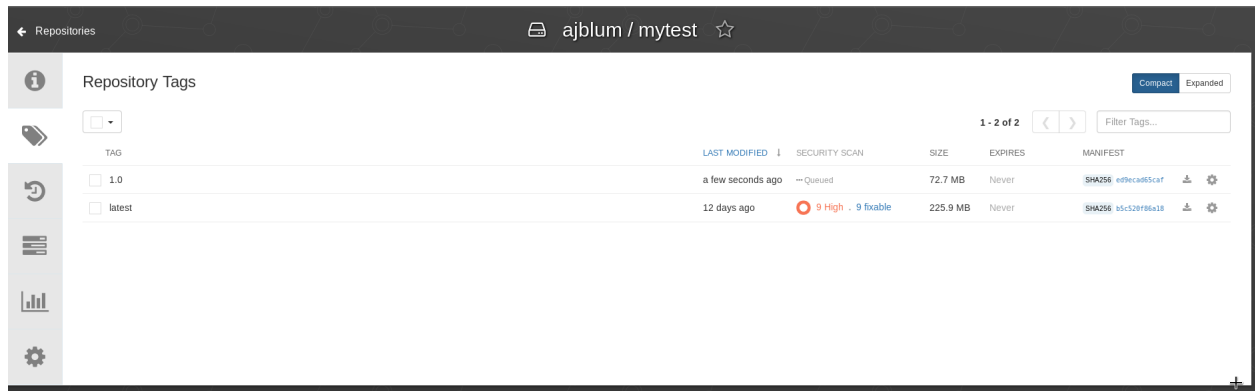
```
Copying config 085e96ca8c done
```

```
Writing manifest to image destination
```

```
Copying config 085e96ca8c done
```

```
Writing manifest to image destination
```

```
Storing signatures
```



```
[student@workstation mytmp]$ podman rmi mytest:1.0
```

```
Untagged: localhost/mytest:1.0
```

```
Deleted: 085e96ca8c163fdd011c2d246984b3cb3f781c90e04ff2fa381c646d664eb417
```

```
[student@workstation mytmp]$
```

```
[student@workstation mytmp]$
```

```
[student@workstation mytmp]$ podman images
```

```
REPOSITORY          TAG      IMAGE ID      CREATED      SIZE
(missing mytest:1.0)
```

```
[student@workstation mytmp]$ podman run quay.io/ajblum/mytest:13.0 cat
abum_was_here
```

```
Trying to pull quay.io/ajblum/mytest:13.0...
```

```
Getting image source signatures
```

```
Copying blob ebd715ce2661 skipped: already exists
```

```
Copying blob b2af222bc479 skipped: already exists
```

```
Copying blob e1bb39f7b07a done
```

```
Copying config 085e96ca8c done
```

```
Writing manifest to image destination
```

```
Storing signatures
```

```
test
```

What services are running on a registry ?

```
[root@workstation ~]# ssh root@registry.lab.example.com
```

```
[root@services ~]# netstat -tunap | grep -i 443
```

```
tcp      0      0 172.25.250.13:443  0.0.0.0:*        LISTEN      3247/registry
```



```
[root@services ~]# ps -ef | grep 3247
```

```
root    3247   1 0 14:27 ?    00:00:00 /usr/bin/registry serve
/etc/docker-distribution/registry/config.yml
root    4041  3993 0 15:16 pts/0    00:00:00 grep --color=auto 3247
```

```
[root@services ~]# systemctl | grep docker
```

```
docker-distribution.service                                loaded active running v2
Registry server for Docker
```

```
[root@services ~]# rpm -qf /usr/lib/systemd/system/docker-distribution.service
```

```
docker-distribution-2.6.2-2.git48294d9.el7.x86_64
```

```
[root@services ~]# rpm -qc docker-distribution
```

```
/etc/docker-distribution/registry/config.yml
```

```
version: 0.1
```

```
log:
```

```
fields:
```

```
    service: registry
```

```
storage:
```

```
    cache:
```

```
    layerinfo: inmemory
```

```
    filesystem:
```

```
    rootdirectory: /var/lib/registry
```

```
    delete:
```

```
    enabled: true
```

```
http:
```

```
    addr: registry.lab.example.com:443
```

```
    host: https://registry.lab.example.com
```

```
    tls:
```

```
    certificate: /etc/pki/tls/certs/example.com.crt
```

```
    key: /etc/pki/tls/private/example.com.key
```

```
[root@services ~]# rpm -qd docker-distribution
```

```
/usr/share/doc/docker-distribution-2.6.2/AUTHORS
```

```
/usr/share/doc/docker-distribution-2.6.2/CONTRIBUTING.md
```

```
/usr/share/doc/docker-distribution-2.6.2/LICENSE
```

```
/usr/share/doc/docker-distribution-2.6.2/MAINTAINERS
```

```
/usr/share/doc/docker-distribution-2.6.2/README.md
```

<https://github.com/docker/distribution/blob/master/docs/spec/api.md>

NOTE1: Red Hat Quay is the future !

NOTE2: skopeo is a better tool (vs podman pull/push/save/load) for sync'ing content around

CHAPTER 5 CREATING CUSTOM IMAGES

BUILDING CUSTOM CONTAINER IMAGES WITH DOCKERFILES

Why do I get permission denied building images after running su - (or sudo su -) ?

```
[root@workstation echo1]# podman build -t mytest:1.0 .
```

```
STEP 1: FROM registry.redhat.io/rhel7:latest
```

```
Error: error creating build container: Error initializing source
```

```
docker://registry.redhat.io/rhel7:latest: unable to retrieve auth token: invalid username/password:
```

```
unauthorized: Please login to the Red Hat Registry using your Customer Portal credentials.
```

```
Further instructions can be found here: https://access.redhat.com/RegistryAuthentication
```

Check:

```
[root@workstation echo1]# echo ${XDG_RUNTIME_DIR}
```

```
[root@workstation echo1]# man pam_systemd
```

“ On login, this module — in conjunction with systemd-logind.service — ensures the following:

1. If it does not exist yet, the user runtime directory /run/user/\$UID is either created or mounted as new "tmpfs" file system with quota applied, and its ownership changed to the user that is logging in.

2. The \$XDG_SESSION_ID environment variable is initialized. If auditing is available and pam_loginuid.so was run before this module (which is highly recommended), the variable is initialized from the auditing session id (/proc/self/sessionid). Otherwise, an independent session counter is used.”

Workaround is just pass the auth file path in the build command:

```
[root@workstation echo1]# podman build -t mytest:1.0 --authfile  
/run/user/0/containers/auth.json .
```

STEP 1: FROM registry.redhat.io/rhel7:latest

Getting image source signatures

ENTRYPOINT VS CMD

Understanding this with example:

```
[root@workstation ~]# mkdir echo1  
[root@workstation ~]# cd echo1  
[root@workstation echo1]# vim Containerfile  
FROM ubi8  
ENTRYPOINT ["/usr/bin/echo", "Hello world!"]  
[root@workstation date1]# podman build -t ubi8:echo1 .  
[root@workstation echo1]# podman run ubi8:echo1  
Hello world!  
[root@workstation echo1]# podman inspect 3a97b163f35a --format  
'{{.ContainerConfig.Entrypoint}}'  
[/usr/bin/echo Hello world!]
```

Now, lets add in a CMD instruction:

```
[root@workstation echo1]# vim Containerfile
```

```
FROM ubi8
ENTRYPOINT ["/usr/bin/echo"]
CMD ["Hello world!"]
```

```
[root@workstation echo1]# podman build -t ubi8:echo2 .
[root@workstation echo1]# podman run ubi8:echo2
Hello world!
```

Let's override the CMD built in the image when running these two, just to understand the impact they have:

```
[root@workstation echo1]# podman run rhel7:echo2 hello foo
hello foo
```

```
[root@workstation echo1]# podman run rhel7:echo1 hello foo
Hello world! hello foo
```

```
[root@workstation echo1]# podman run rhel7:echo2 -ne "hello foo"
hello foo[root@workstation echo1]#
```

```
[root@workstation echo1]# podman run rhel7:echo1 -ne "hello foo"
Hello world! -ne hello foo
```

```
[root@workstation echo1]# podman ps -a --no-trunc --format 'table {{.Image}}
{{.Command}}'
IMAGE          COMMAND
localhost/rhel7:echo1.1  /usr/bin/echo -ne hello foo
localhost/rhel7:echo1  /usr/bin/echo Hello world! -ne hello foo
```

Notice how the commands built vary due to the differences in the entrypoint.

EXEC vs SHELL

Now, let's understand "exec" form vs "shell" form.

```
[root@workstation echo1]# vim Dockerfile
FROM rhel7
```

```
ENV FOO "Hello World!"
ENTRYPOINT ["/usr/bin/echo","$FOO"]
```

```
[root@workstation echo1]# podman build -t rhel7:echo1.2 .
[root@workstation echo1]# podman run rhel7:echo1.2
$FOO
```

Did it do what we were expecting ? Why not ?
Let's override the entrypoint and check out the environment:

```
[root@workstation echo1]# podman run --entrypoint "env" rhel7:echo1.2
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
TERM=xterm
HOSTNAME=8102a385729a
container=oci
FOO=Hello World!
HOME=/root
```

The reason ? "Exec" form vs "shell" form. With exec form, the entrypoint executable isn't run inside a shell. Environment variables are available for use inside a shell. So, let's change to use the shell form and observe the behavior:

```
[root@workstation echo1]# vim Dockerfile
FROM rhel7
ENV FOO "Hello World!"
ENTRYPOINT /usr/bin/echo $FOO
[root@workstation echo1]# podman build -t rhel7:echo1.3 .
[root@workstation echo1]# podman run rhel7:echo1.3
Hello World!
```

```
[root@workstation echo1]# podman ps -a --no-trunc
```

CONTAINER ID	IMAGE	COMMAND
CREATED	STATUS	PORTS NAMES IS
INFRA		
173d53f94eb8b48149407456512800f27355c925302e9ce9bba789e5d501ea9a		
localhost/rhel7:echo1.4	/bin/sh -c /usr/bin/echo \$FOO	About a minute ago Exited
(0) About a minute ago	nifty_mccarthy	false

Notice how the entrypoint is wrapped with /bin/sh -c " ". This allows the shell to expand the variable with the Hello World! string.

Also, now lets see what happens when we pass cmd arguments to the podman run:

```
[root@workstation echo1]# podman run rhel7:echo1.3 whoot  
Hello World!  
[root@workstation echo1]# podman run rhel7:echo1.3 -ne  
Hello World!
```

Notice how the CMD is ignored here when using the shell form for entrypoint

So, what's preferred exec or shell ? exec preferred in most cases.

Exec PRO: natively allows for the executable to handle signals sent to it

Shell PRO: prevents any CMD from affecting the way the entrypoint runs

Stopping containers:: If you are using the shell format, start the ENTRYPOINT with the exec command or the process wont be able to handle the SIGTERM and then a SIGKILL will be sent.

RUN AND LAYERS

```
[root@workstation ~]# mkdir myls  
[root@workstation myls]# vim Dockerfile  
FROM rhel7  
ENTRYPOINT ["/usr/bin/ls"]  
RUN touch /var/tmp/data  
[root@workstation myls]# podman build -t rhel7:mys1 .  
[root@workstation myls]# podman run rhel7:mys1 /var/tmp/  
data  
[root@workstation myls]# podman run rhel7:mys1 -l /var/tmp/  
total 0  
-rw-r--r--. 1 root root 0 Aug 15 17:02 data
```

```
[root@workstation myls]# vim Dockerfile
```

```
FROM rhel7  
ENTRYPOINT ["/usr/bin/ls"]  
RUN touch /var/tmp/data  
RUN touch /var/tmp/data1  
RUN touch /var/tmp/data2  
RUN touch /var/tmp/data3  
RUN touch /var/tmp/data4  
RUN touch /var/tmp/data5  
[root@workstation myls]# podman build -t rhel7:mys1.1 .
```

```
[root@workstation myls]# podman run rhel7:mys1.1 -l /var/tmp/  
total 0
```

```
-rw-r--r--. 1 root root 0 Aug 15 17:02 data  
-rw-r--r--. 1 root root 0 Aug 15 17:03 data1  
-rw-r--r--. 1 root root 0 Aug 15 17:03 data2  
-rw-r--r--. 1 root root 0 Aug 15 17:03 data3  
-rw-r--r--. 1 root root 0 Aug 15 17:03 data4  
-rw-r--r--. 1 root root 0 Aug 15 17:03 data5
```

```
[root@workstation myls]# podman inspect rhel7:mys1.1 --format  
'{{.GraphDriver.Data.LowerDir}}'
```

```
"/var/lib/containers/storage/overlay/bac6c053b9c1da5358c139ebdf4c16560e6db6a64f09592fcc  
8b1054b195b059/diff:/var/lib/containers/storage/overlay/e6064a9d0531666c471742660e06913  
105e1de507ba72b2692ca86bd4ea145dc/diff:/var/lib/containers/storage/overlay/7730fed410205  
03148ad5f44d8c0b9e324ae6f8c3ad47712986f6f6edb5c3e52/diff:/var/lib/containers/storage/ove  
rlay/b98054dd86a3c833e6d55689e5849fb09590484a6e3cddcbbe87c39a37a86515/diff:/var/lib/c  
ontainers/storage/overlay/66e7744fbe7bd7ef81ae7ede769dbb35e5390923755c323380ded1e7  
bdf0cf19/diff:/var/lib/containers/storage/overlay/f9c4bab32ab9a1e25673c7fc3f9f88386767f69e0  
db0157b775ffb9477e4b267/diff:/var/lib/containers/storage/overlay/3a5127d99f58d4abfcd3dc42c  
6cdf3bbe7dd3a6d9140a92eb8cd54b94997ff82/diff:/var/lib/containers/storage/overlay/3451cffb7  
8092cbf2877e44b1fa2774cae891125b5752c5e1c02303bc4ab61a4/diff",
```

```
[root@workstation myls]# podman inspect rhel7:mys1 --format  
'{{.GraphDriver.Data.LowerDir}}'
```

```
"/var/lib/containers/storage/overlay/f9c4bab32ab9a1e25673c7fc3f9f88386767f69e0db0157b775  
ffb9477e4b267/diff:/var/lib/containers/storage/overlay/3a5127d99f58d4abfcd3dc42c6cdf3bbe7d  
d3a6d9140a92eb8cd54b94997ff82/diff:/var/lib/containers/storage/overlay/3451cffb78092cbf287  
7e44b1fa2774cae891125b5752c5e1c02303bc4ab61a4/diff",
```

To avoid these extra layers, combine RUN instructions into one:

```
[root@workstation myls]# vim Dockerfile
```

```
FROM rhel7
```

```
ENTRYPOINT ["/usr/bin/ls"]
```

```
RUN touch /var/tmp/data && \  
    touch /var/tmp/data1 && \  
    touch /var/tmp/data2 && \  
    touch /var/tmp/data3 && \  
    touch /var/tmp/data4 && \  
    touch /var/tmp/data5
```

```
[root@workstation myls]# podman build -t rhel7:mysls1.2 .
[root@workstation myls]# podman inspect rhel7:mysls1.2 --format
'{{.GraphDriver.Data.LowerDir}}'
/var/lib/containers/storage/overlay/f9c4bab32ab9a1e25673c7fc3f9f88386767f69e0db0157b775ff
b9477e4b267/diff:/var/lib/containers/storage/overlay/3a5127d99f58d4abfcd3dc42c6cdf3bbe7dd
3a6d9140a92eb8cd54b94997ff82/diff:/var/lib/containers/storage/overlay/3451cffb78092cbf2877
e44b1fa2774cae891125b5752c5e1c02303bc4ab61a4/diff
```

When building we are also keeping around these intermediate layers:

```
[root@workstation myls]# podman images -a
```

To build without saving the intermediate (build) layers:

```
$ podman build --layers=false -t <name:tag> <dir>
```

```
[root@workstation myls]# man podman-build
--layers
```

Cache intermediate images during the build process (Default is true).

Can you squash the excessive layers when building ?

Yes. Layers can be squashed using **--squash** or **--squash-all**

```
[student@workstation myls]$ podman build --help | grep squash
--squash                squash newly built layers into a single new layer
--squash-all            Squash all layers into a single layer
```

```
[student@workstation myls]$ cat Containerfile
FROM registry.redhat.io/ubi8:latest
ENTRYPOINT ["/usr/bin/ls"]
RUN touch /var/tmp/data0
RUN touch /var/tmp/data1
RUN touch /var/tmp/data2
RUN touch /var/tmp/data3
RUN touch /var/tmp/data4
RUN touch /var/tmp/data5
RUN touch /var/tmp/data6
```



```
[student@workstation myls]$ podman build -t ubi8:squash --squash .
STEP 1/9: FROM registry.redhat.io/ubi8:latest
STEP 2/9: ENTRYPOINT ["/usr/bin/ls"]
STEP 3/9: RUN touch /var/tmp/data0
STEP 4/9: RUN touch /var/tmp/data1
STEP 5/9: RUN touch /var/tmp/data2
STEP 6/9: RUN touch /var/tmp/data3
STEP 7/9: RUN touch /var/tmp/data4
STEP 8/9: RUN touch /var/tmp/data5
STEP 9/9: RUN touch /var/tmp/data6
COMMIT ubi8:squash
Getting image source signatures
Copying blob b38cb9259677 skipped: already exists
Copying blob 23e15b9ab3f0 skipped: already exists
Copying blob f4cb19500042 done
Copying config 1a958042d3 done
Writing manifest to image destination
Storing signatures
--> 1a958042d30
Successfully tagged localhost/ubi8:squash
1a958042d30d08789a566e09578d503d300b0dcb0e0b1b03ed39aaff885b12e4
```

```
[student@workstation myls]$ podman inspect ubi8:squash --format
'{{.GraphDriver.Data.LowerDir}}'
/home/student/.local/share/containers/storage/overlay/f4a999c201294a0a171618e413f1ea7628
562e96d8ac148e00708e421aee56a8/diff:/home/student/.local/share/containers/storage/overlay/
b38cb92596778e2c18c2bde15f229772fe794af39345dd456c3bf6702cc11eef/diff
```

NOTE: There is some storage savings you will get from keeping some of the common layers shared across your container runtime. If you use --squash-all then you will be left with no shared layers missing out on page cache and potentially increasing the overall storage use on your runtime host.

```
[student@workstation myls]$ podman build -t ubi8:squashall --squash-all .
[student@workstation myls]$ podman inspect ubi8:squashall --format
'{{.GraphDriver.Data.LowerDir}}'
<no value>
```

All the data in this case is in the image's UpperDir which will become a unique LowerDir when running the squashall image.

How can you use “OR” logic instead of “AND” logic like the &&

```
[root@workstation ~]#  
[root@workstation ~]# mkdir test1  
[root@workstation ~]# cd test1/  
[root@workstation test1]# vim Dockerfile  
FROM rhel7  
ENTRYPOINT ["/usr/bin/ls"]  
RUN touch /var/tmp/test1 && touch /var/tmp/test2  
RUN touch /var/tmp/foo/bar/test3 || touch /var/tmp/test3  
[root@workstation test1]# podman build -t myls:1.0 .  
STEP 1: FROM rhel7  
Getting image source signatures  
Skipping fetch of repeat blob  
sha256:00f17e0b37b0515380a4aece3cb72086c0356fc780ef4526f75476bea36a2c8b  
Skipping fetch of repeat blob  
sha256:305d73a95c8fece2b53a34e040df1c97eb6b7f7cc4e0a7933465f0b7325e3d72  
Copying config  
sha256:55a1f4beaf8e2d27982b38e3ecfd458c66753cbfd3a09bcf562877fe60255157  
6.44 KB / 6.44 KB [=====]  
0s  
Writing manifest to image destination  
Storing signatures  
STEP 2: ENTRYPOINT ["/usr/bin/ls"]  
ERRO[0002] HOSTNAME is not supported for OCI image format, hostname 9d3d66a8bfcc will  
be ignored. Must use `docker` format  
--> 9daeef6d672b0d61990b3ce73a025ae19437d52bee563bbeb45be9527bd6eeb7  
STEP 3: FROM 9daeef6d672b0d61990b3ce73a025ae19437d52bee563bbeb45be9527bd6eeb7  
STEP 4: RUN touch /var/tmp/test1 && touch /var/tmp/test2  
--> 71900ead142213818c31b3a2c5f8e0b65c3d118cc36a93f5ef73ce977e0baac7  
STEP 5: FROM 71900ead142213818c31b3a2c5f8e0b65c3d118cc36a93f5ef73ce977e0baac7  
STEP 6: RUN touch /var/tmp/foo/bar/test3 || touch /var/tmp/test3  
touch: cannot touch '/var/tmp/foo/bar/test3': No such file or directory  
--> b8d4bad38f669d1031fff9719daede8998fd0404b3ec03d686ea671b2caf3400  
STEP 7: COMMIT myls:1.0  
  
[root@workstation ~]# podman run myls:1.0 /var/tmp  
test1  
test2  
test3
```

Now, try with &&:

```
[root@workstation test2]# vim Dockerfile
FROM rhel7
ENTRYPOINT ["/usr/bin/ls"]
RUN touch /var/tmp/test1 && touch /var/tmp/test2
RUN touch /var/tmp/foo/bar/test3 && touch /var/tmp/test3
```

```
[root@workstation test2]# podman build -t myls:2.0 .
STEP 1: FROM rhel7
STEP 2: ENTRYPOINT ["/usr/bin/ls"]
--> Using cache 9daee6d672b0d61990b3ce73a025ae19437d52bee563bbbeb45be9527bd6eeb7
STEP 3: FROM 9daee6d672b0d61990b3ce73a025ae19437d52bee563bbbeb45be9527bd6eeb7
STEP 4: RUN touch /var/tmp/test1 && touch /var/tmp/test2
--> Using cache 71900ead142213818c31b3a2c5f8e0b65c3d118cc36a93f5ef73ce977e0baac7
STEP 5: FROM 71900ead142213818c31b3a2c5f8e0b65c3d118cc36a93f5ef73ce977e0baac7
STEP 6: RUN touch /var/tmp/foo/bar/test3 && touch /var/tmp/test3
touch: cannot touch '/var/tmp/foo/bar/test3': No such file or directory
error building at step
{Env:[PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin container=oci]
Command:run Args:[touch /var/tmp/foo/bar/test3 && touch /var/tmp/test3] Flags:[]
Attrs:map[] Message:RUN touch /var/tmp/foo/bar/test3 && touch /var/tmp/test3
Original:RUN touch /var/tmp/foo/bar/test3 && touch /var/tmp/test3}: error while running
runtime: exit status 1
```

COPY VS ADD

- They both add files to the container filesystem
- COPY only copies local files, ADD can do that plus decompress .tar or files from a URL
- Use COPY when you don't need ADD
- Both the ADD and COPY instructions copy the files, retaining permissions, with root as the owner, even if the USER instruction is specified. Red Hat recommends using a RUN instruction after the copy to change the owner and avoid "permission denied" errors.

```
[root@workstation ~]# mkdir mycat
[root@workstation ~]# cd mycat
[root@workstation mycat]# echo "helloworld" > important
```

```
[root@workstation mycat]# vim Containerfile
FROM ubi8
ENTRYPOINT ["/bin/cat"]
COPY ./important /tmp/
```

```
[root@workstation mycat]# podman build -t ubi8:mycat1 .
[root@workstation myecho]# podman run ubi8:mycat1 /tmp/important
helloworld
```

Now, lets try and use ADD. First we create a tarball

```
[root@workstation mycat]# tar cvf important.tar important
important
[root@workstation mycat]# vim Containerfile
FROM ubi8
ENTRYPOINT ["/bin/cat"]
ADD ./important.tar /tmp/
```

```
[root@workstation mycat]# podman build -t ubi8:mycat1.1 .
[root@workstation mycat]# podman run ubi8:mycat1.1 /tmp/important
helloworld
[root@workstation mycat]# podman run -it --entrypoint /bin/bash ubi8:mycat1.1
[root@d3f45bdf01ea /]# ls -l /tmp
total 8
-rw-r--r--. 1 root root 11 Aug 15 17:21 important
```

Watch out for permissions on these files using ADD and COPY. ADD will assign the owner to be the first user created (ie uid=1000). This might result in an inability to modify files.

```
[root@workstation mycat]# vim Containerfile
FROM ubi8
ENTRYPOINT ["/bin/cat"]
RUN useradd test1
RUN useradd test2
RUN useradd foo
ADD ./important.tar /tmp/
USER foo
```

```
[student@workstation mycat]$ podman build -t ubi8:mycat1.2 .
[student@workstation mycat]$ podman run -it --entrypoint /bin/bash ubi8:mycat1.2
[foo@93b3b0afa711 /]$ whoami
```

foo

```
[foo@93b3b0afa711 /]$ ls -l /tmp/important
```

```
-rw-rw-r--. 1 test1 test1 10240 Sep 23 10:40 /tmp/important
```

With COPY, the ownership of the copied file(s) will be root:

```
[root@workstation mycat]# vim Containerfile
```

```
FROM ubi8
```

```
ENTRYPOINT ["/bin/cat"]
```

```
RUN useradd test1
```

```
RUN useradd test2
```

```
RUN useradd foo
```

```
COPY important /tmp
```

```
USER foo
```

```
[student@workstation mycat]$ podman build -t ubi8:mycat1.3 .
```

```
[student@workstation mycat]$ podman run -it --entrypoint /bin/bash ubi8:mycat1.3
```

```
[foo@6e6bd10818a0 /]$ ls -l /tmp
```

```
total 20
```

```
-rw-rw-r--. 1 root root 10240 Sep 23 10:40 important
```

```
-rwx-----. 1 root root  701 Sep 14 16:20 ks-script-6x37t6sn
```

```
-rwx-----. 1 root root  291 Sep 14 16:20 ks-script-hlgzt1pz
```

Let's fix this with another RUN instruction to change owner. Only, be careful where the USER instruction is in comparison to the chown. B/c we won't be able to chown a file owned by root as the megatron user:

```
[student@workstation mycat]$ vim Containerfile
```

```
FROM ubi8
```

```
ENTRYPOINT ["/bin/cat"]
```

```
RUN useradd test1
```

```
RUN useradd test2
```

```
RUN useradd foo
```

```
ADD important.tar /tmp
```

```
RUN chown foo:foo /tmp/important*
```

```
USER foo
```

ADDING METADATA: EXPOSE, LABEL, MAINTAINER

```
[student@workstation mycat]$ vim Containerfile
FROM ubi8
LABEL myenv=dev \
      site=dc \
      org=gss
MAINTAINER Andrew Blum <ablum@redhat.com>
EXPOSE 8080
ENTRYPOINT ["/bin/cat"]
RUN useradd test1
RUN useradd test2
RUN useradd foo
ADD important.tar /tmp
RUN chown foo:foo /tmp/important*
USER foo
WORKDIR /var
```

```
[student@workstation mycat]$ podman build -t mycat:latest .
[student@workstation mycat]$ podman inspect mycat:latest | less
```

```
"Config": {
  "User": "foo",
  "ExposedPorts": {
    "8080/tcp": {}
  }
}
```

```
"Labels": {
  "myenv": "dev",
  "name": "ubi8",
  "org": "gss",
```

```
"Author": "Andrew Blum \u003cablum@redhat.com\u003e",
```

EXTRA PRACTICE (if needed)

Let's build a simple Hello World webserver:

```
[student@workstation ~]$ mkdir webhello
[student@workstation ~]$ cd webhello/
[student@workstation webhello]$ vim index.html
<!DOCTYPE html>
<html>
  <title>Welcome to the web Hello World</title>
  <body>
    <h1>Hello World!</h1>
  </body>
</html>
```

```
[student@workstation webhello]$ vi Dockerfile
```

```
FROM registry.access.redhat.com/ubi8:latest
MAINTAINER Andrew Blum <ablum@redhat.com>
LABEL stage=dev \
      version=1.0
ENV myport=8080

RUN yum install httpd -y
RUN sed -i "s/Listen 80/Listen ${myport}/" /etc/httpd/conf/httpd.conf
COPY index.html /var/www/html/index.html

EXPOSE ${myport}
ENTRYPOINT ["httpd", "-D", "FOREGROUND"]
```

```
[student@workstation webhello]$ podman run -d -p 9090:8080 webhello:1.0
3d057c001849eaa83be7497b945f04018f268188ccf95df159aabd6ebb3891fa
[student@workstation webhello]$ curl localhost:9090
<!DOCTYPE html>
<html>
  <title>Welcome to the web Hello World</title>
  <body>
    <h1>Hello World!</h1>
  </body>
```

</html>

[student@workstation webhello]\$ **podman inspect webhello:1.0 | less**

```
"ExposedPorts": {  
  "8080/tcp": {}  
  
  "stage": "dev",  
  "summary": "Provides the latest release of Red Hat Universal Base Image 8.",  
  "url":  
"https://access.redhat.com/containers/#/registry.access.redhat.com/ubi8/images/8.4-206",  
  "vcs-ref": "ed5adf70c28eb951940c72f4173fa32c4bca2165",  
  "vcs-type": "git",  
  "vendor": "Red Hat, Inc.",  
  "version": "1.0"  
  
  "Author": "Andrew Blum \u003cablum@redhat.com\u003e",
```

(USE ONLY IF EXTRA EXAMPLE IS NEEDED)

Let's get a simple webserver created:

```
[root@workstation ~]# mkdir myhttpd  
[root@workstation ~]# cd myhttpd/  
[root@workstation myhttpd]# cp /etc/yum.repos.d/rhel_dvd.repo .  
[root@workstation myhttpd]# vim Dockerfile  
FROM rhel7  
ENTRYPOINT ["httpd","-D","FOREGROUND"]  
COPY ./rhel_dvd.repo /etc/yum.repos.d/  
RUN yum install httpd -y  
  
[root@workstation myhttpd]# podman build -t myhttpd:1.0 .  
[root@workstation myhttpd]# podman run -d myhttpd:1.0  
c5bd745f3bb3b30a9a098d1bfd27fe840c678496c565139baf897e951f2889e3  
[root@workstation myhttpd]# podman inspect c5 | grep -i ipaddress  
  "SecondaryIPAddresses": null,  
  "IPAddress": "10.88.0.37",  
[root@workstation myhttpd]# curl 10.88.0.37 | head
```



```

% Total      % Received % Xferd Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100 3985 100 3985 0      0 961k      0 --:--:-- --:--:-- --:--:-- 1297k
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.1//EN"
"http://www.w3.org/TR/xhtml11/DTD/xhtml11.dtd">

<html xmlns="http://www.w3.org/1999/xhtml" xml:lang="en">
  <head>
    <title>Test Page for the Apache HTTP Server on Red Hat Enterprise Linux</title>

```

Ok, lets modify this container such that a user `megatron` serves up content from his home directory (ie public_html):

```
[root@workstation myhttpd]# echo helloworld > index.html
```

```

[root@workstation myhttpd]# vim Dockerfile
FROM rhel7
ENTRYPOINT ["httpd","-D","FOREGROUND"]
COPY ./rhel_dvd.repo /etc/yum.repos.d/
RUN yum install httpd -y
RUN sed -i 's/UserDir disabled/#UserDir disabled/' /etc/httpd/conf.d/userdir.conf
RUN sed -i 's/#UserDir public_html/Userdir public_html/' /etc/httpd/conf.d/userdir.conf
RUN useradd megatron
RUN mkdir /home/megatron/public_html
COPY ./index.html /home/megatron/public_html
RUN chown -R megatron /home/megatron
RUN chmod 755 /home/megatron/

```

```

[root@workstation myhttpd]# podman build -t myhttpd:2.0 .
[root@workstation myhttpd]# podman run -d myhttpd:2.0
0331755fd6480fc81b6f97b9504f0bde342bb056cdad058773cb582597cd4349
[root@workstation myhttpd]# podman inspect 033 | grep -i ipaddress
    "SecondaryIPAddresses": null,
    "IPAddress": "10.88.0.38",
[root@workstation myhttpd]# curl 10.88.0.38/~megatron/
helloworld

```

Now EXPOSE:

```
[root@workstation myhttpd]# vim Dockerfile
```

```
FROM rhel7
ENTRYPOINT ["httpd","-D","FOREGROUND"]
ENV myport 8080
COPY ./rhel_dvd.repo /etc/yum.repos.d/
RUN yum install httpd -y
RUN sed -i 's/UserDir disabled/^#UserDir disabled/' /etc/httpd/conf.d/userdir.conf
RUN sed -i 's/^#UserDir public_html/Userdir public_html/' /etc/httpd/conf.d/userdir.conf
RUN sed -i "s/Listen 80/Listen ${myport}/" /etc/httpd/conf/httpd.conf
RUN useradd megatron
RUN mkdir /home/megatron/public_html
COPY ./index.html /home/megatron/public_html
RUN chown -R megatron /home/megatron
RUN chmod 755 /home/megatron/
EXPOSE ${myport}
```

```
[root@workstation myhttpd]# podman build -t myhttpd:3.0 .
```

```
[root@workstation myhttpd]# podman inspect myhttpd:3.0 | less
```

(show the exposed port is now in the metadata)

```
[root@workstation myhttpd]# podman run -d myhttpd:3.0
6ea9a0759467f3d7618fb420ebdca8dda24ecf4882822e5f92361c1473aa4a7e
[root@workstation myhttpd]# podman inspect 6ea | grep -i ipaddr
    "SecondaryIPAddresses": null,
    "IPAddress": "10.88.0.39",
[root@workstation myhttpd]# curl 10.88.0.39:8080
```

Now add in author and a label:

```
[root@workstation myhttpd]# vim Dockerfile
```

```
FROM rhel7
MAINTAINER Andrew Blum <ablum@redhat.com>
LABEL myttpd dev
ENTRYPOINT ["httpd","-D","FOREGROUND"]
```

```

ENV myport 8080
COPY ./rhel_dvd.repo /etc/yum.repos.d/
RUN yum install httpd -y
RUN sed -i 's/UserDir disabled/\#UserDir disabled/' /etc/httpd/conf.d/userdir.conf
RUN sed -i 's/\#UserDir public_html/Userdir public_html/' /etc/httpd/conf.d/userdir.conf
RUN sed -i "s/Listen 80/Listen ${myport}/" /etc/httpd/conf/httpd.conf
RUN useradd megatron
RUN mkdir /home/megatron/public_html
COPY ./index.html /home/megatron/public_html
RUN chown -R megatron /home/megatron
RUN chmod 755 /home/megatron/
EXPOSE ${myport}

```

```

[root@workstation myhttpd]# podman build -t myhttpd:4.0 .
[root@workstation myhttpd]# podman inspect myhttpd:4.0 | less

```

Finally, lets clean it up and add in comments:

```

# This file provides a webserver that serves static content from a non-root user's public_html
directory
FROM rhel7
MAINTAINER Andrew Blum <ablum@redhat.com>
LABEL myttpd dev
ENTRYPOINT ["httpd","-D","FOREGROUND"]
ENV myport 8080
EXPOSE ${myport}
# use classroom yum repo and install httpd
COPY ./rhel_dvd.repo /etc/yum.repos.d/
RUN yum install httpd -y
# modify default httpd configuration to allow for local user's public_html and to listen on a
different port
RUN sed -i 's/UserDir disabled/\#UserDir disabled/' /etc/httpd/conf.d/userdir.conf && \
    sed -i 's/\#UserDir public_html/Userdir public_html/' /etc/httpd/conf.d/userdir.conf && \
    sed -i "s/Listen 80/Listen ${myport}/" /etc/httpd/conf/httpd.conf
# add a user and set permissions to allow apache user to serve up public_html
RUN useradd megatron && \
    mkdir /home/megatron/public_html && \
    chown -R megatron /home/megatron && \
    chmod 755 /home/megatron/
# copy webcontent to public_html directory
COPY ./index.html /home/megatron/public_html

```

```
[root@workstation myhttpd]# podman build -t myhttpd:latest .
```

```
[student@workstation ~]$ mkdir myserver
student@workstation ~]$ cd myserver/
[student@workstation myserver]$ echo helloworld > index.html
[student@workstation myserver]$
[student@workstation myserver]$ vim Dockerfile
# This file provides a webserver that serves static content from a non-root user's public_html
directory
FROM ubi8
MAINTAINER Andrew Blum <ablum@redhat.com>
LABEL myhttpd dev
ENTRYPOINT ["httpd","-D","FOREGROUND"]
ENV myport 8080
EXPOSE ${myport}
RUN yum install httpd -y
# modify default httpd configuration to allow for local user's public_html and to listen on a
different port
RUN sed -i 's/UserDir disabled/#UserDir disabled/' /etc/httpd/conf.d/userdir.conf && \
    sed -i 's/#UserDir public_html/Userdir public_html/' /etc/httpd/conf.d/userdir.conf && \
    sed -i "s/Listen 80/Listen ${myport}/" /etc/httpd/conf/httpd.conf
# add a user and set permissions to allow apache user to serve up public_html
RUN useradd megatron && \
    mkdir /home/megatron/public_html && \
    chown -R megatron /home/megatron && \
    chmod 755 /home/megatron/
# copy webcontent to public_html directory
COPY ./index.html /home/megatron/public_html
```

```
[student@workstation myserver]$ podman build -t myhttpd:latest .
[student@workstation myserver]$ podman run -d -p 8083:8080 myhttpd
[student@workstation myserver]$ curl localhost:8083/~megatron/
helloworld
```

```
[student@workstation public_html]$ mkdir ~/public_html
[student@workstation public_html]$ sudo chcon -R -t container_file_t ~/public_html/
[student@workstation public_html]$ podman run -d -p 8084:8080 -v
~/public_html:/home/megatron/public_html myhttpd
```

```
[student@workstation public_html]$ curl localhost:8084/~megatron/
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 3.2 Final//EN">
<html>
<head>
<title>Index of /~megatron</title>
</head>
<body>
<h1>Index of /~megatron</h1>
<table>
<tr><th valign="top"></th><th><a
href="?C=N;O=D">Name</a></th><th><a href="?C=M;O=A">Last modified</a></th><th><a
href="?C=S;O=A">Size</a></th><th><a href="?C=D;O=A">Description</a></th></tr>
<tr><th colspan="5"><hr></th></tr>
<tr><td valign="top"></td><td><a
href="/">Parent Directory</a>      </td><td>&nbsp;</td><td align="right"> -
</td><td>&nbsp;</td></tr>
<tr><th colspan="5"><hr></th></tr>
</table>
</body></html>
```

```
[student@workstation public_html]$ echo hello_from_host_volume >
~/public_html/index.html
[student@workstation public_html]$ curl localhost:8083/~megatron/
hello_from_host_volume
```

How to install software from Red Hat repositories?

<https://access.redhat.com/solutions/1443553>

enable repos in container: <https://access.redhat.com/solutions/1443553>

""

Entitlement information from the host is injected into the container when the first yum command in the container is run.

Thus, containers are not entitled, but they can access any repository the host can access based on those entitlements, even if the repositories are disabled on the host

NOTE: Until the first yum command is run, /etc/yum.repos.d/redhat.repo contains no repositories, so yum-config-manager will not enable/disable anything.

""

```
RUN yum repolist --disablerepo=* && \
    yum-config-manager --disable '*' > /dev/null && \
    yum-config-manager --enable rhel-7-server-rpms > /dev/null
```

Otherwise, ubi.repo are available if you are using ubi container images:

```
[root@8213cb08ef88 /]# yum repolist
```

Updating Subscription Management repositories.

Unable to read consumer identity

This system is not registered to Red Hat Subscription Management. You can use subscription-manager to register.

Red Hat Universal Base Image 8 (RPMs) - BaseOS

120 kB/s | 759 kB 00:06

Red Hat Universal Base Image 8 (RPMs) - AppStream

1.5 MB/s | 3.1 MB 00:02

Red Hat Universal Base Image 8 (RPMs) - CodeReady Builder

5.9 kB/s | 9.1 kB 00:01

repo id	status	repo name
ubi-8-appstream		Red Hat Universal Base Image 8 (RPMs) -
AppStream	797	
ubi-8-baseos		Red Hat Universal Base Image 8 (RPMs) -
BaseOS	663	
ubi-8-codeready-builder		Red Hat Universal Base Image 8 (RPMs) -
CodeReady Builder	12	

```
[root@8213cb08ef88 /]# yum search openjdk
```

Updating Subscription Management repositories.

Unable to read consumer identity

This system is not registered to Red Hat Subscription Management. You can use subscription-manager to register.

Last metadata expiration check: 0:00:12 ago on Thu Dec 12 21:55:15 2019.

```
=====
Name & Summary Matched: openjdk
=====
```

```
java-11-openjdk.x86_64 : OpenJDK Runtime Environment 11
java-1.8.0-openjdk.x86_64 : OpenJDK Runtime Environment 8
java-11-openjdk-devel.x86_64 : OpenJDK Development Environment 11
java-1.8.0-openjdk-devel.x86_64 : OpenJDK Development Environment 8
java-11-openjdk-headless.x86_64 : OpenJDK Headless Runtime Environment 11
java-1.8.0-openjdk-headless.x86_64 : OpenJDK Headless Runtime Environment 8
```

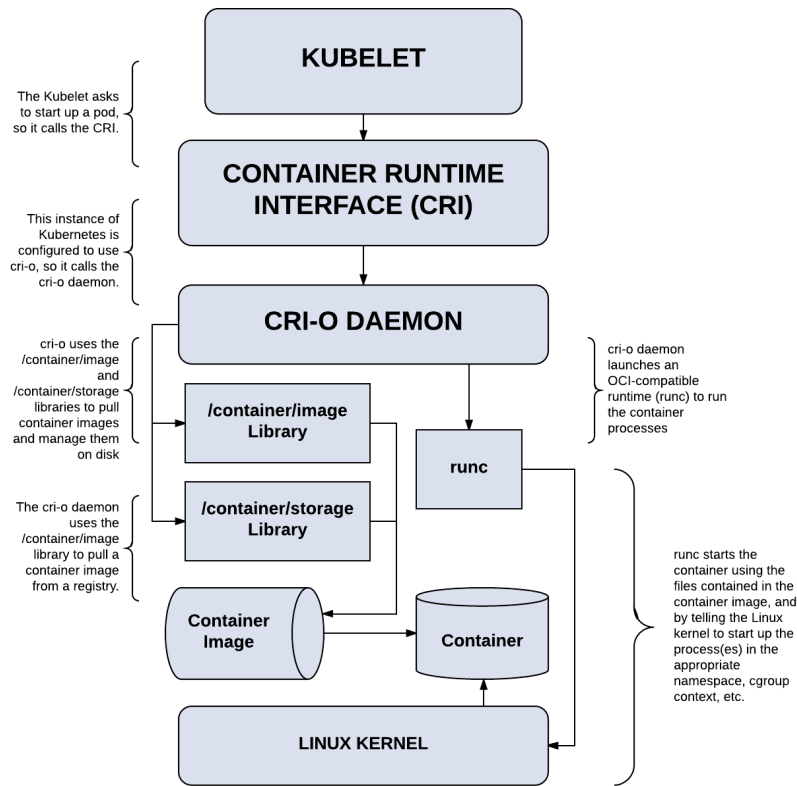
```
[root@8213cb08ef88 /]# ls /etc/yum.repos.d/
redhat.repo  ubi.repo
```

CHAPTER 6: DEPLOYING CONTAINERIZED APPLICATIONS ON OPENSIFT

DESCRIBING KUBERNETES AND OPENSIFT ARCHITECTURE

What is crio ?

Cri-o: container runtime used in k8s/ocp. It's what k8s depends on to run containers. see <https://www.redhat.com/en/blog/introducing-cri-o-10>



(Using DO280 environment)

Classroom Identifier vwormreqvvomqez200819
 Cluster ID bb8bca34-2206-4483-8e6f-2b14a70b1547
 Cluster Username kubeadmin
 Cluster Password IWx4R-Q8qMj-yTImo-A4Zms
 Cluster API URL api.ocp-vwormreqvvomqez200819.do280.rht-na.nextcle.com:6443
 State installing
 Classroom Status OpenShift Installation Completed

```
[ablum@badger ocp4]$ oc login -u kubeadmin -p IWx4R-Q8qMj-yTImo-A4Zms
api.ocp-vwormreqvvomqez200819.do280.rht-na.nextcle.com:6443
[ablum@badger ocp4]$ oc debug node/ip-10-0-136-0.ec2.internal
Starting pod/ip-10-0-136-0ec2internal-debug ...
To use host binaries, run `chroot /host`
Pod IP: 10.0.136.0
If you don't see a command prompt, try pressing enter.
sh-4.2#
sh-4.2# chroot /host /bin/bash
```



```
[root@master0 /]# systemctl status kubelet
```

- kubelet.service - Kubernetes Kubelet
Loaded: loaded (/etc/systemd/system/kubelet.service; enabled; vendor preset: enabled)

```
[root@master0 /]# systemctl status crio
```

- crio.service - Open Container Initiative Daemon
Loaded: loaded (/usr/lib/systemd/system/crio.service; disabled; vendor preset: disabled)

```
[root@master0 /]# podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

```
[root@master0 /]#
```

```
[root@master0 /]#
```

```
[root@master0 /]# podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

```
[root@master0 /]# crictl -h
```

NAME:

crictl - client for CRI

```
[root@master0 /]# crictl ps
```

<https://blog.openshift.com/crictl-vs-podman/>

NOTE1: crictl is not as fully featured as podman. It can't build nor restart containers. They both use runc, but differently so they won't "see" the same running containers. Crictl sees what k8s sees.

NOTE2: podman and cri-o do share the same image storage libraries so a podman rmi will remove the image from local storage on that node from cri-o (will have to re-pull the image from registry)

What makes a control node ?

- Kube-apiserver
- Kube-controller-manager - The Controller Manager Server watches etcd for changes to objects such as replication, namespace, and serviceaccount controller objects, and then uses the API to enforce the specified state. Several such processes create a cluster with one active leader at a time.
- Kube-scheduler

- Etcd - etcd stores the persistent master state while other components watch etcd for changes to bring themselves into the specified state.

```
[root@cluster-master-0 ~]# crictl ps | grep etcd
a9e431dc52745
643c21638c1c966fe18ca1cc8547dd401df70e85d83ca6de76b9a7957703b993
    About an hour ago    Running          etcd-member
```

```
[root@cluster-master-0 ~]# crictl ps | grep kube
```

<https://kubernetes.io/docs/concepts/overview/components/>

Podman and skopeo are still available, point out how to use the authfile:

```
[root@master0 /]# skopeo inspect docker://registry.redhat.io/rhel7
FATA[0000] unable to retrieve auth token: invalid username/password
[root@master0 /]# skopeo inspect --authfile=/var/lib/kubelet/config.json
docker://registry.redhat.io/rhel7
```

(useful for recovering from control plane certificate expiry:

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.2/html/backup_and_restore/disaster-recovery#dr-scenario-3-recovering-expired-certs_dr-recovering-expired-certs)

CREATING KUBERNETES RESOURCES

GUIDED PRACTICE: Deploying a Database Server on OpenShift

```
[student@workstation ~]$ lab openshift-resources start
```

Setting up workstation for the Guided Exercise: Deploying a Database Server on OpenShift

Verifying the OpenShift cluster is running:

- Log in on OpenShift..... SUCCESS
- Check the internal registry is up and running..... SUCCESS
- Ensuring the 'rhn-support-ablum-mysql-openshift' project is absent SUCCESS

What is oc ?

First, Let's understand the "oc" or openshift client cli better. <https://github.com/openshift/oc>

Navigate to <https://mirror.openshift.com/pub/openshift-v4/> then clients > oc. Click on the appropriate version (should try to match major versions to cluster version).

```
[student@workstation ~]$ mkdir work
```

```
[student@workstation ~]$ cd work/
```

```
[student@workstation work]$ wget
```

```
https://mirror.openshift.com/pub/openshift-v4/clients/oc/4.6/linux/oc.tar.gz
```

```
[student@workstation work]$ tar xvfz oc.tar.gz
```

```
oc
```

```
[student@workstation work]$ file oc
```

```
oc: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked (uses shared libs),  
for GNU/Linux 2.6.32, BuildID[sha1]=b175d93b9990cdeb07fc6a7c603dd3d8ded14c6f, stripped
```

```
[student@workstation work]$ ./oc version
```

```
Client Version: 4.6.44
```

```
Kubernetes Version: v1.19.0+8d12420
```

```
[student@workstation work]$ oc completion bash > oc_completion
```

```
[student@workstation work]$ head oc_completion
```

```
# bash completion for oc
```

```
-*- shell-script -*-
```

```
__oc_debug()
```

```
{
```

```
    if [[ -n ${BASH_COMP_DEBUG_FILE} ]]; then
```

```
        echo "$*" >> "${BASH_COMP_DEBUG_FILE}"
```

If you wanted to use this version you could replace the one that comes with our training environment. Let's not do that. Here is what we get with our environment:

```
[student@workstation work]$ cd ..
[student@workstation ~]$ rm -rf work/
[student@workstation ~]$
```

```
[student@workstation work]$ which oc
/usr/local/bin/oc
[student@workstation work]$ oc version
Client Version: 4.5.4
Kubernetes Version: v1.18.3+002a51f
```

```
[student@workstation work]$ head /etc/bash_completion.d/oc
```

```
# bash completion for oc                                -*- shell-script -*-
```

```
__oc_debug()
{
    if [[ -n ${BASH_COMP_DEBUG_FILE} ]]; then
        echo "$*" >> "${BASH_COMP_DEBUG_FILE}"
    fi
}
```

There are a lot of subcommands used with oc:

```
[student@workstation ~]$ oc [TAB][TAB]
```

adm	cancel-build	delete	extract	logs	policy	
rollback	set					
annotate	cluster-info	describe	get	new-app	port-forward	rollout
	start-build					
api-resources	completion	diff	idle	new-build	process	rsh
	status					
api-versions	config	edit	image	new-project	project	rsync
	tag					
apply	convert	ex	import-image	observe	projects	run
	version					
attach	cp	exec	label	options	proxy	scale
	wait					
auth	create	explain	login	patch	registry	secrets
	whoami					
autoscale	debug	expose	logout	plugin	replace	
serviceaccounts						

Getting help on any of them:

```
[student@workstation ~]$ oc logs --help
```

Print the logs for a resource

Why not man pages ? We didn't install oc via rpm - many customers don't use RHEL as their base laptop or workstation. The oc client is available for Windows and MacOS as well.

Let's login into the shared OCP cluster provisioned by our learning environment:

```
[student@workstation ~]$ cat /usr/local/etc/ocp4.config
```

```
[student@workstation ~]$ source /usr/local/etc/ocp4.config
```

```
[student@workstation ~]$ oc login --help
```

Log in to the given server with the given credentials (will not prompt interactively)

```
oc login localhost:8443 --username=myuser --password=mypass
```

```
[student@workstation ~]$ oc login -u ${RHT_OCP4_DEV_USER} -p
```

```
${RHT_OCP4_DEV_PASSWORD} ${RHT_OCP4_MASTER_API}
```

Login successful.

For the future, I'll just the alias we setup in bashrc:

```
[student@workstation ~]$ alias dl
```

```
alias dl='oc login -u ${RHT_OCP4_DEV_USER} -p ${RHT_OCP4_DEV_PASSWORD}  
${RHT_OCP4_MASTER_API}'
```

```
[student@workstation ~]$ dl
```

Login successful.

You don't have any projects. You can try to create a new project, by running

```
oc new-project <projectname>
```

How to create my first openshift project ?

Well, we don't have any projects...ok, let's change that:

```
[student@workstation ~]$ oc new-project ${RHT_OCP4_DEV_USER}-mysql-openshift  
[student@workstation ~]$ oc new-app -h
```

Point out::

Use a MySQL image in a private registry to create an app and override application artifacts' names

oc new-app --docker-image=myregistry.com/mycompany/mysql --name=private

-e, --env=[]: Specify a key-value pair for an environment variable to set into each container.

--as-deployment-config

Starting in OCP4.5, oc new-app will create k8s deployments (by default) instead of deploymentconfigs.

In our new project:

```
[student@workstation ~]$ oc project
```

Using project "gwohys-mysql-openshift" on server "https://api.na46.prod.nextcle.com:6443".

How to create my first application in openshift?

Let's start with a simple example: <https://hub.docker.com/r/openshift/hello-openshift/>

```
[student@workstation ~]$ skopeo inspect docker://docker.io/openshift/hello-openshift
```

The code is available here:

https://github.com/openshift/origin/blob/master/examples/hello-openshift/hello_openshift.go

Deploying using the one from docker.io might result in the following errors:

```
ERRO[0008] error searching registry "docker.io": couldn't search registry "docker.io": error  
pinging docker registry index.docker.io: Get https://index.docker.io/v2/: dial tcp: lookup  
index.docker.io on 172.25.250.254:53: server misbehaving
```

Or

W0119 08:12:47.671222 13832 dockerimagelookup.go:237] container image registry lookup failed: docker.io/openshift/hello-openshift:latest: toomanyrequests: You have reached your pull rate limit. You may increase the limit by authenticating and upgrading:
<https://www.docker.com/increase-rate-limit>
error: unable to locate any local docker images with name
"docker.io/openshift/hello-openshift:latest"

Then, use personal one cloned from docker.io: to quay.io:

```
[student@workstation ~]$ oc new-app --image=quay.io/ajblum/hello-openshift:latest
```

```
[student@workstation ~]$ oc status
```

In project ablum-mytest on server <https://api.na46.prod.nextcle.com:6443>

svc/hello-openshift - 172.30.182.91 ports 8080, 8888
deployment/hello-openshift deploys istag/hello-openshift:latest
deployment #2 running for 27 seconds - 1 pod
deployment #1 deployed 28 seconds ago

1 info identified, use 'oc status --suggest' to see details.

```
[student@workstation ~]$ oc get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod /hello-openshift-7b74767ff6-28hf8	1/1	Running	0	39s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service /hello-openshift	ClusterIP	172.30.182.91	<none>	8080/TCP,8888/TCP	40s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment .apps/hello-openshift	1/1	1	1	40s

NAME	DESIRED	CURRENT	READY	AGE
replicaset .apps/hello-openshift-7b74767ff6	1	1	1	39s
replicaset.apps/hello-openshift-9c8f5d9f	0	0	0	40s

NAME	IMAGE REPOSITORY
imagestream .image.openshift.io/hello-openshift	
default-route-openshift-image-registry.apps.na46.prod.nextcle.com/ablum-mytest/hello-openshift	
latest	39 seconds ago

Looking at any one of these resources:

```
[student@workstation ~]$ oc describe pod/hello-openshift-57749cf8d4-6cx4c
```

Labels: deployment=hello-openshift
pod-template-hash=57749cf8d4

containers:

hello-openshift:

Container ID:

cri-o://6acc97c88bb6062d85f55296da532b44b4ab881fb1a2a1359bd636d56bed3680

Image:

docker.io/openshift/hello-openshift@sha256:aaea76ff622d2f8bcb32e538e7b3cd0ef6d291953f3e7c9f556c1ba5baf47e2e

Image ID:

docker.io/openshift/hello-openshift@sha256:aaea76ff622d2f8bcb32e538e7b3cd0ef6d291953f3e7c9f556c1ba5baf47e2e

Ports: 8080/TCP, 8888/TCP

Host Ports: 0/TCP, 0/TCP

State: Running

Started: Thu, 16 Sep 2021 15:24:26 -0400

Ready: True

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	Scheduled	29s	default-scheduler	Successfully assigned zrsqwh-mysql-openshift/hello-openshift-57749cf8d4-6cx4c to na46-5m9nf-worker-0-fbm4
Normal	AddedInterface	27s	multus	Add eth0 [10.128.4.218/23]
Normal	Pulled	27s	kubelet	Container image "docker.io/openshift/hello-openshift@sha256:aaea76ff622d2f8bcb32e538e7b3cd0ef6d291953f3e7c9f556c1ba5baf47e2e" already present on machine
Normal	Created	26s	kubelet	Created container hello-openshift
Normal	Started	26s	kubelet	Started container hello-openshift

To see a yaml representation (instead of this “pretty” format):

```
[student@workstation ~]$ oc get pods hello-openshift-57749cf8d4-6cx4c -o yaml
```


WHOA!!! That's a lot of stuff....I don't understand any of that !

What are PODS?

```
[student@workstation ~]$ oc explain pods
```

KIND: Pod


VERSION: v1

DESCRIPTION:

Pod is a collection of containers that can run on a host. This resource is created by clients and scheduled onto hosts.

```
[student@workstation ~]$ oc explain pods.spec
```

All the containers in a pod share the same networking namespace. Different pods are connected to each other by the k8s pod sdn, a networking overlay implemented by opensvswitch (ovs) running on each node:



What are those managedFields ?

<https://access.redhat.com/solutions/5332041>

```
managedFields:
- apiVersion: v1
  fieldsType: FieldsV1
  fieldsV1:
    f:metadata:
    f:annotations:
      .: {}
    f:openshift.io/generated-by: {}
    f:generateName: {}
    f:labels:
      .: {}
    f:deployment: {}
    f:pod-template-hash: {}
    f:ownerReferences:
      .: {}
    k:{"uid":"b06524c9-cae7-41c4-88bc-9113c9ff21ed"}:
      .: {}
    f:apiVersion: {}
    f:blockOwnerDeletion: {}
    f:controller: {}
    f:kind: {}
    f:name: {}
    f:uid: {}
```

Interesting discussion on this issue upstream:

<https://github.com/kubernetes/kubernetes/issues/90066>

Where do we specify the pod that we want to see deployed ? There are several ways to do that, but one that's used here by the `oc new-app` command is a deployment:

```
[student@workstation ~]$ oc get deployment
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
hello-openshift 1/1    1           1           7m52s
```

```
[student@workstation ~]$ oc explain deployment
```

```
[student@workstation ~]$ oc explain deployment.spec.template
```

KIND: Deployment

VERSION: apps/v1

RESOURCE: template <Object>

DESCRIPTION:

Template describes the pods that will be created.

PodTemplateSpec describes the data a pod should have when created from a template

```
[student@workstation ~]$ oc get deployment hello-openshift -o yaml
```

template:

metadata:

annotations:

openshift.io/generated-by: OpenShiftNewApp

creationTimestamp: null

labels:

deployment: hello-openshift

spec:

containers:

- image:

docker.io/openshift/hello-openshift@sha256:aaea76ff622d2f8bcb32e538e7b3cd0ef6d291953f3e7c9f556c1ba5baf47e2e

imagePullPolicy: IfNotPresent

name: hello-openshift

ports:

- containerPort: 8080

protocol: TCP

- containerPort: 8888

protocol: TCP

resources: {}

terminationMessagePath: /dev/termination-log

terminationMessagePolicy: File

Notice the pod running here inherited the label defined in the template spec of the deployment:

```
[student@workstation ~]$ oc get pods --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
------	-------	--------	----------	-----	--------

```
hello-openshift-57749cf8d4-6cx4c 1/1      Running 0          11m
deployment=hello-openshift,pod-template-hash=57749cf8d4
```

You can use this label to select matching resources:

```
[student@workstation ~]$ oc get pods -l deployment=hello-openshift
NAME                                READY STATUS   RESTARTS AGE
hello-openshift-57749cf8d4-6cx4c 1/1      Running 0          13m
```

What are DEPLOYMENT CONFIGURATION vs DEPLOYMENT?

<https://docs.openshift.com/container-platform/4.5/applications/deployments/what-deployments-are.html#what-deployments-are>

OpenShift Container Platform deployments from DeploymentConfigs also provide the ability to transition from an existing deployment of an image to a new one and also define hooks to be run before or after creating the ReplicationController.

One important difference between Deployments and DeploymentConfigs is the properties of the [CAP theorem](#) that each design has chosen for the rollout process. DeploymentConfigs prefer consistency, whereas Deployments take availability over consistency.

For DeploymentConfigs, if a node running a deployer Pod goes down, it will not get replaced. The process waits until the node comes back online or is manually deleted. Manually deleting the node also deletes the corresponding Pod. This means that you can not delete the Pod to unstick the rollout, as the kubelet is responsible for deleting the associated Pod.

However, Deployments rollouts are driven from a controller manager. The controller manager runs in high availability mode on masters and uses leader election algorithms to value availability over consistency. During a failure it is possible for other masters to act on the same Deployment at the same time, but this issue will be reconciled shortly after the failure occurs

How do you delete this application ? Let's try:

```
[student@workstation ~]$ oc delete pod/hello-openshift-7b74767ff6-28hf8
```

(short pause)

```
[student@workstation ~]$ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-openshift-7b74767ff6-529h6	1/1	Running	0	11s

It's back ! This is due to the replicaset (controller) insuring that our desired state is met:

```
[student@workstation ~]$ oc get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
hello-openshift-57749cf8d4	1	1	1	14m
hello-openshift-5df9dfbb6c	0	0	0	14m

```
[student@workstation ~]$ oc describe replicaset hello-openshift-57749cf8d4
```

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	SuccessfulCreate	14m	replicaset-controller	Created pod: hello-openshift-57749cf8d4-6cx4c
Normal	SuccessfulCreate	75s	replicaset-controller	Created pod: hello-openshift-57749cf8d4-x48z

What are REPLICATION CONTROLLERS vs REPLICASETS?

DeploymentConfigs involve one or more *ReplicationControllers*, which contain a point-in-time record of the state of a DeploymentConfig as a Pod template.

Similarly, Deployments involve one or more *ReplicaSets*.

The difference between a ReplicaSet and a ReplicationController is that a ReplicaSet supports set-based selector requirements whereas a replication controller only supports equality-based selector requirements. *Set-based* label requirements allow filtering keys according to a set of values.

Consider this set-based selector,

environment in (production, qa)

Selects all resources with key equal to environment and value equal to production or qa.

Consider this equality-based selector:

Environment = production

selects all resources with key equal to environment and value equal to production

So, how do we delete this application ?

It was the deployment that defined what we wanted, this is what we must delete:

```
[student@workstation ~]$ oc delete deployment.apps/hello-openshift
deployment.apps "hello-openshift" deleted
```

There are some other resources left as well (route, service, imagestream):

```
[student@workstation ~]$ oc get all --show-labels
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
LABELS					
service/hello-openshift	ClusterIP	172.30.182.91	<none>	8080/TCP,8888/TCP	21m
app.kubernetes.io/component=hello-openshift,app.kubernetes.io/instance=hello-openshift,app=hello-openshift					

```
[student@workstation ~]$ oc delete all -l app=hello-openshift
service "hello-openshift" deleted
imagestream.image.openshift.io "hello-openshift" deleted
route.route.openshift.io "hello-openshift" deleted
```

Let's try to deploy again using --as-deployment-config

```
[student@workstation ~]$ oc new-app --image=quay.io/ajblum/hello-openshift:latest
--as-deployment-config
```

```
[student@workstation ~]$ oc get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/hello-openshift-1-deploy	0/1	Completed	0	27s
pod/hello-openshift-1-gn6kd	1/1	Running	0	24s

This time, the deployment is rolled out using a "deploy" pod running our namespace. This represents a fundamental difference between a deployment and a deploymentconfig.

See

<https://docs.openshift.com/container-platform/4.7/applications/deployments/what-deployments-are.html>


What are SERVICES ?

Every pod is assigned a unique IP addressed by the pod SDN. When a new pod is created, it will use a different IP (depending on which node it is assigned to):

```
[student@workstation ~]$ oc get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
hello-openshift-1-7sln2	1/1	Running	0	34s	10.128.5.87	
na46-5m9nf-worker-0-fbmq4	<none>	<none>				
hello-openshift-1-deploy	0/1	Completed	0	37s	10.128.5.86	
na46-5m9nf-worker-0-fbmq4	<none>	<none>				

What is the best way for applications in different pods to communicate ? They could use the pod IP, but that would be tough for the applications to keep track of. A service makes this easier, but will require a whole different SDN:



```
[student@workstation ~]$ oc get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-openshift	ClusterIP	172.30.106.151	<none>	8080/TCP,8888/TCP	37m

```
[student@workstation ~]$
[student@workstation ~]$
[student@workstation ~]$ oc describe service hello-openshift
Name:          hello-openshift
Namespace:      zrsqwh-mysql-openshift
Labels:         app=hello-openshift
                 app.kubernetes.io/component=hello-openshift
                 app.kubernetes.io/instance=hello-openshift
Annotations:    openshift.io/generated-by: OpenShiftNewApp
Selector:       deploymentconfig=hello-openshift
```

```
[student@workstation ~]$ oc get endpoints
NAME          ENDPOINTS          AGE
hello-openshift 10.128.5.87:8080,10.128.5.87:8888 40m
```

```
[student@workstation ~]$ oc get pods -l deploymentconfig=hello-openshift
NAME          READY STATUS  RESTARTS AGE
hello-openshift-1-7sln2 1/1 Running 0         44m
[student@workstation ~]$ oc get pods -l deploymentconfig=hello-openshift -o wide
NAME          READY STATUS  RESTARTS AGE IP      NODE
NOMINATED NODE READINESS GATES
hello-openshift-1-7sln2 1/1 Running 0         44m 10.128.5.87
na46-5m9nf-worker-0-fbmq4 <none> <none>
```

If we delete the pod, watch what happens...The endpoint will change to use the new pod's IP, but the service IP won't change:

```
[student@workstation ~]$ oc get endpoints
NAME          ENDPOINTS          AGE
hello-openshift 10.128.5.115:8080,10.128.5.115:8888 46m
[student@workstation ~]$ oc get pods -l deploymentconfig=hello-openshift -o wide
NAME          READY STATUS  RESTARTS AGE IP      NODE
NOMINATED NODE READINESS GATES
hello-openshift-1-dhsvj 1/1 Running 0         56s 10.128.5.115
na46-5m9nf-worker-0-fbmq4 <none> <none>
[student@workstation ~]$ oc get service
NAME          TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
hello-openshift ClusterIP 172.30.106.151 <none>      8080/TCP,8888/TCP 46m
```


What is an IMAGESTREAM ?

[student@workstation ~]\$ **oc explain is**

KIND: ImageStream

VERSION: image.openshift.io/v1

DESCRIPTION:

An ImageStream stores a mapping of tags to images, metadata overrides that are applied when images are tagged in a stream, and an optional reference to a container image repository on a registry. Users typically update the spec.tags field to point to external images which are imported from container registries using credentials in your namespace with the pull secret type, or to existing image stream tags and images which are immediately accessible for tagging or pulling. The history of images applied to a tag is visible in the status.tags field and any user who can view an image stream is allowed to tag that image into their own image streams. Access to pull images from the integrated registry is granted by having the "get imagestreams/layers" permission on a given image stream. Users may remove a tag by deleting the imagestreamtag resource, which causes both spec and status for that tag to be removed. Image stream history is retained until an administrator runs the prune operation, which removes references that are no longer in use. To preserve a historical image, ensure there is a tag in spec pointing to that image by its digest.

[student@workstation ~]\$ **oc explain istag**

KIND: ImageStreamTag

VERSION: image.openshift.io/v1

DESCRIPTION:

ImageStreamTag represents an Image that is retrieved by tag name from an ImageStream. Use this resource to interact with the tags and images in an image stream by tag, or to see the image details for a particular tag. The image associated with this resource is the most recently successfully tagged, imported, or pushed image (as described in the image stream status.tags.items list for this tag). If an import is in progress or has failed the previous image will be shown. Deleting an image stream tag clears both the status and spec fields of an image stream. If no image can be retrieved for a given tag, a not found error will be returned.

```

[student@workstation ~]$ oc get is
NAME          IMAGE REPOSITORY
TAGS  UPDATED
hello-openshift
default-route-openshift-image-registry.apps.na46.prod.nextcle.com/zrsqwh-mysql-openshift/hello-openshift latest 47 minutes ago
[student@workstation ~]$ oc describe is hello-openshift
Name:          hello-openshift
Namespace:     zrsqwh-mysql-openshift
Created:       47 minutes ago
Labels:        app=hello-openshift
               app.kubernetes.io/component=hello-openshift
               app.kubernetes.io/instance=hello-openshift
Annotations:   openshift.io/generated-by=OpenShiftNewApp
               openshift.io/image.dockerRepositoryCheck=2021-09-16T19:40:15Z
Image Repository:
default-route-openshift-image-registry.apps.na46.prod.nextcle.com/zrsqwh-mysql-openshift/hello-openshift
Image Lookup:   local=false
Unique Images:  1
Tags:           1

latest
tagged from docker.io/openshift/hello-openshift

*
docker.io/openshift/hello-openshift@sha256:aaea76ff622d2f8bcb32e538e7b3cd0ef6d291953f3e7c9f556c1ba5baf47e2e
47 minutes ago

```

These abstractions are useful for handling changes in image registries. They can be used to trigger builds and deployments when images are pushed to registries.

NOTE: These do not contain the actual image data (ie the lower dirs). They are pointers.

https://docs.openshift.com/container-platform/4.8/openshift_images/images-understand.html

What are PERSISTENT VOLUMES and PERSISTENT VOLUME CLAIMS ?

```
[student@workstation ~]$ oc explain pv
```

```
KIND: PersistentVolume
```

```
VERSION: v1
```

DESCRIPTION:

PersistentVolume (PV) is a storage resource provisioned by an administrator. It is analogous to a node. More info:

<https://kubernetes.io/docs/concepts/storage/persistent-volumes>

```
[student@workstation ~]$ oc explain pvc
```

```
KIND: PersistentVolumeClaim
```

```
VERSION: v1
```

DESCRIPTION:

PersistentVolumeClaim is a user's request for and claim to a persistent volume

When making the request for storage for your application (ie the **pvc**), you will need to define characteristics that the persistent volume controller (a control loop that runs within the kube-controller-manager) uses to match the available persistent volumes (the **pv**).

Once a match is made it is said to be “bound” and can be used (ie mounted) within a container in a pod.

NOTE: Developers do not have permissions to manipulate or view persistent volumes (pv) only the claims (pvc)

So, how do we test our hello-world application ?

One way to test this application is with `oc port-forward`

```
[student@workstation ~]$ nohup oc port-forward hello-openshift-5fffbb958-nxwgn 8080:8080 &
```

```
[1] 70971
```

```
[student@workstation ~]$ nohup: ignoring input and appending output to 'nohup.out'
```

```
[student@workstation ~]$ curl localhost:8080
Hello OpenShift!
[student@workstation ~]$ fg
nohup oc port-forward hello-openshift-5fffbfb958-nxwgn 8080:8080
^C[student@workstation ~]$
```

For this application, it's not possible to create an interactive shell. Why ?

```
[student@workstation ~]$ oc rsh hello-openshift-1-l68pp
ERRO[0000] exec failed: unable to start container process: exec: "/bin/sh": stat /bin/sh: no such
file or directory
command terminated with exit code 255
```

In this case, consider the Containerfile (Dockerfile) used to create this image:

<https://github.com/openshift/origin/blob/master/examples/hello-openshift/Dockerfile>

FROM scratch

There is no /bin/sh so no way to run a command that doesn't exist. You would need to either rebuild this image or use a different image for debugging like:

```
[student@workstation ~]$ oc debug pod/hello-openshift-1-l68pp --image
registry.access.redhat.com/ubi8:latest
Starting pod/hello-openshift-1-l68pp-debug, command was: /hello-openshift
Pod IP: 10.129.8.120
If you don't see a command prompt, try pressing enter.
sh-4.4$ ls
sh-4.4$ env | grep SERVICE
HELLO_OPENSHIFT_SERVICE_PORT_8080_TCP=8080
HELLO_OPENSHIFT_SERVICE_HOST=172.30.91.2
KUBERNETES_SERVICE_PORT_HTTPS=443
HELLO_OPENSHIFT_SERVICE_PORT=8080
KUBERNETES_SERVICE_PORT=443
HELLO_OPENSHIFT_SERVICE_PORT_8888_TCP=8888
KUBERNETES_SERVICE_HOST=172.30.0.1
sh-4.4$ whoami
1001180000
```

Let's clean up and complete this exercise:

```
[student@workstation ~]$ oc delete deploymentconfig.apps.openshift.io/hello-openshift
deploymentconfig.apps.openshift.io "hello-openshift" deleted
[student@workstation ~]$ oc delete all -l app=hello-openshift
service "hello-openshift" deleted
imagestream.image.openshift.io "hello-openshift" deleted
```

How to deploy mysql database in openshift using a template?

Now, let's deploy a mysql database using a prebuilt template:

```
[student@workstation ~]$ oc new-app --template=mysql-persistent -p MYSQL_USER=user1
-p MYSQL_PASSWORD=mypa55 -p MYSQL_DATABASE=testdb -p
MYSQL_ROOT_PASSWORD=r00tpa55 -p VOLUME_CAPACITY=10Gi
[student@workstation ~]$ oc status -h
[student@workstation ~]$ oc status
In project rhn-support-ablum-mysql-openshift on server
https://api.ocp-na2.prod.nextcle.com:6443
```

```
svc/mysql-openshift - 172.30.192.167:3306
  dc/mysql-openshift deploys istag/mysql-openshift:latest
    deployment #1 deployed 45 seconds ago - 1 pod
```

3 infos identified, use 'oc status --suggest' to see details.

```
[student@workstation ~]$ oc get all
NAME                                READY STATUS  RESTARTS  AGE
pod/mysql-openshift-1-deploy        0/1   Completed  0         3m40s
pod/mysql-openshift-1-zlzbr         1/1   Running    0         3m37s
```

```
NAME                                DESIRED CURRENT  READY AGE
replicationcontroller/mysql-openshift-1 1      1        1     3m41s
```

```
NAME                                TYPE           CLUSTER-IP      EXTERNAL-IP  PORT(S)    AGE
service/mysql-openshift             ClusterIP       172.30.139.239  <none>       3306/TCP   3m42s
```

```
NAME                                REVISION  DESIRED  CURRENT  TRIGGERED BY
deploymentconfig.apps.openshift.io/mysql-openshift 1         1        1
config,image(mysql-openshift:latest)
```

NAME	IMAGE REPOSITORY	TAGS	UPDATED
imagestream.image.openshift.io/mysql-openshift	default-route-openshift-image-registry.apps.na45.prod.nextcle.com/rhn-support-ablum-mysql-op	enshfit/mysql-openshift	latest About an hour ago

Maybe, you see this failure:

```
Warning FailedAttachVolume 37s (x4 over 3m44s) attachdetach-controller
AttachVolume.Attach failed for volume "pvc-b253a9e9-7eee-4a3b-bc74-93c99380f070" :
Volume "8e5a6442-df79-4127-87b9-8b824d0c3662" failed to be attached within the allotted time
Warning FailedMount 25s (x2 over 2m42s) kubelet Unable to attach or
mount volumes: unmounted volumes=[mysql-data], unattached volumes=[mysql-data
default-token-pdcdd]: timed out waiting for the condition
```

Retry with:

```
[student@workstation ~]$ oc rollout latest dc/mysql
deploymentconfig.apps.openshift.io/mysql rolled out
```

If this still fails with storage try a different deployment using:
<https://catalog.redhat.com/software/containers/search?q=mysql>

```
[student@workstation ~]$ oc new-app --as-deployment-config
--docker-image=registry.redhat.io/rhel8/mysql-80 --name=mysql-openshift -e
MYSQL_USER=user1 -e MYSQL_PASSWORD=mypa55 -e MYSQL_DATABASE=testdb -e
MYSQL_ROOT_PASSWORD=r00tpa5
```

OR

```
[student@workstation ~]$ oc get templates -n openshift | grep mysql
mysql-ephemeral MySQL database service, without persistent
storage. For more information abou... 8 (3 generated) 3
```

```
[student@workstation ~]$ oc delete all --all
pod "mysql-2-bzjgl" deleted
pod "mysql-2-deploy" deleted
replicationcontroller "mysql-1" deleted
replicationcontroller "mysql-2" deleted
```

```
service "mysql" deleted
deploymentconfig.apps.openshift.io "mysql" deleted
[student@workstation ~]$ oc delete secrets mysql
secret "mysql" deleted
[student@workstation ~]$ oc new-app --template=mysql-ephemeral -p
MYSQL_USER=user1 -p MYSQL_PASSWORD=mypa55 -p MYSQL_DATABASE=testdb -p
MYSQL_ROOT_PASSWORD=r00tpa55
```

```
[student@workstation ~]$ oc get pods
NAME          READY STATUS   RESTARTS AGE
mysql-1-deploy 0/1   Completed 0       2m26s
mysql-1-kgx5g  1/1   Running    0       2m23s
```

```
[student@workstation ~]$ oc logs mysql-1-kgx5g
[student@workstation ~]$ nohup oc port-forward mysql-1-kgx5g 3306:3306 &
[1] 3626
[student@workstation ~]$ nohup: ignoring input and appending output to 'nohup.out'
```

```
[student@workstation ~]$ mysql -uuser1 -pmypa55 --protocol tcp -h localhost
mysql: [Warning] Using a password on the command line interface can be insecure.
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 35
Server version: 8.0.21 Source distribution
```

Copyright (c) 2000, 2019, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> show databases;
+-----+
| Database          |
+-----+
| information_schema |
| testdb            |
+-----+
2 rows in set (0.00 sec)
```

```
mysql> exit
Bye
```

```
[student@workstation ~]$ fg  
nohup oc port-forward mysql-1-kgx5g 3306:3306  
^C[student@workstation ~]$
```

CLEANUP

```
[student@workstation ~]$ oc delete project ${RHT_OCP4_DEV_USER}-mysql-openshift  
project.project.openshift.io "rhn-support-ablum-mysql-openshift" deleted  
[student@workstation ~]$ lab openshift-resources finish
```

Completing the Guided Exercise: Deploying a Database Server on OpenShift

- Deleting the 'rhn-support-ablum-mysql-openshift' project.... SUCCESS

How can we better understand the different networking layers in k8s ?

WITHIN A CONTAINER

containers within a Pod can all reach each other's ports on localhost.

This also means that containers within a Pod must coordinate port usage, but this is no different from processes in a VM


```
[student@workstation php-helloworld]$ oc rsh php-helloworld-1-wk49d
```

```
sh-4.2$ lsof -i -P
```

COMMAND	PID	USER	FD	TYPE	DEVICE	SIZE/OFF	NODE	NAME
httpd	1	1008920000	3u	IPv4	54259547	0t0	TCP	*:8080 (LISTEN)
httpd	1	1008920000	4u	IPv4	54259553	0t0	TCP	*:8443 (LISTEN)

```
sh-4.2$ curl localhost:8080
```

```
Hello, World! php version is 7.3.11
```

```
sh-4.2$ curl -k https://localhost:8443
```

```
Hello, World! php version is 7.3.11
```

Processes running within the same container share the same networking in the same way as processes running in the same virtual machine:

```
sh-4.2$ ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
1008920+	1	0	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	35	1	0	18:35	?	00:00:00	/usr/bin/cat
1008920+	36	1	0	18:35	?	00:00:00	/usr/bin/cat
1008920+	37	1	0	18:35	?	00:00:00	/usr/bin/cat
1008920+	38	1	0	18:35	?	00:00:00	/usr/bin/cat
1008920+	39	1	0	18:35	?	00:00:01	httpd -D FOREGROUND
1008920+	46	1	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	55	1	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	56	1	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	57	1	0	18:35	?	00:00:01	httpd -D FOREGROUND
1008920+	58	1	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	66	1	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	69	1	0	18:35	?	00:00:00	httpd -D FOREGROUND
1008920+	103	1	0	19:20	?	00:00:00	httpd -D FOREGROUND
1008920+	109	0	0	19:20	pts/0	00:00:00	/bin/sh
1008920+	215	109	0	19:42	pts/0	00:00:00	ps -ef

Let's modify the deploymentconfig and add another container running in the same pod:

```
[student@workstation php-helloworld]$ oc edit
```

```
deploymentconfig.apps.openshift.io/php-helloworld
```

```
- image:
```

```
image-registry.openshift-image-registry.svc:5000/rhn-support-ablum-route/php-helloworld@sha256:40891a7ac25f42cb7d90e6b6959a054cca04e28c9bafbe77ffdfb60327d00f32
```

```
imagePullPolicy: Always
```

```
name: php-helloworld
ports:
- containerPort: 8080
  protocol: TCP
- containerPort: 8443
  protocol: TCP
resources: {}
terminationMessagePath: /dev/termination-log
terminationMessagePolicy: File
- image: registry.access.redhat.com/rhel7
  imagePullPolicy: Always
  name: test1
  command: ["/usr/bin/sleep", "5000"]
  resources: {}
dnsPolicy: ClusterFirst
terminationGracePeriodSeconds: 30
```

```
[student@workstation php-helloworld]$ oc describe pod php-helloworld-3-9gl4f
IP: 10.129.21.132
```

```
[student@workstation php-helloworld]$ oc rsh -c test1 pod/php-helloworld-3-9gl4f
sh-4.2$
sh-4.2$
```

```
sh-4.2$ curl 10.129.21.132:8080
Hello, World! php version is 7.1.30
```

```
sh-4.2$ curl localhost:8080
Hello, World! php version is 7.1.30
```

Wait....what ? The httpd process isn't running here....

```
sh-4.2$ ps -ef
UID    PID  PPID  C STIME TTY          TIME CMD
1008920+   1    0  0 20:11 ?    00:00:00 /usr/bin/sleep 5000
1008920+  28    0  0 20:21 pts/0    00:00:00 /bin/sh
1008920+  39   28  0 20:22 pts/0    00:00:00 ps -ef
```

Again. Containers in the same pod are like processes in the same VM

POD TO POD

Let's fire up a new pod. This can help with port management (different procs trying to listen on the same port inside the same pod) along with providing scaling benefits:

```
[student@workstation ~]$ oc new-app --name mypod  
--docker-image="registry.access.redhat.com/rhel7"
```

We will need to edit the pod to add a command that

```
[student@workstation ~]$ oc edit deployment mypod  
spec:  
  containers:  
  - image:  
    registry.access.redhat.com/rhel7@sha256:d8e52fab67bc27384fe5f4022e8c5e5d83a7901b83df  
    7ed41d4a216aea57b44c  
    imagePullPolicy: Always  
    name: mypod  
    command: ["/usr/bin/sleep","5000"]
```

```
[student@workstation ~]$ oc get pods -l app=mypod -o wide  
NAME          READY STATUS  RESTARTS AGE   IP             NODE  
NOMINATED NODE READINESS GATES  
mypod-2-6hgpn 1/1 Running 0         2m1s 10.130.21.64   ip-10-0-220-55.ec2.internal  
<none>        <none>
```

```
[student@workstation ~]$ oc get pods -l app=php-helloworld -o wide  
NAME          READY STATUS  RESTARTS AGE   IP             NODE  
NOMINATED NODE READINESS GATES  
php-helloworld-3-9gl4f 2/2 Running 0         34m 10.129.21.132  
ip-10-0-218-157.ec2.internal <none>    <none>
```

So, two separate pods, with different ips but both in the "Pod SDN".

Let's test Pod to Pod communications by connecting to the php-helloworld's IP from a shell in mypod's only container:

```
[student@workstation ~]$ oc rsh mypod-2-6hgpn  
sh-4.2$ curl 10.129.21.132:8080
```

Hello, World! php version is 7.1.30

This works. So, pod-to-pod communication work using the pod SDN.

But, what happens when pods are restarted...

```
[student@workstation ~]$ oc delete pod php-helloworld-3-9gl4f
pod "php-helloworld-3-9gl4f" deleted
[student@workstation ~]$ oc get pods -l app=php-helloworld -o wide
NAME                                READY STATUS   RESTARTS  AGE  IP              NODE
NOMINATED NODE  READINESS GATES
php-helloworld-3-5vx7d  2/2 Running    0          43s  10.129.21.183  ip-10-0-218-157.ec2.internal  <none>      <none>
```

The IP associated with the php-helloworld pod has changed.

```
[student@workstation ~]$ oc rsh mypod-2-6hgpn
sh-4.2$ curl 10.129.21.132:8080
curl: (7) Failed connect to 10.129.21.132:8080; No route to host
```

So, it would be difficult for containers running in different pods to keep track of the podIP as it might change over time as pods are restarted or scaled.

What can help ? Services:

POD TO SERVICE

```
[student@workstation php-helloworld]$ oc get svc
NAME          TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
php-helloworld ClusterIP      172.30.237.166  <none>       8080/TCP,8443/TCP 110m
[student@workstation php-helloworld]$ oc rsh -c test1 pod/php-helloworld-3-9gl4f

sh-4.2$ curl 172.30.237.166:8080
Hello, World! php version is 7.1.30
```

Pods are able to use the service IP which will load balance the requests to pods that are labeled the same as the service selector.

Under the hood, a kubernetes service called “kube-proxy” handles this using iptables -L -t nat rules on the schedulable nodes.

How can you discover a service ?

1. Using oc commands (ie using the openshift API)

```
[student@workstation ~]$ oc get service
NAME          TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
php-helloworld ClusterIP      172.30.237.166  <none>       8080/TCP,8443/TCP 164m
[student@workstation ~]$ oc get endpoints
NAME          ENDPOINTS          AGE
php-helloworld 10.129.21.183:8443,10.129.21.183:8080 165m
```

2. DNS each service is dynamically assigned:

SVC_NAME.PROJECT_NAME.svc.cluster.local

```
[student@workstation ~]$ oc get svc
NAME          TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
php-helloworld ClusterIP      172.30.237.166  <none>       8080/TCP,8443/TCP 179m
```

```
[student@workstation ~]$ oc project
Using project "rhn-support-ablum-route" on server
"https://api.ocp-na2.prod.nextcle.com:6443".
```

```
[student@workstation ~]$ oc rsh mypod-2-6hgpn
sh-4.2$ curl php-helloworld.rhn-support-ablum-route.svc.cluster.local:8080
Hello, World! php version is 7.1.30
```

```
sh-4.2$ cat /etc/resolv.conf
search rhn-support-ablum-route.svc.cluster.local svc.cluster.local cluster.local ec2.internal
nameserver 172.30.0.10
options ndots:5
```

3. Environment variables

For each service inside an OpenShift project, the following environment variables are automatically defined and injected into containers for all pods inside the same project:

- SVC_NAME_SERVICE_HOST is the service IP address.
- SVC_NAME_SERVICE_PORT is the service TCP port.

```
sh-4.2$ env | grep PHP_HELLOWORLD
PHP_HELLOWORLD_PORT=tcp://172.30.237.166:8080
PHP_HELLOWORLD_SERVICE_PORT_8443_TCP=8443
```

```

PHP_HELLOWORLD_SERVICE_PORT_8080_TCP=8080
PHP_HELLOWORLD_PORT_8443_TCP_ADDR=172.30.237.166
PHP_HELLOWORLD_PORT_8080_TCP_PORT=8080
PHP_HELLOWORLD_SERVICE_HOST=172.30.237.166
PHP_HELLOWORLD_SERVICE_PORT=8080
PHP_HELLOWORLD_PORT_8443_TCP=tcp://172.30.237.166:8443
PHP_HELLOWORLD_PORT_8080_TCP_ADDR=172.30.237.166
PHP_HELLOWORLD_PORT_8443_TCP_PORT=8443
PHP_HELLOWORLD_PORT_8080_TCP=tcp://172.30.237.166:8080
PHP_HELLOWORLD_PORT_8443_TCP_PROTO=tcp
PHP_HELLOWORLD_PORT_8080_TCP_PROTO=tcp

```

```
sh-4.2$ curl
```

```

${PHP_HELLOWORLD_SERVICE_HOST}:${PHP_HELLOWORLD_SERVICE_PORT}
Hello, World! php version is 7.1.30

```

NOTE: The pod must be created AFTER a given service or environment variables won't be available. For example,

```

[student@workstation ~]$ oc create service clusterip myservice --tcp=12345:8080
service/myservice created

```

```
[student@workstation ~]$ oc get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
myservice	ClusterIP	172.30.96.95	<none>	12345/TCP	11s
php-helloworld	ClusterIP	172.30.237.166	<none>	8080/TCP,8443/TCP	3h7m

Let's see if there are any environment variables with MYSERVICE:

```
[student@workstation ~]$ oc rsh mypod-2-6hgpn
```

```
sh-4.2$ env | grep MYSERVICE
```

```
sh-4.2$
```

```
sh-4.2$ exit
```

```
exit
```

```
command terminated with exit code 1
```

Let's try again after restarting a pod:

```
[student@workstation ~]$ oc delete pod mypod-2-6hgpn
```

```
pod "mypod-2-6hgpn" deleted
```

```
[student@workstation ~]$ oc get pods -l app=mypod
```

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

mypod-2-s8wqw 1/1 Running 0 58s

```
[student@workstation ~]$ oc exec mypod-58d86ccb46-prw7k -- env | grep MYSERVICE
```

```
MYSERVICE_PORT_12345_TCP=tcp://172.30.96.95:12345
```

```
MYSERVICE_SERVICE_HOST=172.30.96.95
```

```
MYSERVICE_SERVICE_PORT_12345_8080=12345
```

```
MYSERVICE_PORT=tcp://172.30.96.95:12345
```

```
MYSERVICE_PORT_12345_TCP_ADDR=172.30.96.95
```

```
MYSERVICE_PORT_12345_TCP_PORT=12345
```

```
MYSERVICE_PORT_12345_TCP_PROTO=tcp
```

```
MYSERVICE_SERVICE_PORT=12345
```

(of course there are no endpoints so http traffic wont actually get handled)

```
[student@workstation ~]$ oc get endpoints
```

NAME	ENDPOINTS	AGE
myservice	<none>	5m23s
php-helloworld	10.129.21.183:8443,10.129.21.183:8080	3h12m

EXTERNAL TO SERVICE (ROUTE)

The “Ingress router” uses DNS wildcards to expose a service external to the OCP cluster. The route created points to a service (that is managed by kube-proxy).

```
[student@workstation ~]$ oc get routes
```

NAME	HOST/PORT	PATH	SERVICES	PORT
	TERMINATION	WILDCARD		
php-helloworld	php-helloworld-rhn-support-ablum-route		apps.ocp-na2.prod.nextcle.com	
php-helloworld	8080-tcp	None		

```
[student@workstation ~]$ oc describe route php-helloworld
```

```
Name: php-helloworld
Namespace: rhn-support-ablum-route
Created: 3 hours ago
Labels: app=php-helloworld
Annotations: openshift.io/host.generated=true
Requested Host: php-helloworld-rhn-support-ablum-route.apps.ocp-na2.prod.nextcle.com
                exposed on router default (host apps.ocp-na2.prod.nextcle.com) 3 hours ago
Path: <none>
```

TLS Termination: <none>
Insecure Policy: <none>
Endpoint Port: 8080-tcp

Service: php-helloworld
Weight: 100 (100%)
Endpoints: 10.129.21.183:8443, 10.129.21.183:8080

Suppose we want to change the default so that a different hostname is used (maybe one we control the DNS records for). To configure the ingress router to use that name to point to a given service lets use:

```
[student@workstation ~]$ source /usr/local/etc/ocp4.config  
[student@workstation ~]$ oc expose service php-helloworld  
--name=${RHT_OCP4_DEV_USER}-xyz  
route.route.openshift.io/rhn-support-ablum-xyz expose
```

```
[student@workstation ~]$ curl  
rhn-support-ablum-xyz-rhn-support-ablum-route.apps.ocp-na2.prod.nextcle.com  
Hello, World! php version is 7.1.30
```

Of course, you can't just use ANY fqdn, even those you don't manage (or own). Consider:

```
[student@workstation ~]$ oc expose service php-helloworld  
--hostname=www.example.com  
route.route.openshift.io/php-helloworld exposed
```

```
[student@workstation ~]$ oc get routes
```

NAME	HOST/PORT	PATH
SERVICES	PORT	TERMINATION WILDCARD
php-helloworld	www.example.com	
php-helloworld	8080-tcp	None
rhn-support-ablum-xyz		
rhn-support-ablum-xyz-rhn-support-ablum-route.apps.ocp-na2.prod.nextcle.com		
php-helloworld	8080-tcp	None

```
[student@workstation ~]$ oc describe route php-helloworld
```

Name: php-helloworld
Namespace: rhn-support-ablum-route
Created: About a minute ago
Labels: app=php-helloworld
Annotations: <none>
Requested Host: www.example.com


```
[student@workstation ~]$ curl www.example.com
curl: (6) Could not resolve host: www.example.com; Unknown error
```

For routes secured using x509 certificates consider:

```
[student@workstation ~]$ oc create route -h
```

```
[student@workstation php-helloworld]$ oc expose service php-helloworld --name ablum-test
--hostname=rhn-support-ablum-foo.apps.na45.prod.nextcle.com
route.route.openshift.io/ablum-test exposed
```

To clean up run:

```
[student@workstation ~]$ lab openshift-routes finish
```

In SUMMARY:

<https://kubernetes.io/docs/concepts/cluster-administration/networking/>

There are really 4 networking problems to solve:

1. Highly-coupled container-to-container communications: this is solved by [pods](#) and localhost communications.
2. Pod-to-Pod communications: this is the primary focus of this document.
3. Pod-to-Service communications: this is covered by [services](#).
4. External-to-Service communications: this is covered by [services](#).

CREATING APPLICATIONS WITH SOURCE-TO-IMAGE

Study figure 6.8

1. Start a container from a base container image called the builder image, which includes a programming language runtime and essential development tools such as compilers and package managers.

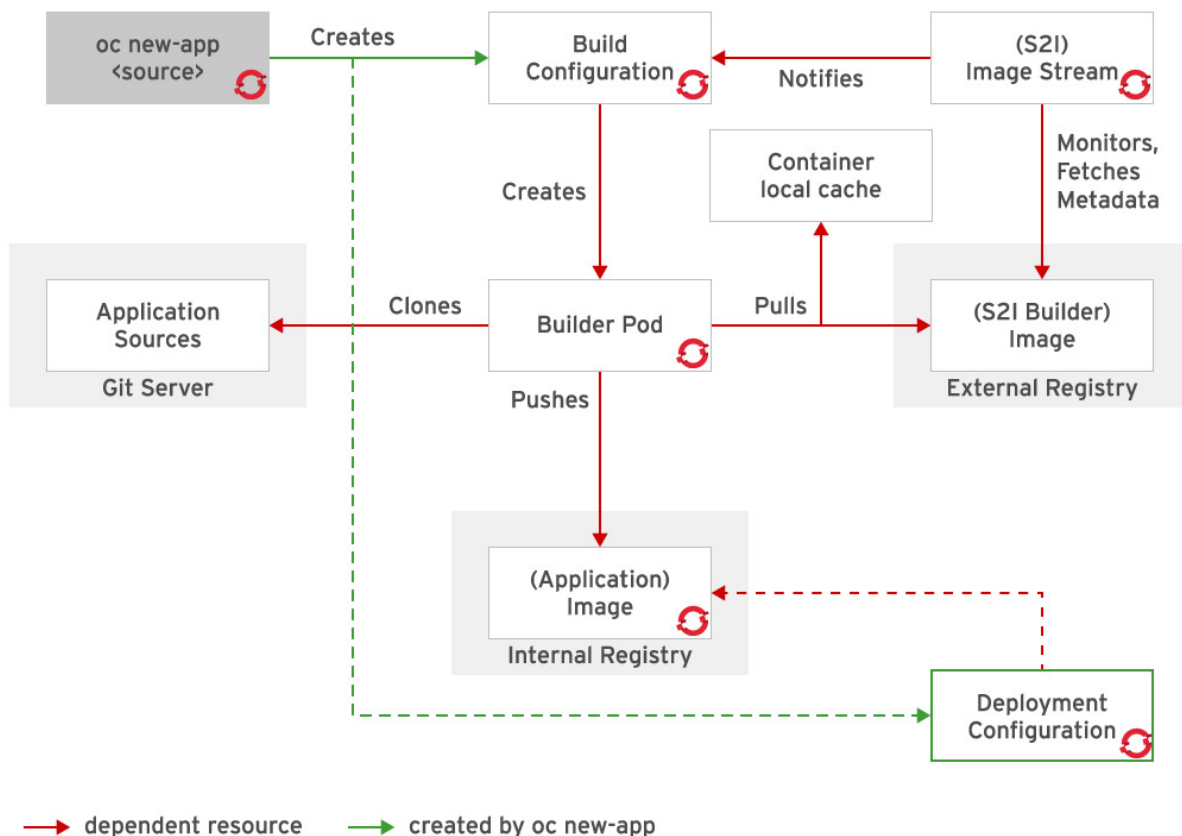
2. Fetch the application source code, usually from a Git server, and send it to the container.

3. Build the application binary files inside the container.

Build tools are run at this point:

- Yum
- Bundler (for ruby)
- Maven (for java)
- Npm (nodejs)
- Gcc (for C programs)

4. Save the container, after some clean up, as a new container image, which includes the programming language runtime and the application binaries.



GUIDED PRACTICE

```
[student@workstation ~]$ lab openshift-s2i start
```

The first part of the s2i process involves building a container.

How to inspect a builder image ?

Let's inspect a sample "builder" image that assists in building a container for a webserver serving up a web application written in php.

Navigate to <https://catalog.redhat.com/software/containers/explore> and search for php, check the "builder" facet in the search

Look for the card rhsc/php-73-rhel7. Click get this image.

Let's pull it with podman so we can easily inspect it. We will have to log in to get this one:

```
[root@workstation ~]# podman pull registry.access.redhat.com/rhsc/php-73-rhel7
Trying to pull registry.access.redhat.com/rhsc/php-73-rhel7...
unsupported: This repo requires terms acceptance and is only available on registry.redhat.io
Error: error pulling image "registry.access.redhat.com/rhsc/php-73-rhel7": unable to pull
registry.access.redhat.com/rhsc/php-73-rhel7: unable to pull image: Error initializing source
docker://registry.access.redhat.com/rhsc/php-73-rhel7:latest: Error reading manifest latest in
registry.access.redhat.com/rhsc/php-73-rhel7: unsupported: This repo requires terms
acceptance and is only available on registry.redhat.io
```

```
[root@workstation ~]# podman login registry.redhat.io
Username: rhn-support-ablum
Password:
Login Succeeded!
[root@workstation ~]# podman pull registry.redhat.io/rhsc/php-73-rhel7
```

```
[root@workstation ~]# podman inspect fd | less
```

```
"STI_SCRIPTS_URL=image:///usr/libexec/s2i",
```

```

"STI_SCRIPTS_PATH=/usr/libexec/s2i",

"Cmd": [
    "/bin/sh",
    "-c",
    "$STI_SCRIPTS_PATH/usage"
]

```

[root@workstation ~]# **podman run fd**
 This is a S2I PHP-7.3 rhel base image:

To use it in Openshift, run:

```
oc new-app php:7.3~https://github.com/sclorg/cakephp-ex.git
```

To access the application:

```
oc get pods
oc exec <pod> -- curl 127.0.0.1:8080
```

Alternatively, to run the image directly using podman or docker, or how to use it as a parent image in a Dockerfile, see documentation at <https://github.com/sclorg/s2i-php-container/blob/master/7.3/README.md>

```
[root@workstation ~]# podman run -it --entrypoint /bin/bash --user 0 fd
```

```

bash-4.2# cd /usr/libexec/s2i/
bash-4.2# ls
assemble run save-artifacts usage
bash-4.2# cat usage
bash-4.2# cat run
#!/bin/bash
(this is what will be run later in the deployed application)

```

```
bash-4.2# cat assemble
```

```
#!/bin/bash
```

```
set -e
```

```
shopt -s dotglob
```

```
echo "---> Installing application source..."
```

```
mv /tmp/src/* ./ <-- Moves the php source code to the  
container working dir ..... THINK COPY Dockerfile instruction
```

```
if [ -f composer.json ]; then
```

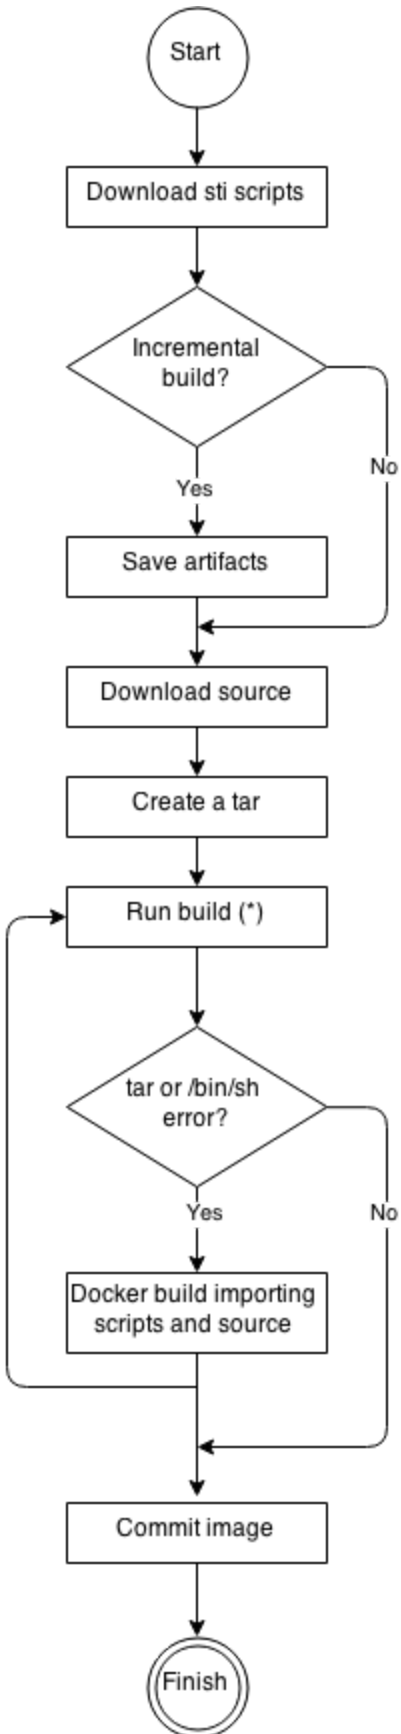
```
echo "Found 'composer.json', installing dependencies using composer.phar..."
```

```
^-- Installs dependencies needed by the application .... THINK various RUN instructions
```

```
...SNIP...
```

From

https://docs.openshift.com/container-platform/4.2/builds/build-strategies.html#builds-strategy-s2i-override-builder-image-scripts_build-strategies



```
[student@workstation 5.6]$ cd ~/DO180-apps/
[student@workstation DO180-apps]$ git checkout master
Already on 'master'
[student@workstation DO180-apps]$
[student@workstation DO180-apps]$
[student@workstation DO180-apps]$ git checkout -b s2i
Switched to a new branch 's2i'
[student@workstation DO180-apps]$ git push -u origin s2i
Username for 'https://github.com': ajblum
Password for 'https://ajblum@github.com':
Total 0 (delta 0), reused 0 (delta 0)
[student@workstation DO180-apps]$ cat php-helloworld/index.php
<?php
print "Hello, World! php version is " . PHP_VERSION . "\n";
?>
```

```
[student@workstation DO180-apps]$ source /usr/local/etc/ocp4.config
```

```
[student@workstation DO180-apps]$ oc login -u ${RHT_OCP4_DEV_USER} -p
${RHT_OCP4_DEV_PASSWORD} ${RHT_OCP4_MASTER_API}
Login successful.
```

You don't have any projects. You can try to create a new project, by running

```
oc new-project <projectname>
```

```
[student@workstation DO180-apps]$ oc new-project ${RHT_OCP4_DEV_USER}-s2i
```

```
[student@workstation DO180-apps]$ oc new-app -h
```

-i, --image-stream=[]: Name of an image stream to use in the app.

--context-dir="": Context directory to be used for the build.

--name="": Set name to use for generated application artifacts

```
[student@workstation DO180-apps]$ oc get is -n openshift
```

NOTE: image streams don't actually hold the container image data but they are more powerful in that they are aware of changes in the repo that can be used as a trigger for (re)deployment of an application.

OpenShift components such as builds and deployments can watch an image stream to receive notifications when new images are added and react by performing a build or a deployment.

<https://blog.openshift.com/image-streams-faq/>

```
[student@workstation DO180-apps]$ oc get is -n openshift php
NAME IMAGE REPOSITORY TAGS
UPDATED
php default-route-openshift-image-registry.apps.ocp-na2.prod.nextcle.com/openshift/php
7.0,7.1,7.2,latest 2 months ago
```

```
[student@workstation DO180-apps]$ oc describe is php -n openshift
```

7.3 (latest)

tagged from registry.redhat.io/rhscpl/php-73-rhel7:latest
prefer registry pullthrough when referencing this tag

Build and run PHP 7.3 applications on RHEL 7. For more information about using this builder image, including OpenShift considerations, see

<https://github.com/sclorg/s2i-php-container/blob/master/7.3/README.md>

7.2

tagged from registry.redhat.io/rhscpl/php-72-rhel7:latest
prefer registry pullthrough when referencing this tag

Build and run PHP 7.2 applications on RHEL 7. For more information about using this builder image, including OpenShift considerations, see

<https://github.com/sclorg/s2i-php-container/blob/master/7.2/README.md>

Let's use the latest image stream tag 7.3:

```
[student@workstation DO180-apps]$ oc new-app -i php:7.3 --name=php-helloworld
https://github.com/${RHT_OCP4_GITHUB_USER}/DO180-apps#s2i --context-dir
php-helloworld
```

```
[student@workstation DO180-apps]$ oc get pods
NAME READY STATUS RESTARTS AGE
```


php-helloworld-1-build 1/1 Running 0 18s

```
[student@workstation DO180-apps]$ oc get bc
NAME          TYPE FROM          LATEST
php-helloworld Source Git@s2i 1
[student@workstation DO180-apps]$ oc logs -f bc/php-helloworld
```

...

Writing manifest to image destination

Storing signatures

Successfully pushed

image-registry.openshift-image-registry.svc:5000/rhn-support-ablum-s2i/php-helloworld@sha256:d39d3b8e0e1d5b95c7ef06b755bdbd538bae0b8139e12ef6b138c5fc5238c51f

Push successful

```
[student@workstation DO180-apps]$ oc logs deployment.apps/php-helloworld -f
-> Cgroups memory limit is set, using HTTPD_MAX_REQUEST_WORKERS=102
..SNIP..
[Fri Jan 24 18:15:47.579737 2020] [mpm_prefork:notice] [pid 1] AH00163: Apache/2.4.34 (Red Hat) OpenSSL/1.0.2k-fips configured -- resuming normal operations
[Fri Jan 24 18:15:47.579776 2020] [core:notice] [pid 1] AH00094: Command line: 'httpd -D FOREGROUND'
```

How did the builder pod build the image ?

Let's look at that builder pod and what containers ran there:

```
[student@workstation DO180-apps]$ oc get pod php-helloworld-1-build -o
go-template='{{range .spec.InitContainers}}{{.name}}{{end}}'
sti-build
```

But, there also "init containers" in there responsible for pulling the source code (git) and creating the Dockerfile that includes in its FROM the builder image we gave on the command line:

```
[student@workstation DO180-apps]$ oc get pod php-helloworld-1-build -o
go-template='{{range .spec.initContainers}} {{.name}} {{end}}'
git-clone manage-dockerfile
```

The image used for these are all the same.

If we wanted to inject special environment variables into the build we could build include those in our builder image or customize the assemble scripts. But, if we want just a certain env variable set, we could use from `oc new-app -h`:

--build-env=[]: Specify a key-value pair for an environment variable to set into each build image.

Useful for programs like npm (package manager for nodejs) to customize the registry used.

```
[student@workstation DO180-apps]$ oc get service
NAME                TYPE          CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
php-helloworld      ClusterIP     172.30.201.129  <none>       8080/TCP,8443/TCP 4m10s
[student@workstation DO180-apps]$ oc expose service php-helloworld --name
${RHT_OCP4_DEV_USER}-helloworld
route.route.openshift.io/rhn-support-ablum-helloworld exposed
[student@workstation DO180-apps]$ oc get route
NAME                HOST/PORT          PATH
SERVICES    PORT      TERMINATION  WILDCARD
rhn-support-ablum-helloworld
rhn-support-ablum-helloworld-rhn-support-ablum-s2i.apps.na45.prod.nextcle.com
php-helloworld  8080-tcp      None
[student@workstation DO180-apps]$ curl
rhn-support-ablum-helloworld-rhn-support-ablum-s2i.apps.na45.prod.nextcle.com
Hello, World! php version is 7.3.11
```

Let's change the source code of our php application and start a new s2i build.

```
[student@workstation DO180-apps]$ cd ~/DO180-apps/php-helloworld/
[student@workstation php-helloworld]$ vim index.php
```

```
<?php
print "Hello, World! php version is " . PHP_VERSION . "<br>";
print "ablum was here";
?>
```

```
[student@workstation php-helloworld]$ git add .
[student@workstation php-helloworld]$ git commit -m 'Added a message to index page
content.'
[s2i e5fde79] Added a message to index page content.
```

1 file changed, 1 insertion(+)

```
[student@workstation php-helloworld]$ git push origin s2i
```

Counting objects: 7, done.

Delta compression using up to 2 threads.

Compressing objects: 100% (3/3), done.

Writing objects: 100% (4/4), 419 bytes | 0 bytes/s, done.

Total 4 (delta 1), reused 0 (delta 0)

remote: Resolving deltas: 100% (1/1), completed with 1 local object.

To <https://github.com/ajblum/DO180-apps.git>

f7cd896..e5fde79 s2i -> s2i

Now, kick off another build:

```
student@workstation php-helloworld]$ oc get builds
```

NAME	TYPE	FROM	STATUS	STARTED	DURATION
php-helloworld-1	Source	Git@f7cd896	Complete	23 minutes ago	2m38s

```
[student@workstation php-helloworld]$ oc get bc
```

NAME	TYPE	FROM	LATEST
php-helloworld	Source	Git@s2i	1

```
[student@workstation php-helloworld]$ oc start-build php-helloworld
```

build.build.openshift.io/php-helloworld-2 started

```
[student@workstation php-helloworld]$ oc get builds
```

NAME	TYPE	FROM	STATUS	STARTED	DURATION
php-helloworld-1	Source	Git@f7cd896	Complete	23 minutes ago	2m38s
php-helloworld-2	Source	Git@s2i	Pending		

```
[student@workstation php-helloworld]$ oc logs -f bc/php-helloworld
```

...SNIP...

Successfully pushed

image-registry.openshift-image-registry.svc:5000/rhn-support-ablum-s2i/php-helloworld@sha256:2eb9366fb146a7bdd9b8ae08c77203b95f3fb5127ab8c321a87d279c3a62e58d

Push successful

The new build result in a new image to the

default-route-openshift-image-registry.apps.ocp-na2.prod.nextcle.com/rhn-support-ablum-s2i/php-helloworld repository.

Consider the imagestream:

```
[student@workstation php-helloworld]$ oc describe is php-helloworld
```

As a result of this change a new deployment is triggered:

```
[student@workstation php-helloworld]$ oc get dc
NAME          REVISION  DESIRED  CURRENT  TRIGGERED BY
php-helloworld 2        1        1        config,image/php-helloworld:latest
[student@workstation php-helloworld]$ oc logs -f dc/php-helloworld
```

And a new pod is running:

```
[student@workstation php-helloworld]$ oc get pods
NAME                READY  STATUS   RESTARTS  AGE
php-helloworld-1-build 0/1    Completed 0         26m
php-helloworld-1-deploy 0/1    Completed 0         24m
php-helloworld-2-build 0/1    Completed 0         3m9s
php-helloworld-2-deploy 0/1    Completed 0         40s
php-helloworld-2-p48dv 1/1    Running   0         30s
```

What about the message ?

```
[student@workstation php-helloworld]$ curl
rhn-support-ablum-helloworld-rhn-support-ablum-s2i.apps.ocp-na2.prod.nextcle.com
Hello, World! php version is 7.2.24
ablum was here
```

Let's clean up:

```
[student@workstation php-helloworld]$ lab openshift-s2i finish
```

EXTRA PRACTICE WITH S2I

```
[student@workstation ~]$ oc new-project ablum-s2i-practice
```

```
[student@workstation ~]$ oc new-app -h
```

Create a Ruby application based on the provided [image]~[source code] combination
oc new-app centos/ruby-25-centos7~https://github.com/sclorg/ruby-ex.git

Navigate to <https://github.com/sclorg/ruby-ex> and show \$ oc new-app
openshift/ruby:25~https://github.com/< yourusername >/ruby-ex

“These steps assume your OpenShift deployment has the default set of ImageStreams defined”

Where are the default imagestream resources ? In the openshift project:

```
[student@workstation ~]$ oc get is -n openshift
```

```
[student@workstation ~]$ oc get is ruby -n openshift
```

```
NAME IMAGE REPOSITORY
```

```
TAGS
```

```
UPDATED
```

```
ruby default-route-openshift-image-registry.apps.ocp-na2.prod.nextcle.com/openshift/ruby  
2.3,2.4,2.5,latest 2 months ago
```

```
[student@workstation ~]$ oc describe is ruby -n openshift
```

```
2.5 (latest)
```

```
tagged from registry.redhat.io/rhsc/ruby-25-rhel7:latest
```

```
prefer registry pullthrough when referencing this tag
```

This is a “builder” image. See

<https://catalog.redhat.com/software/containers/search?q=rhsc/ruby>

Builder Image

Platform for building and running Ruby 2.5 applications

```
[student@workstation ~]$ oc -o json new-app
```

```
openshift/ruby:2.5~https://github.com/sclorg/ruby-ex.git
```

Study figure 6.8 while inspecting

```
[student@workstation ~]$ oc new-app
```

```
openshift/ruby:2.5~https://github.com/sclorg/ruby-ex.git
```

```
[student@workstation ~]$ oc get all
```

```
[student@workstation ~]$ oc logs pod/ruby-ex-1-build -f
```

```
Getting image source signatures
```

```
Copying blob
```

```
sha256:8ba884070f611d31cb2c42eddb691319dc9facf5e0ec67672fcfa135181ab3df
```

```
[student@workstation ~]$ oc logs bc/ruby-ex -f
```

Refer again to figure 6.8 during the build

<https://github.com/sclorg/ruby-ex>

To trigger a new build:

```
$ oc get bc
```

```
$ oc start-build <build_name>
```

```
[student@workstation ~]$ oc expose service/ruby-ex
```

```
[student@workstation ~]$ oc get routes
```

NAME	HOST/PORT	PATH	SERVICES	PORT
TERMINATION	WILDCARD			
ruby-ex	ruby-ex-ablum-s2i-practice.apps.ocp-na2.prod.nextcle.com		ruby-ex	
8080-tcp	None			

<http://ruby-ex-ablum-s2i-practice.apps.ocp-na2.prod.nextcle.com/>

Welcome to your Ruby application on OpenShift

Deploying code changes

The source code for this application is available to be forked from the [OpenShift GitHub repository](#). You can configure a webhook in your repository to make OpenShift automatically start a build whenever you push your code:

1. From the Web Console homepage, navigate to your project
2. Click on Browse > Builds

Managing your application

Documentation on how to manage your application from the Web Console or Command Line is available at the [Developer Guide](#).

Web Console

You can use the Web Console to view the state of your application components and launch new builds.

Let's look at the builder image that was used in myfourthapp:

```
[student@workstation ~]$ oc describe buildconfig.build.openshift.io/ruby-ex
```

```
[student@workstation ~]$ oc describe build.build.openshift.io/ruby-ex-1
```

Strategy: Source

URL: https://github.com/sclorg/ruby-ex.git

Commit: c00ecd7 (Merge pull request #25 from pvalena/master)

Author/Committer: Honza Horak / GitHub

From Image: DockerImage

image-registry.openshift-image-registry.svc:5000/openshift/ruby@sha256:9866398704db9207862bdb930b1dba4139dbaf71c6eaa6d084ea036478b28de9

```
[student@workstation ~]$ oc get is -n openshift
```

```
[student@workstation ~]$ oc describe is ruby -n openshift
```

Image Repository: image-registry.openshift-image-registry.svc:5000/openshift/ruby
2.5 (latest)

tagged from [docker.io/centos/ruby-25-centos7](https://registry.hub.docker.io/r/centos/ruby-25-centos7):latest
prefer registry pullthrough when referencing this tag

Build and run Ruby 2.5 applications on CentOS 7. For more information about using this builder image, including OpenShift considerations, see <https://github.com/sclorg/s2i-ruby-container/blob/master/2.5/README.md>.

Tags: builder, ruby

Supports: ruby:2.5, ruby

Example Repo: <https://github.com/sclorg/ruby-ex.git>

*

[docker.io/centos/ruby-25-centos7@sha256:9866398704db9207862bdb930b1dba4139dbaf71c6eaa6d084ea036478b28de9](https://registry.hub.docker.io/r/centos/ruby-25-centos7@sha256:9866398704db9207862bdb930b1dba4139dbaf71c6eaa6d084ea036478b28de9)

3 months ago

Let's pull this image from [docker.io](https://registry.hub.docker.io/r/centos/ruby-25-centos7) and take a look at it:

```
[student@workstation ~]$ sudo podman pull
registry.redhat.io/rhsc/ruby-25-rhel7@sha256:d773b37c133b1a59dd11e69c801a8021bfcc6c065d1dcb8c19d2a42402596235
```

(maybe you need to login)

```
[student@workstation ~]$ sudo podman images
```

```
registry.redhat.io/rhsc/ruby-25-rhel7  none          7ca2412cdd19  3 months ago  543MB
```

```
[student@workstation ~]$ sudo podman inspect 7ca2412cdd19
```

...

```
    "STI_SCRIPTS_URL=image:///usr/libexec/s2i",
    "STI_SCRIPTS_PATH=/usr/libexec/s2i",
```

...

```
    "Entrypoint": [
        "container-entrypoint"
    ],
    "Cmd": [
        "/bin/sh",
        "-c",
        "$STI_SCRIPTS_PATH/usage"
    ],
```

```
[student@workstation ~]$ sudo podman run -it --entrypoint /bin/bash -u 0 7ca2412cdd19
```

```
bash-4.2$ cd /usr/libexec/s2i
```

```
bash-4.2$ ls
```

```
assemble run usage
```

```
bash-4.2$ cat assemble
```

```
...
```

```
echo "---> Running 'bundle install ${ADDTL_BUNDLE_ARGS}' ..."
```

```
bundle install --path ./bundle ${ADDTL_BUNDLE_ARGS}
```

```
...
```

Let's go back to the bc logs and compare:

```
[student@workstation ~]$ oc logs bc/ruby-ex | grep Running
```

```
---> Running 'bundle install --retry 2 --deployment --without development:test' ...
```

```
Running `bundle clean --verbose` with bundler 1.16.1
```

How to write your own s2i scripts?

How to create custom s2i builder images: <https://blog.openshift.com/create-s2i-builder-image/>

Inspecting a builder image, making a small change:

```
[root@workstation ~]# podman run -it -u root registry.redhat.io/ubi7/python-27 /bin/bash
```

```
(app-root)
```

```
(app-root)cd /usr/libexec/s2i/
```

```
(app-root)ls
```

```
assemble init-wrapper run usage
```

```
(app-root)vi /usr/libexec/s2i/assemble
```

```
# set permissions for any installed artifacts
```

```
fix-permissions /opt/app-root
```

```
touch /opt/app-root/ablum_was_here
```

```
(app-root)exit
```

```
exit
```

```
[root@workstation ~]# podman ps -a
```


CONTAINER ID	IMAGE	COMMAND	CREATED
14fe9771dd0c	registry.redhat.io/ubi7/python-27:latest	container-entrypoin...	3 minutes ago
Exited (127) 4 seconds ago	distracted_chandrasekhar	false	
b162b86d5f05	registry.redhat.io/ubi7/python-27:latest	container-entrypoin...	15 minutes ago
Exited (0) 4 minutes ago	happy_bohr	false	
17c997ad2b33	registry.redhat.io/ubi8/python-27:latest	container-entrypoin...	28 minutes ago
Exited (127) 20 minutes ago	fervent_euclid	false	
33365bae4853	registry.redhat.io/ubi8/python-27:latest	container-entrypoin...	30 minutes ago
Exited (0) 29 minutes ago	upbeat_aryabhata	false	

[root@workstation ~]# **podman commit 14fe9771dd0c do180:latest**

[root@workstation ~]# **podman push c487c0a98379 quay.io/ajblum/mytest:latest**

[student@workstation ~]\$ **oc new-project s2i-fun**

[student@workstation ~]\$ **oc new-app**
quay.io/ajblum/mytest~https://github.com/OpenShiftDemos/os-sample-python.git

https://docs.openshift.com/container-platform/4.2/builds/build-strategies.html#builds-strategy-s2i-override-builder-image-scripts_build-strategies

1. Provide an **assemble**, **run**, or **save-artifacts** script in the **.s2i/bin** directory of your application source repository, or
2. Provide a URL of a directory containing the scripts as part of the strategy definition. For example:

strategy:

sourceStrategy:

from:

kind: "ImageStreamTag"

name: "builder-image:latest"

scripts: "http://somehost.com/scripts_directory"

https://docs.openshift.com/container-platform/4.2/builds/build-strategies.html#images-create-s2i_build-strategies

Also, Look for ""**Writing s2i scripts**""

How can you use the s2i build images from redhat's registry?

\$ oc new-project javafun

Try:

```
[student@workstation ~]$ oc run openjdk11
```

```
--image=registry.redhat.io/openjdk/openjdk-11-rhel7
```

```
[student@workstation ~]$ oc get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/openjdk11-1-deploy	1/1	Running	0	5m
pod/openjdk11-1-fj8bt	0/1	ImagePullBackOff	0	5m

Create a "Registry Service Account" on <https://access.redhat.com/terms-based-registry/#/token/ablum-rhel8-training/openshift-secret>

```
[student@workstation ~]$ vi ablum-secret.yaml
```

```
secret/1979710-ablum-rhel8-training-pull-secret created
```

```
[student@workstation ~]$ oc create -f ablum-secret.yaml
```

```
[student@workstation ~]$ oc edit deploymentconfig.apps.openshift.io/openjdk11
```

...

spec:

containers:

- image: registry.redhat.io/openjdk/openjdk-11-rhel7

imagePullPolicy: Always

name: openjdk11

resources: {}

terminationMessagePath: /dev/termination-log

terminationMessagePolicy: File

dnsPolicy: ClusterFirst

restartPolicy: Always

schedulerName: default-scheduler

securityContext: {}

terminationGracePeriodSeconds: 30

imagePullSecrets:

- name: 1979710-ablum-rhel8-training-pull-secret

```
[student@workstation ~]$ oc get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/openjdk11-1-deploy	0/1	Error	0	21m
pod/openjdk11-2-hd9nb	0/1	CrashLoopBackOff	4	2m

```
[student@workstation ~]$ oc logs pod/openjdk11-2-hd9nb
```

Starting the Java application using /opt/jboss/container/java/run/run-java.sh ...

ERROR Neither \$JAVA_MAIN_CLASS nor \$JAVA_APP_JAR is set and 0 JARs found in /deployments (1 expected)

There isn't anything compiled (ie no jar) to execute. Why ? This is a "builder" image see:

<https://access.redhat.com/containers/?tab=images#/registry.access.redhat.com/openjdk/openjdk-11-rhel7>

So, let's use s2i to build an application using it. Let's try one from the openshift-quickstarts:

```
[student@workstation ~]$ oc delete deploymentconfig.apps.openshift.io/openjdk11
deploymentconfig.apps.openshift.io "openjdk11" deleted
[student@workstation ~]$ oc new-app
registry.redhat.io/openjdk/openjdk-11-rhel7~https://github.com/jboss-openshift/openshift-quickstarts --context-dir=undertow-servlet
```

W1002 12:26:19.215464 4811 dockerimagelookup.go:236] Docker registry lookup failed: Get https://registry.redhat.io/v2/openjdk/openjdk-11-rhel7/manifests/latest: unauthorized: Please login to the Red Hat Registry using your Customer Portal credentials. Further instructions can be found here: https://access.redhat.com/articles/3399531

error: unable to locate any images in image streams, local docker images with name "registry.redhat.io/openjdk/openjdk-11-rhel7"

```
[student@workstation ~]$ oc get serviceaccounts
```

NAME	SECRETS	AGE
builder 2	13m	
default 2	13m	
deployer 2	13m	

```
[student@workstation ~]$ oc get secrets
```

NAME	TYPE	DATA	AGE
1979710-ablum-rhel8-training-pull-secret	kubernetes.io/dockerconfigjson	1	27m
builder-token-db4fb	kubernetes.io/service-account-token	3	12m
builder-token-gc6dw	kubernetes.io/service-account-token	3	12m
default-dockercfg-62t9g	kubernetes.io/dockercfg	1	12m
default-token-5vkjr	kubernetes.io/service-account-token	3	12m
default-token-hf589	kubernetes.io/service-account-token	3	12m
deployer-dockercfg-ttl54	kubernetes.io/dockercfg	1	12m
deployer-token-9jt9s	kubernetes.io/service-account-token	3	12m
deployer-token-p9krp	kubernetes.io/service-account-token	3	12m

```
[student@workstation ~]$ oc get sa default -o yaml
```

```
apiVersion: v1
```

```
imagePullSecrets:
```

```
- name: default-dockercfg-62t9g
```

Go to Download [ablum-rhel8-training-secret.yaml](#) and paste contents into a .yaml:

```
$ vi ablum-pull-secret.yaml
```

```
apiVersion: v1
```

```
kind: Secret
```

```
metadata:
```

```
  name: 1979710-ablum-rhel8-training-pull-secret
```

```
data:
```

```
  .dockerconfigjson: eyJhdXRob3R5L...
```

```
type: kubernetes.io/dockerconfigjson
```

```
[student@workstation ~]$ oc create -f ablum-pull-secret.yaml
```

```
secret/1979710-ablum-rhel8-training-pull-secret created
```

```
[student@workstation ~]$ oc secrets link --for=pull default
```

```
1979710-ablum-rhel8-training-pull-secret
```

```
[student@workstation ~]$ oc secrets link --for=mount builder
```

```
1979710-ablum-rhel8-training-pull-secret
```

```
[student@workstation ~]$ oc get sa default -o yaml
```

```
apiVersion: v1
```

```
imagePullSecrets:
```

```
- name: default-dockercfg-62t9g
```

```
- name: 1979710-ablum-rhel8-training-pull-secret
```

```
[student@workstation ~]$ oc new-app
```

```
registry.redhat.io/openjdk/openjdk-11-rhel7~https://github.com/jboss-openshift/openshift-quickstarts --context-dir=undertow-servlet
```

```
--> Found Docker image 781deed (4 weeks old) from registry.redhat.io for  
"registry.redhat.io/openjdk/openjdk-11-rhel7"
```

```
[student@workstation ~]$ oc expose service/openshift-quickstarts
```

```
[student@workstation ~]$ curl openshift-quickstarts-javafun.apps.cluster.lab.example.com
```

```
Hello World
```

How to import your own images ?

```
[student@workstation ~]$ oc new-project myproject
```

```
[student@workstation ~]$ oc import-image myproject/mytest --from=quay.io/ajblum/mytest  
--confirm --all --scheduled=true
```

```
[student@workstation ~]$ oc get all
```

NAME	IMAGE REPOSITORY	TAGS	UPDATED
------	------------------	------	---------

imagestream.image.openshift.io/mytest

image-registry.openshift-image-registry.svc:5000/myproject/mytest 1.0,latest 2 minutes ago

Now, try to push a new image to the public registry:

```
[root@workstation ~]# podman login quay.io
```

Username: ajblum

Password:

Login Succeeded!

```
[root@workstation ~]# podman push 089c3c916a95 quay.io/ajblum/mytest:latest
```

Monitor [student@workstation ~]\$ **oc describe imagestream.image.openshift.io/mytest** for changes to the sha for the istag "latest"

```
[student@workstation ~]$ oc new-app myproject/mytest:latest
```

Now, from a private repo:

https://docs.openshift.com/container-platform/4.2/openshift_images/managing-images/using-image-pull-secrets.html

(show private docker.io repo docker.io/ajblum/do180)

```
[student@workstation ~]$ oc new-project test1
```

```
[student@workstation ~]$ vim secret
```

```
[student@workstation ~]$ oc create secret docker-registry ajblum
--docker-server=docker.io --docker-username=ajblum --docker-password=`cat secret`
```

```
[student@workstation ~]$ oc secrets link default ajblum --for=pull
```

```
[student@workstation ~]$ oc describe sa default
```

```
Name:          default
Namespace:      test1
Labels:         <none>
Annotations:    <none>
Image pull secrets: default-dockercfg-8g2q6
                  ajblum
Mountable secrets: default-token-v4g2f
                  default-dockercfg-8g2q6
Tokens:         default-token-blq4
                  default-token-v4g2f
Events:         <none>
```

```
[student@workstation ~]$ oc new-app --docker-image=docker.io/ajblum/do180:latest
```

```
[student@workstation ~]$ oc expose service/do180
```

```
[student@workstation ~]$ curl do180-test1.apps.cluster.lab.example.com
```

Hello OpenShift!

How can you push images directly into the internal image registry ?

<https://docs.openshift.com/container-platform/4.7/registry/securing-exposing-registry.html>

```
[student@workstation ~]$ kubectl
```

```
[student@workstation ~]$ oc edit configs.imageregistry.operator.openshift.io
```

spec:

```
  defaultRoute: true
```

```
[student@workstation ~]$ oc get routes
```

```
[student@workstation ~]$ sudo vi /etc/containers/registries.conf
```

```
[registries.insecure]
```

```
registries = ['default-route-openshift-image-registry.apps.ocp4.example.com']
```

```
[student@workstation ~]$ oc whoami -t
```

```
[student@workstation ~]$ podman login -u kubeadmin -p $(oc whoami -t)
```

```
default-route-openshift-image-registry.apps.ocp4.example.com
```

Now, deploy an application from an imagestream

```
[student@workstation ~]$ oc project developer-route
[student@workstation ~]$ oc login -u developer -p developer
[student@workstation ~]$ oc create imagestream php-hello-dockerfile
[student@workstation ~]$ oc login -u admin -p redhat
[student@workstation ~]$ oc whoami -t
[student@workstation ~]$ oc get routes -A
[student@workstation ~]$ podman push 4dc9cf1e0b6d
default-route-openshift-image-registry.apps.ocp4.example.com/developer-route/php-hello-dockerfile
[student@workstation ~]$ oc login -u developer -p developer
[student@workstation ~]$ oc get is
[student@workstation ~]$ oc new-app -i php-hello-dockerfile --name php-helloworld
[student@workstation ~]$ oc expose service/php-helloworld
[student@workstation ~]$ oc get routes
[student@workstation ~]$ curl php-helloworld-developer-route.apps.ocp4.example.com
[student@workstation ~]$ podman logout
default-route-openshift-image-registry.apps.ocp4.example.com
```

Can you use a build strategy based on a Containerfile/Dockerfile ?

https://docs.openshift.com/container-platform/4.8/cicd/builds/build-strategies.html#build-strategy-docker-build_build-strategies

```
[student@workstation ~]$ oc new-app --help
```

--strategy=: Specify the build strategy to use if you don't want to detect (docker|pipeline|source).

```
[student@workstation ~]$ mkdir mytest
[student@workstation ~]$ cd mytest
[student@workstation mytest]$ vim Dockerfile
FROM registry.access.redhat.com/ubi8:latest
MAINTAINER Andrew Blum <ablum@redhat.com>
ENTRYPOINT ["/usr/bin/sleep","5000"]
```

```
[student@workstation mytest]$ cd ~
[student@workstation ~]$ oc new-project ablum-mytest
[student@workstation ~]$ oc new-app --name mytest --strategy=docker
/home/student/mytest
```

```
[student@workstation ~]$ oc start-build buildconfig.build.openshift.io/mytest --from-dir /home/student/mytest/
```

Uploading directory "/home/student/mytest" as binary input for the build ...

```
[student@workstation ~]$ oc logs buildconfig.build.openshift.io/mytest
```

Receiving source from STDIN as archive ...

Replaced Dockerfile FROM image registry.access.redhat.com/ubi8:latest

Caching blobs under "/var/cache/blobs".

Pulling image

registry.access.redhat.com/ubi8@sha256:5e334d76fc059f7b44ee8fc2da6a2e8b240582d0214364c8c88596d20b33d7f1 ...

Getting image source signatures

Copying blob

sha256:262268b65bd5f33784d6a61514964887bc18bc00c60c588bc62bfae7edca46f1

Copying blob

sha256:06038631a24a25348b51d1bfc7d0a0ee555552a8998f8328f9b657d02dd4c64c

Copying config

sha256:53ce4390f2adb1681eb1a90ec8b48c49c015e0a8d336c197637e7f65e365fa9e

Writing manifest to image destination

Storing signatures

Adding transient rw bind mount for /run/secrets/rhsm

STEP 1: FROM

registry.access.redhat.com/ubi8@sha256:5e334d76fc059f7b44ee8fc2da6a2e8b240582d0214364c8c88596d20b33d7f1

STEP 2: MAINTAINER Andrew Blum <ablum@redhat.com>

--> de87c382d1b

STEP 3: ENTRYPOINT ["/usr/bin/sleep","5000"]

--> 2cf98c8c01a

STEP 4: ENV "OPENSIFT_BUILD_NAME"="mytest-2"

"OPENSIFT_BUILD_NAMESPACE"="ablum-mytest"

--> 111fc391266

STEP 5: LABEL "io.openshift.build.name"="mytest-2"

"io.openshift.build.namespace"="ablum-mytest"

STEP 6: COMMIT temp.builder.openshift.io/ablum-mytest/mytest-2:f6fb5ea5

```
[student@workstation ~]$ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
mytest-2-build	0/1	Completed	0	41s
mytest-656cb4fc4c-dlz5h	1/1	Running	0	24s

```
[student@workstation ~]$
```

```
[student@workstation ~]$
```

```
[student@workstation ~]$ oc rsh mytest-656cb4fc4c-dlz5h
```

```
sh-4.4$ ls /proc
```


1

```
sh-4.4$ cat /proc/1/cmdline
```

```
/usr/bin/coreutils--coreutils-prog-shebang=sleep/usr/bin/sleep5000sh-4.4$
```

Now, we can try to make a change:

```
[student@workstation ~]$ vi mytest/Dockerfile
```

```
ENTRYPOINT ["/usr/bin/sleep","22000"]
```

```
[student@workstation ~]$ oc start-build bc/mytest --from-dir /home/student/mytest/
```

```
[student@workstation ~]$ oc rsh mytest-699c4f694c-r7qjr
```

```
sh-4.4$ cat /proc/1/cmdline
```

```
/usr/bin/coreutils--coreutils-prog-shebang=sleep/usr/bin/sleep22000
```

CREATING APPLICATIONS WITH THE OPENSIFT WEB CONSOLE

For extra information regarding the topology view

https://docs.openshift.com/container-platform/4.7/applications/application_life_cycle_management/odc-viewing-application-composition-using-topology-view.html

CHAPTER 7 DEPLOYING MULTI-CONTAINER APPLICATIONS

Objective: Describe considerations for containerizing applications with multiple container images.

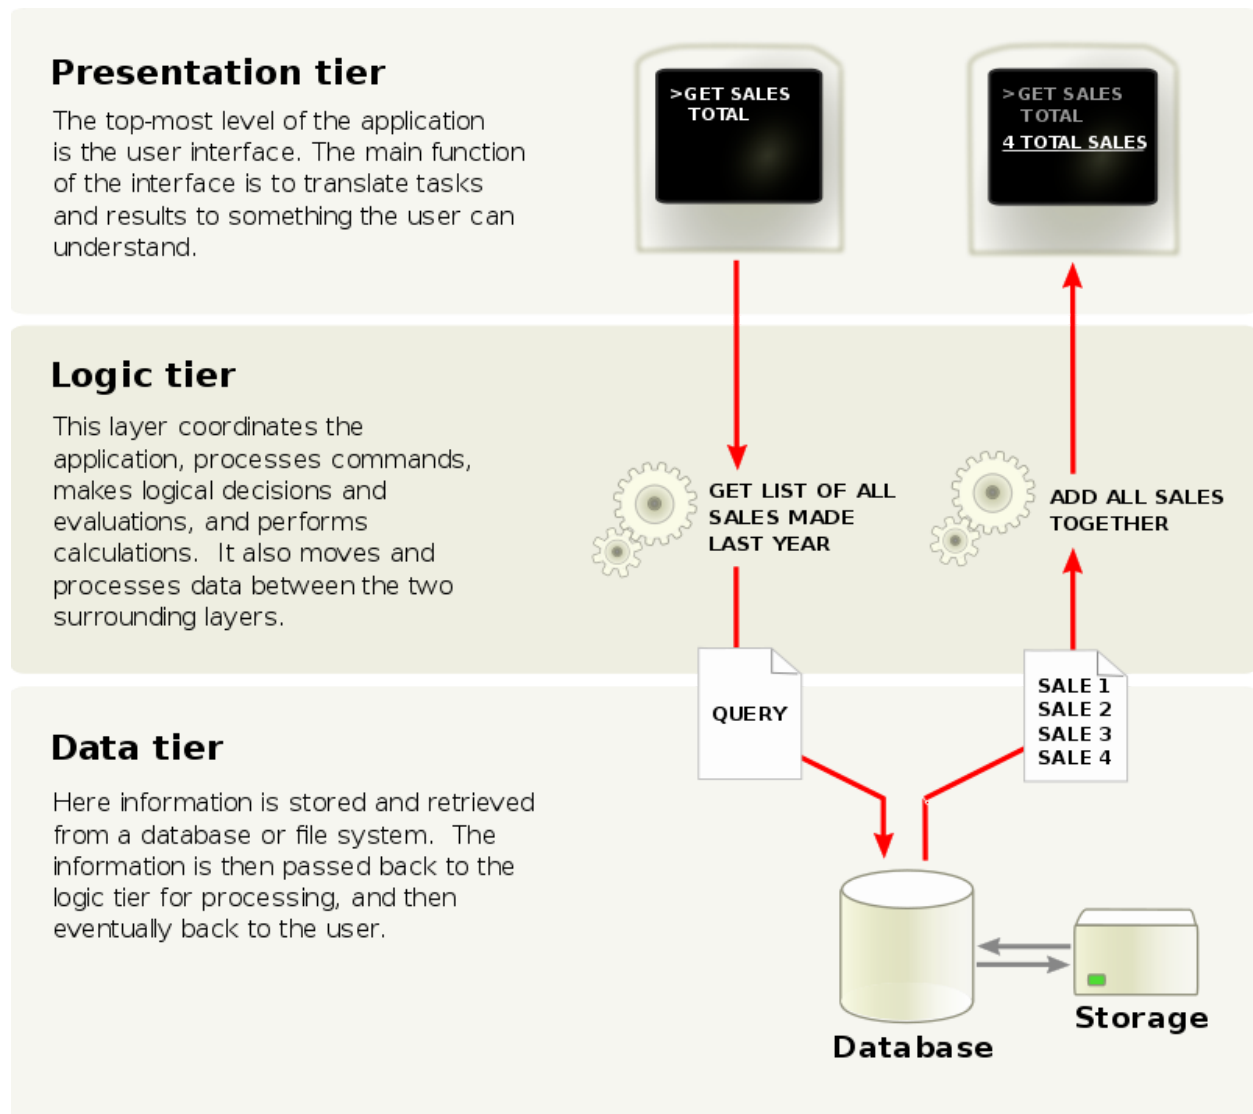
Objective: Deploy a multi-container application on OpenShift using a template.

CONSIDERATIONS FOR MULTI-CONTAINER APPLICATIONS

What are application tiers ?

Multi-tier applications can benefit from being deployed in separate containers (or pods).

https://en.wikipedia.org/wiki/Multitier_architecture



https://upload.wikimedia.org/wikipedia/commons/thumb/5/51/Overview_of_a_three-tier_application_vectorVersion.svg/800px-Overview_of_a_three-tier_application_vectorVersion.svg.png

Applications that are designed in “tiers” make good candidates to separate ie:

- Presentation
- Logic (Business)
- Data (Persistence)

Openshift has advantages over podman with applications like this. See:
Figure 7.1: A restart breaks three-tiered application links

Intro to the “TODO List” application used in this course

Figure 7.2: To Do List application logical architecture

How to use podman pod to deploy multiple containers ?

```
[student@workstation ~]# git clone https://github.com/ajblum/sampleapp.git
```

```
[studnet@workstation sampleapp]# cd sampleapp/
```

```
[student@workstation sampleapp]# podman build -t nodejs:1.0 .
```

STEP 1: FROM ubi8/nodejs-12

STEP 2: ADD app-src .

f170b8ebcb1bf745f7ff56d3ae398e464a33b2ba91c1d7240ee5ff624ddf411c

STEP 3: RUN npm install

npm notice created a lockfile as package-lock.json. You should commit this file.

up to date in 0.296s

found 0 vulnerabilities

fc17e774ec331fbe4d7e34b176e8fe422522ab1b11cf72055ea55dbbc0e6dccc

STEP 4: CMD npm run -d start

STEP 5: COMMIT nodejs:1.0

```
[student@workstation sampleapp]# podman pod create --name myapp1
```

```
[student@workstation sampleapp]# podman pod list
```

POD ID	NAME	STATUS	CREATED	# OF CONTAINERS	INFRA ID
c7cdb8663e1b	myapp1	Created	8 seconds ago	1	ac53be7d166b

```
[student@workstation sampleapp]# podman create --pod myapp1 -d nodejs:1.0
```

```
[student@workstation sampleapp]# podman create --pod myapp1 -d ubi8:latest sleep 5000
```

```
[student@workstation sampleapp]# podman pod start myapp1
```

```
[student@workstation sampleapp]$ podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED
d1b2833b71ad	registry.access.redhat.com/ubi8:latest	sleep 5000	11 seconds ago
Up 4 seconds ago	condescending_archimedes		
199f1fe4eb1e	localhost/nodejs:1.0	/bin/sh -c npm ru...	19 seconds ago
Up 5 seconds ago	boring_hellman		
2ad4c15c175a	k8s.gcr.io/pause:3.1		30 seconds ago
Up 5 seconds ago	ada250d7db62-infra		

```
[root@workstation sampleapp]# podman exec -it d1b2833b71ad /bin/bash
[root@myapp1 /]# yum install procps-ng
root@myapp1 /]# ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
root           1        0  0 15:33 ?    00:00:00 /usr/bin/coreutils --coreutils-prog-shebang=sleep
/usr/bin/sleep 5000
root           6        0  0 15:34 pts/0    00:00:00 /bin/bash
root          42        6  0 15:34 pts/0    00:00:00 ps -ef

[root@myapp1 /]# curl localhost:3000
Hello World[root@myapp1 /]# exit
```

```
[student@workstation sampleapp]$ podman pod list
POD ID      NAME STATUS      CREATED      # OF CONTAINERS  INFRA ID
ada250d7db62 myapp1 Running    4 minutes ago  3                2ad4c15c175a
contains containers and cannot be removed: container already exists
[student@workstation sampleapp]$ podman pod rm ad -f
```

How to deploy the todo list application in pod ?

```
[student@workstation ~]$ lab multicontainer-design start
```

```
[student@workstation networked]$ podman rm -a -f
```

```
[student@workstation networked]$ podman pod create --name todo -p 30081:30080
```

```
[student@workstation networked]$ podman pod list
POD ID      NAME      STATUS      CREATED      INFRA ID      # OF CONTAINERS
93eb660d7472 todo      Created      8 minutes ago 44ae8346607b 1
```

```
[student@workstation networked]$ podman pod inspect todo
(notice the infra pod, lets check to see what image it uses)
```

```
[student@workstation networked]$ podman ps -a
```

44ae8346607b registry.access.redhat.com/ubi8/pause:latest
ago Created 0.0.0.0:30081->30080/tcp 93eb660d7472-infra

10 minutes

```
[student@workstation networked]$ podman images
```

(determine container image id)

```
[student@workstation networked]$ podman inspect d41a45d48071
```

The pause container provides a way to run a container indefinitely. It will wait for interruption signals which terminate its execution.

We need this to hold our shared networking namespace used across the pod.

It runs a program (catatonit -P) <https://github.com/openSUSE/catatonit> “A container init that is so simple it's effectively brain-dead.” haha

```
[student@workstation ~]$ cd
```

```
/home/student/DO180/labs/multicontainer-design/deploy/nodejs/nodejs-source/models
```

```
[student@workstation models]$ vi db.js
```

```
module.exports.params = {  
  dbname: process.env.MYSQL_DATABASE,  
  username: process.env.MYSQL_USER,  
  password: process.env.MYSQL_PASSWORD,  
  params: {  
    host: '127.0.0.1',  
    port: '3306',  
    dialect: 'mysql'  
  }  
};
```

Pods will share networking namespace, thus 127.0.0.1 can be used for the nodetodo application to connect to the mysql database. Next, need to change the name of the image that will be built so as not to conflict with the one built previously

```
[student@workstation nodejs]$ cd
```

```
/home/student/DO180/labs/multicontainer-design/deploy/nodejs
```

Now, build this as before:

```
[student@workstation nodejs]$ ./build.sh
```

Preparing build folder

STEP 1: FROM registry.redhat.io/rhel8/nodejs-12:1

...SNIP...

Storing signatures

0f684b55c295d2b1870842bebc1b992d6b105abdb6cdc5482d72fd7ec78145f0

[student@workstation nodejs]\$

Now, we will configure persistent storage just like what was done in the ./run.sh script:

[student@workstation nodejs]\$ **cd networked/**

[student@workstation nodejs]\$ **cat run.sh**

(we wont use this start but get some reminders about how to create a persistent volume)

[student@workstation networked]\$ **mkdir -p work1/data**

[student@workstation networked]\$ **sudo chcon -Rt container_file_t work1**

[student@workstation networked]\$ **podman unshare chown -R 27:27 work1**

Now, we'll create our two containers but notice now we add in the --pod :

[student@workstation networked]\$ **podman create --pod todo --name mysql -e
MYSQL_DATABASE=items -e MYSQL_USER=user1 -e MYSQL_PASSWORD=mypa55 -e
MYSQL_ROOT_PASSWORD=r00tpa55 -v \$PWD/work1/data:/var/lib/mysql/data
registry.redhat.io/rhel8/mysql-80:1**

cc0fb27e73259c118729beb278cd4c982adca951b89dbb1a16c99d7ab198108

Notice, we don't need to configure mysql to use a host port since the networking will be shared within the pod.

[student@workstation networked]\$ **podman create --pod todo --name todoapi -e
MYSQL_DATABASE=items -e MYSQL_USER=user1 -e MYSQL_PASSWORD=mypa55
do180/todonodejs**

64d59133e1af963fb64b04f96fe0bbc0700802797f84d3325143f8daacd5cb7a

To start the application (and both the containers) use `podman pod start` like:

[student@workstation networked]\$ **podman pod list**

[student@workstation networked]\$ **podman pod inspect 5ae83fc0d4cf**

[student@workstation networked]\$ **podman pod start todo**

[student@workstation networked]\$ **podman ps**

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
cbacd5f9e5d5	localhost/do180/todonodejs-pod:latest	./run.sh	2 minutes ago	Up 2 minutes
0.0.0.0:30081->30080/tcp	todoapi			
362498e6f62e	registry.redhat.io/rhel8/mysql-80:1	run-mysqld	3 minutes ago	Up 2 minutes
0.0.0.0:30081->30080/tcp	mysql			

We will need to create one table in mysql manually for this application to work correctly:

```
[student@workstation networked]$ podman exec -it mysql /bin/bash
bash-4.4$ mysql -uroot
```

```
mysql> use items
Database changed
mysql> CREATE TABLE `Item` (`id` BIGINT not null auto_increment primary key,
`description` VARCHAR(100), `done` BIT);
Query OK, 0 rows affected (0.05 sec)
mysql> INSERT INTO `Item` (`id`,`description`,`done`) VALUES (1,'Pick up newspaper', 0);
Query OK, 1 row affected (0.02 sec)

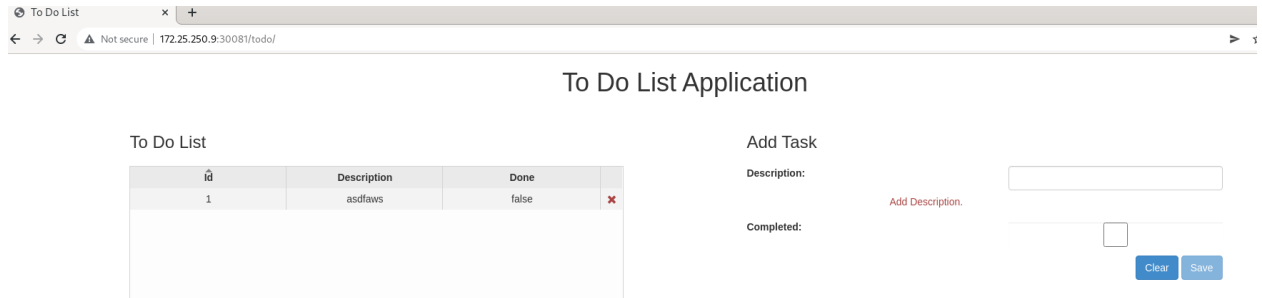
mysql> INSERT INTO `Item` (`id`,`description`,`done`) VALUES (2,'Buy groceries', 1);
Query OK, 1 row affected (0.02 sec)
```

Now, test:

```
[student@workstation networked]$ curl http://127.0.0.1:30081/todo/api/items/1
{"id":1,"description":"Pick up newspaper","done":false}
[student@workstation networked]$ curl http://127.0.0.1:30081/todo/api/items/2
{"id":2,"description":"Buy groceries","done":true}
```

```
[student@workstation ~]$ sudo firewall-cmd --add-port 30081/tcp
success
```

Go to <http://172.25.250.9:30081/todo/> and put some data in.



```
[student@workstation networked]$ curl http://127.0.0.1:30081/todo/api/items/1
{"id":1,"description":"asdfaws","done":false}
```

```
[student@workstation networked]$ podman pod stop todo
21e7a8c60238df5d37ecdab7e5b7dbe87a0fb0fc7af5e4b0b460a2bccaac2362
```

```
[student@workstation networked]$ podman ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
```

```
[student@workstation networked]$ podman ps
CONTAINER ID IMAGE COMMAND CREATED STATUS
PORTS NAMES
cbacd5f9e5d5 localhost/do180/todonodejs-pod:latest ./run.sh 9 minutes ago Up 24
seconds ago 0.0.0.0:30081->30080/tcp todoapi
362498e6f62e registry.redhat.io/rhel8/mysql-80:1 run-mysqld 9 minutes ago Up 24 seconds
ago 0.0.0.0:30081->30080/tcp mysql
```

```
[student@workstation networked]$ podman pod rm todo -f
21e7a8c60238df5d37ecdab7e5b7dbe87a0fb0fc7af5e4b0b460a2bccaac2362
```

TO enable systemd to manage the containers/pod:

https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html-single/building_running_and_managing_containers/index#assembly_porting-containers-to-systemd-using-podman_building-running-and-managing-containers

```
[student@workstation ~]$ lab multicontainer-design finish
```

Cleaning up the lab for Guided Exercise: Connecting Web Application and MySQL Container

- Stopping mysql container..... SUCCESS
- Removing mysql container..... SUCCESS

- Stopping todoapi container..... SUCCESS
- Removing todoapi container..... SUCCESS
- Removing registry.redhat.io/rhel8/mysql-80:1 image..... SUCCESS
- Removing registry.redhat.io/rhel8/nodejs-12:1 image..... SUCCESS
- Removing do180/todonodejs image..... SUCCESS
- Removing the project directory..... SUCCESS
- Removing the solution directory..... SUCCESS

DEPLOYING A MULTI-CONTAINER APPLICATION ON OPENSIFT

DEPLOYING A MULTI-CONTAINER APPLICATION ON OPENSIFT USING A TEMPLATE

Objective: After completing this section, students should be able to deploy a multicontainer application on OpenShift using a template.

Creating applications in Openshift require defining properties (like a name) for many resources. Often they are the same value for many of these resources:

- BuildConfig
- DeploymentConfig
- Service
- Route

Templates provide a way to simplify creation of resources that an application requires by including all the resources together as well as parameters that can be reused - like a name.

Extra TEMPLATE fun

```
[student@workstation ~]$ lab openshift-resources start
```

```
[student@workstation ~]$
```

```
[student@workstation ~]$ oc new-project ablum-template-fun
```

```
[student@workstation ~]$ oc get templates -n openshift
```

```
[student@workstation ~]$ oc get template httpd-example -n openshift
```

NAME	DESCRIPTION	PARAMETERS
OBJECTS		

httpd-example	An example Apache HTTP Server (httpd) application that serves static content...	9 (3 blank) 5
---------------	---	---------------

```
[student@workstation ~]$ oc describe template httpd-example -n openshift
```

```
[student@workstation ~]$ oc get template httpd-example -n openshift -o yaml > httpd.yaml
```

```
[student@workstation ~]$ oc process -f httpd.yaml --parameters
```

NAME	DESCRIPTION
------	-------------

GENERATOR	VALUE
-----------	-------

NAME	The name assigned to all of the frontend objects defined in this template.
	httpd-example

```
[student@workstation ~]$ oc process -f httpd.yaml -p NAME=myserver
```

```
[student@workstation ~]$ oc process -f httpd.yaml -p NAME=myserver -o yaml > httpd-processed.yaml
```

```
[student@workstation ~]$ oc project
```

Using project "ablum-template-fun" on server "https://api.ocp-na2.prod.nextcle.com:6443".

```
[student@workstation ~]$ oc create -f httpd-processed.yaml
```

service/myserver created

route.route.openshift.io/myserver created

imagestream.image.openshift.io/myserver created

buildconfig.build.openshift.io/myserver created

deploymentconfig.apps.openshift.io/myserver created

```
[student@workstation ~]$ oc logs -f bc/myserver
```

```
[student@workstation ~]$ oc logs -f dc/myserver
```

```
[student@workstation ~]$ curl
```

```
myserver-ablum-template-fun.apps.ocp-na2.prod.nextcle.com | grep Welcome
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
				Dload	Upload	Total	Spent Left Speed

```

0      0      0      0      0      0      0      0      0  --:--:-- --:--:-- --:--:--  0 <title>Welcome to
OpenShift</title>
100 37451 100 37451      0      0 355k      0  --:--:-- --:~:~:~ --:~:~:~  358k
<h1>Welcome to your static httpd application on OpenShift</h1>

```

```

[student@workstation ~]$ oc delete project ablum-template-fun
project.project.openshift.io "ablum-template-fun" deleted

```

Additional practice if needed

```

[student@workstation networked]$ oc get template mysql-ephemeral -n openshift
NAME      DESCRIPTION                                     PARAMETERS
OBJECTS
mysql-ephemeral  MySQL database service, without persistent storage. For more information
abou...  8 (3 generated)  3

```

```

[student@workstation networked]$ oc get template mysql-ephemeral -n openshift -o yaml |
less

```

```

apiVersion: template.openshift.io/v1
kind: Template
...SNIP...
metadata:
  annotations:
    tags: database,mysql                                <--- Makes it easier to find in UI
  labels:
    samplesoperator.config.openshift.io/managed: "true"
  name: mysql-ephemeral                                <-- Template name (not the objects)
objects:                                                <-- these are the ocp resources to be created
- apiVersion: v1
  kind: Secret
...SNIP...
  name: ${DATABASE_SERVICE_NAME}  <- notice how this parameter is reused
...SNIP...
parameters:
...SNIP..

```

- description: The name of the OpenShift Service exposed for the database.
 displayName: Database Service Name
name: DATABASE_SERVICE_NAME <- here is defines the value of the parm
 required: true
 value: mysql
 ...SNIP...
- description: Password for the MySQL connection user.
 displayName: MySQL Connection Password
from: '[a-zA-Z0-9]{16}' <- this will generate a password using regex
 generate: expression
 name: MYSQL_PASSWORD
 required: true

Parameters

Parameters in templates have default values, but they are optional and can be replaced when processing a template:

```
[student@workstation template-fun]$ oc describe template mysql-ephemeral -n openshift
```

```

Name:          MYSQL_PASSWORD
Display Name:  MySQL Connection Password
Description:   Password for the MySQL connection user.
Required:     true
Generated:    expression
From:         [a-zA-Z0-9]{16}
  
```

An easy way to view the parameters is using the oc process command with --parameters:

```
[student@workstation template-fun]$ oc process --parameters mysql-ephemeral -n openshift
```

NAME	DESCRIPTION	GENERATOR
VALUE		
MEMORY_LIMIT	Maximum amount of memory the container can use.	
512Mi		
NAMESPACE	The OpenShift Namespace where the ImageStream resides.	
openshift		
DATABASE_SERVICE_NAME	The name of the OpenShift Service exposed for the database.	
mysql		

MYSQL_USER	Username for MySQL user that will be used for accessing the database.
expression	user[A-Z0-9]{3}
MYSQL_PASSWORD	Password for the MySQL connection user.
expression	[a-zA-Z0-9]{16}
MYSQL_ROOT_PASSWORD	Password for the MySQL root user.
expression	[a-zA-Z0-9]{16}
MYSQL_DATABASE	Name of the MySQL database accessed.
	sampledb
MYSQL_VERSION	Version of MySQL image to be used (5.7, or latest).
	5.7

Let's actually use this template to generate values to these parameters:

```
[student@workstation ~]$ oc get template mysql-ephemeral -n openshift -o yaml >
mysql-ephemeral.yaml
```

```
[student@workstation ~]$ oc process -f mysql-ephemeral.yaml -o yaml >
mysql-ephemeral-processed.yaml
```

```
[student@workstation template-fun]$ vim mysql-ephemeral-processed.yaml
```

```
stringData:
  database-name: sampledb
  database-password: tqm8vIBxdxuLU7L0
  database-root-password: U2U20TmNGYQ5jxD4
  database-user: userQLD
```

Suppose we want to use this template, but override some parameters:

```
[student@workstation ~]$ oc process -f mysql-ephemeral.yaml -p MYSQL_USER='dev' -p
MYSQL_PASSWORD='$P4SSD' -p MYSQL_DATABASE='bank' -o yaml >
mysql-ephemeral-processed.yaml
[student@workstation ~]$ vim mysql-ephemeral-processed.yaml
```

```
stringData:
  database-name: bank
  database-password: $P4SSD
  database-root-password: fEjPviDan6DujROp
  database-user: dev
```

To create the actual resources

```
[student@workstation ~]$ oc create -f mysqlprocessed.yaml
secret/mysql created
service/mysql created
deploymentconfig.apps.openshift.io/mysql created
```

```
[student@workstation ~]$ oc get secrets
[student@workstation ~]$ oc describe secrets mysql
```

```
[student@workstation ~]$ oc get all
```

```
[student@workstation template-fun]$ oc logs pod/mysql-1-h996k -f
```

```
=> sourcing 20-validate-variables.sh ...
=> sourcing 25-validate-replication-variables.sh ...
=> sourcing 30-base-config.sh ...
---> 20:53:54 Processing basic MySQL configuration files ...
=> sourcing 60-replication-config.sh ...
=> sourcing 70-s2i-config.sh ...
---> 20:53:54 Processing additional arbitrary MySQL configuration provided by s2i ...
=> sourcing 10-mysql57.cnf ...
=> sourcing 40-paas.cnf ...
=> sourcing 50-my-tuning.cnf ...
---> 20:53:54 Initializing database ...
```

```
[student@workstation template-fun]$ nohup oc port-forward pod/mysql-1-h996k 3306:3306 &
Forwarding from 127.0.0.1:3306 -> 3306
Forwarding from [::1]:3306 -> 3306
```

(open a 2nd terminal)

```
[root@workstation ~]# mysql -u dev -p'$P4SSD' -h localhost --protocol tcp
```

```
MySQL [(none)]> show databases;
```

```
+-----+
| Database          |
+-----+
| information_schema |
| bank              |
+-----+
```

CLEANUP

```
[student@workstation template-fun]$ oc delete project template-fun
project.project.openshift.io "template-fun" deleted
```

Other ways to use templates to process and create them:

```
$ oc new-app --template=mysql-ephemeral
$ oc process openshift/mysql-ephemeral
$ oc process -f mysql-ephemeral-template.yaml
```

Are there best practices guides available for creating OCP templates ?

The community maintains this <https://github.com/openshift/library> which includes several templates as well as a link to this document https://docs.okd.io/latest/dev_guide/templates.html#writing-templates.

My recommendation is to start with an existing template, preferably one already in ocp4 (ie `oc get templates -n openshift`).

Templates themselves can be published to Openshift cluster so other devs can build an app from it.

```
[student@workstation deploy-multicontainer]$ oc create -f todo-template.yaml
[student@workstation deploy-multicontainer]$ oc create -f todo-template.yaml -n openshift
```

Templates vs. Helm charts

<https://github.com/helm/helm/tree/master/docs/examples/nginx/templates>

CONFIGURING PERSISTENT STORAGE FOR OPENSIFT APPLICATIONS [NO LONGER POSSIBLE]

How do we provide persistent storage to an application ? How did we do it using podman ?

Ocp “pools” persistent storage as a cluster-wide resource. So, how can we avoid committing the same storage volume to multiple projects/applications ? With a reservation or **claim**.

PersistentVolumeClaim (pvc) and PersistentVolume (pv)

- AWS Elastic Block Store (EBS)
- Fibre Channel
- HostPath
- iSCSI
- NFS
- VMWare vSphere

```
[student@workstation ~]$ oc get pv
```

Error from server (Forbidden): persistentvolumes is forbidden: User "rhn-support-ablum" cannot list resource "persistentvolumes" in API group "" at the cluster scope

<https://docs.openshift.com/container-platform/4.1/storage/persistent-storage/persistent-storage-nfs.html>

Let's get an NFS server running:

```
[student@workstation ~]$ ssh root@services
```

Last login: Tue Jul 23 16:55:20 2019 from 172.25.250.250

```
[root@services ~]#
```

```
[root@services ~]#
```

```
[root@services ~]# cat /etc/exports
```

```
[root@services ~]# cat /etc/redhat-release
```

Red Hat Enterprise Linux Server release 7.6 (Maipo)

```
#!/bin/bash
```

```
yum install nfs-utils rpcbind -y
```

```
useradd -u 5555 nfsuser
```

```
mkdir -p /opt/openshift
```

```
chown :5555 /opt/openshift
```

```
chmod 2770 /opt/openshift
```

```
ls -ld /opt/openshift
echo '/opt/openshift *(all_squash,anongid=5555,rw,sync)' > /etc/exports
cat /etc/exports
firewall-cmd --zone=public --add-port=2049/tcp --permanent
firewall-cmd --reload
firewall-cmd --list-all
systemctl start rpcbind
systemctl start nfs
systemctl enable nfs-server
exportfs -v
```

```
[root@services ~]# yum install nfs-utils rpcbind
[root@services ~]# useradd -u 5555 nfsuser
[root@services ~]# mkdir -p /opt/openshift
[root@services ~]# chown :5555 /opt/openshift/
[root@services ~]# chmod 2770 /opt/openshift/
[root@services ~]# ls -ld /opt/openshift/
drwxrws---. 2 root nfsuser 6 Aug  9 14:41 /opt/openshift/
[root@services ~]# vi /etc/exports
/opt/openshift *(all_squash,anongid=5555,rw,sync)
[root@services ~]# firewall-cmd --zone=public --add-port=2049/tcp --permanent
success
[root@services ~]# firewall-cmd --reload
success
[root@services ~]# firewall-cmd --list-all
[root@services ~]# systemctl start rpcbind
[root@services ~]# systemctl start nfs
[root@services ~]# systemctl enable nfs-server
```

Test as a client:

```
[root@workstation ~]# mkdir /opt/data
[root@workstation ~]# mount -t nfs services.lab.example.com:/opt/openshift /opt/data
[root@workstation ~]# touch /opt/data/file1
[root@workstation ~]# ssh root@services ls -l /opt/openshift
Warning: Permanently added 'services,172.25.250.13' (ECDSA) to the list of known hosts.
total 0
-rw-r--r--. 1 nfsnobody nfsuser 0 Aug 18 21:45 file1
[root@workstation ~]# umount /opt/data/
```

Create a pv in ocp:

```
[student@workstation ~]$ lab openshift-resources start
```

```
[student@workstation ~]$ cat .kubeadmin
```

```
jSJUP-HbEee-gx4pL-Ller5
```

```
[student@workstation ~]$ oc login -u kubeadmin -p jSJUP-HbEee-gx4pL-Ller5
```

```
[student@workstation template-fun]$ oc new-project pv-fun
```

<https://docs.openshift.com/container-platform/4.1/storage/persistent-storage/persistent-storage-nfs.html>

```
[student@workstation template-fun]$ vi pv01.yaml
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv01
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
  nfs:
    path: /opt/openshift
    server: 172.25.250.13
  persistentVolumeReclaimPolicy: Retain
```

```
[student@workstation ~]$ oc create -f pv01.yaml
persistentvolume/pv01 created
```

```
[student@workstation ~]$ oc get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
pv01	5Gi	RWO	Retain	Available	31s

```
[student@workstation ~]$ oc describe pv pv01
```

Name: pv01

Labels: <none>
Annotations: <none>
Finalizers: [kubernetes.io/pv-protection]
StorageClass:
Status: Available
Claim:
Reclaim Policy: Retain
Access Modes: RWO
Capacity: 5Gi
Node Affinity: <none>
Message:
Source:
 Type: NFS (an NFS mount that lasts the lifetime of a pod)
 Server: 172.25.250.13
 Path: /opt/openshift
 ReadOnly: false
Events: <none>

Now, let's try to use this storage:

```
[student@workstation ~]$ oc new-app  
--docker-image=registry.access.redhat.com/rhsc1/mysql-57-rhel7 -e  
MYSQL_ROOT_PASSWORD=passwd
```

```
[student@workstation template-fun]$ vim nfs-claim.yaml
```

```
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: nfs-claim1  
spec:  
  accessModes:  
    - ReadWriteOnce  
  resources:  
    requests:  
      storage: 5Gi
```

```
[student@workstation ~]$ oc get pv
```

```
[student@workstation template-fun]$ oc create -f nfs-claim.yaml
```

persistentvolumeclaim/nfs-claim1 created

```
[student@workstation ~]$ oc get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
pv01	5Gi	RWO	Retain	Bound	pv-fun/nfs-claim1

```
[student@workstation ~]$ oc get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
nfs-claim1	Bound	pv01	5Gi	RWO		14s

```
[student@workstation ~]$ oc edit deploymentconfig.apps.openshift.io/mysql-57-rhel7
```

spec:

containers:

- env:

- name: MYSQL_ROOT_PASSWORD

value: passwd

image:

registry.access.redhat.com/rhsc1/mysql-57-rhel7@sha256:225ccc7a059a8e615ce3e5f4aa6fae507a9e03b4562dccf3a98b09ff69bcdeb7

imagePullPolicy: Always

name: mysql-57-rhel7

ports:

- containerPort: 3306

protocol: TCP

resources: {}

terminationMessagePath: /dev/termination-log

terminationMessagePolicy: File

volumeMounts:

- mountPath: "/opt/data"

name: mypv

dnsPolicy: ClusterFirst

restartPolicy: Always

schedulerName: default-scheduler

securityContext: {}

terminationGracePeriodSeconds: 30

volumes:

- name: mypv

persistentVolumeClaim:

claimName: nfs-claim1

```
[student@workstation ~]$ oc rsh pod/mysql-57-rhel7-3-rl4sp
```

```
sh-4.2$ df /opt/data
```

```
Filesystem            1K-blocks    Used Available Use% Mounted on
172.25.250.13:/opt/openshift 104845312 5324544 99520768 6% /opt/data
```

```
sh-4.2$ touch /opt/data/test1
```

```
touch: cannot touch '/opt/data/test1': Permission denied
```

```
sh-4.2$ ls -ld /opt/data
```

```
drwxrws---. 3 root 5555 22 Aug 20 14:33 /opt/data
```

```
sh-4.2$ id
```

```
uid=1000410000 gid=0(root) groups=0(root),1000410000
```

```
[root@cluster-worker-1 ~]# journalctl -a | grep -i avc
```

```
Aug 23 16:05:18 cluster-worker-1 kernel: type=1400 audit(1566576318.390:8): avc: denied {
write } for pid=38574 comm="touch" name="openshift" dev="0:465" ino=92276291
scontext=system_u:system_r:container_t:s0:c0,c21 tcontext=system_u:object_r:nfs_t:s0
tclass=dir permissive=0
```

```
[root@cluster-worker-1 ~]# setsebool -P virt_use_nfs 1
```

```
[root@workstation ~]# ssh core@worker0 sudo setsebool -P virt_use_nfs 1
```

```
[root@workstation ~]# ssh core@master0 sudo setsebool -P virt_use_nfs 1
```

Changes made to nodes could be handled by the Node Tuning Operator:

https://docs.openshift.com/container-platform/4.2/scalability_and_performance/using-node-tuning-operator.html

```
[student@workstation ~]$ oc delete pod/mysql-57-rhel7-2-z7vk4
```

```
pod "mysql-57-rhel7-2-z7vk4" deleted
```

```
sh-4.2$ touch /opt/data/test1
```

Other considerations::

Security Context Constraints (scc)

```
[student@workstation ~]$ oc adm policy add-scc-to-user anyuid -z default
```

Regular filesystem permissions (adding a group to the user running in the container):

```
[student@workstation ~]$ oc edit deploymentconfig.apps.openshift.io/mysql-57-rhel7
```

```
spec:
```

```
  containers:
```

```
    - env:
```

```
      - name: MYSQL_ROOT_PASSWORD
```

```
        value: passwd
```

```
    image:
```

```
registry.access.redhat.com/rhscv/mysql-57-rhel7@sha256:225ccc7a059a8e615ce3e5f4aa6fae507a9e03b4562dccf3a98b09ff69bcdeb7
```

```
    imagePullPolicy: Always
```

```
    name: mysql-57-rhel7
```

```
    ports:
```

```
      - containerPort: 3306
```

```
        protocol: TCP
```

```
    resources: {}
```

```
    terminationMessagePath: /dev/termination-log
```

```
    terminationMessagePolicy: File
```

```
    volumeMounts:
```

```
      - mountPath: /opt/data
```

```
        name: mypv
```

```
  securityContext:
```

```
    supplementalGroups: [5555]
```

CHAPTER 8: TROUBLESHOOTING

How would one collect an application core from an application that crashes while running in a container ?

KCS indicates that the recommended approach is to create a PVC for the pod with the crashing application. Then, add a mount inside the pod to use the PV. The application would also need to be configured to dump to that location. The underlying container image might need a few more yum install to get symbols for some common libraries installed. Here are two examples from OCP3 that should work in OCP4:

<https://access.redhat.com/solutions/3124061>

<https://access.redhat.com/solutions/3374631>

<https://developers.redhat.com/blog/2020/01/09/debugging-applications-within-red-hat-openshift-containers>

I doubt you'd want to actually do any debugging (ie gdb) in the container. So, after the core is generated copying it off with scp or rsync would be what I'd suggest. I asked an SME in sbr-shift about this and here was his response:

""

Application cores have always been a problem, we really should be using a tool like FAF <https://retrace.fedoraproject.org/faf/summary/> and ABRT to collect this information, review and analyze it.

""

How to check node module version using npm

```
[ablum@badger ~]$ npm view express@4.14.2
```

```
--registry="http://nexus-common.apps.na45.prod.nextcle.com/repository/npm-proxy"
```

```
[ablum@badger ~]$ npm view express@4.14.1
```

```
--registry="http://nexus-common.apps.na45.prod.nextcle.com/repository/npm-proxy"
```

```
express@4.14.1 | MIT | deps: 26 | versions: 264
```

```
[ablum@badger ~]$ npm view express versions
```

```
--registry="http://nexus-common.apps.na45.prod.nextcle.com/repository/npm-proxy"
```

This indicates that 4.17.1 is the latest available from this registry while 4.14.0 and 4.14.1 is the only available in 4.14.x.

Note this really just a proxy to the public npm registry: `npm view express --registry="<https://registry.npmjs.org>"`

After discussing with the developer, any 4.14.x version of express dependency will do. So, we can change:

```
[student@workstation DO180-apps]$ vi ~/DO180-apps/nodejs-helloworld/package.json
```

```
"dependencies": {  
  "express": "4.14.x"  
}
```

```
[student@workstation DO180-apps]$ git commit -am "Fixed Express release"
```

...output omitted...

1 file changed, 1 insertion(+), 1 deletion(-)

```
[student@workstation DO180-apps]$ git push
```

```
[student@workstation DO180-apps]$ oc start-build bc/nodejs-hello
```

```
[student@workstation DO180-apps]$ oc logs -f bc/nodejs-hello
```

```
[student@workstation DO180-apps]$ oc get pods
```

How to use public nexus registry instead of the one in DO180

Sometimes we are having an issue with the nexus repo in the training environment:

```
$ npm view express
```

```
--registry="http://nexus-common.apps.ocp-na2.prod.nexctl.com/repository/npm-proxy"
```

```
npm ERR! Unexpected token < in JSON at position 0 while parsing near '<!DOCTYPE html  
PUBLI...'
```

```
$ npm view express --registry="https://registry.npmjs.org"
```

```
express@4.17.1 | MIT | deps: 30 | versions: 264
```

```
Fast, unopinionated, minimalist web framework
```

```
[student@workstation DO180-apps]$ oc edit bc nodejs-hello
```

```
strategy:
  sourceStrategy:
    env:
      - name: npm_config_registry
        value: https://registry.npmjs.org
    from:
```

```
[student@workstation DO180-apps]$ oc start-build bc/nodejs-hello
```

```
[student@workstation DO180-apps]$ oc logs bc/nodejs-hello
```

Successfully pushed

image-registry.openshift-image-registry.svc:5000/rhn-support-ablum-nodejs/nodejs-hello@sha256:7f19fae9e58909e0aaf918ba6b5b6f54598265d2514193da67258db741fb15ae

Push successful

Build was successful ! What about the deployment ?

```
[student@workstation nodejs-helloworld]$ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nodejs-hello-1-build	0/1	Error	0	15m
nodejs-hello-2-build	0/1	Completed	0	114s
nodejs-hello-85b47f4f4c-58zgb	0/1	CrashLoopBackOff	3	80s

```
[student@workstation nodejs-helloworld]$ oc logs nodejs-hello-85b47f4f4c-58zgb
```

Environment:

DEV_MODE=false

NODE_ENV=production

DEBUG_PORT=5858

Launching via npm...

npm info it worked if it ends with ok

npm info using npm@6.14.8

npm info using node@v12.19.1

npm ERR! missing script: start

npm timing npm Completed in 82ms

npm ERR! A complete log of this run can be found in:

npm ERR! /opt/app-root/src/.npm/_logs/2021-07-02T14_40_07_931Z-debug.log

```
[student@workstation DO180-apps]$ vi ~/DO180-apps/nodejs-helloworld/package.json
```

```
"description": "Hello World!",  
"main": "app.js",  
  "scripts": {"start": "node app.js"},
```

```
[student@workstation DO180-apps]$ git commit -am "Added start up script"
```

```
...output omitted...
```

```
1 file changed, 3 insertions(+)
```

```
[student@workstation DO180-apps]$ git push
```

```
[student@workstation DO180-apps]$ oc start-build bc/nodejs-hello
```

```
[student@workstation DO180-apps]$ oc logs dc/nodejs-hello
```

```
> nodejs-helloworld@1.0.0 start /opt/app-root/src
```

```
> node app.js
```

Example app listening on port 8080!

```
[student@workstation DO180-apps]$ oc get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nodejs-hello	ClusterIP	172.30.208.10	<none>	8080/TCP	48m

```
[student@workstation nodejs-helloworld]$ oc expose service nodejs-hello
```

```
route.route.openshift.io/nodejs-hello exposed
```

```
[student@workstation nodejs-helloworld]$ oc get routes
```

NAME	HOST/PORT	PATH	SERVICES	PORT
	TERMINATION	WILDCARD		

nodejs-hello	nodejs-hello-rhn-support-ablum-nodejs.apps.na45.prod.nextcle.com
--------------	---

nodejs-hello	8080-tcp	None
--------------	----------	------

```
[student@workstation nodejs-helloworld]$ curl
```

```
nodejs-hello-rhn-support-ablum-nodejs.apps.na45.prod.nextcle.com
```

```
Hello World!
```

```
[student@workstation nodejs-helloworld]$
```

```
[student@workstation nodejs-helloworld]$ cd ~
```

```
[student@workstation ~]$ lab troubleshoot-s2i finish
```

EXTRA Troubleshooting TIPS and TRICKS

Port-forward

For applications that are serving up content (listening on certain ports) the **podman port-forward** or even the **oc port-forward** can be used to verify the application endpoint directly independent of service/routes.

Another good use for port-forward is to gather additional diagnostics from an application. Some applications can be configured to allow debugging through a network port (ie java vi Java Debug Wire Protocol or JDWP). Once the app is configured (via standalone.conf) to enable remote debugging, the **oc port-forward** command can be used to establish a connection to the app running in the container.

```
[student@workstation ~]$ oc new-project ablum-mytest
```

```
[student@workstation ~]$ oc new-app --docker-image=openshift/hello-openshift --name hello
```

```
[student@workstation ~]$ oc port-forward pod/hello-7764b7f5f8-kpmzr 9091:8888
```

```
Forwarding from 127.0.0.1:9091 -> 8888
```

```
Forwarding from [::1]:9091 -> 8888
```

```
^C[student@workstation ~]$ nohup oc port-forward pod/hello-7764b7f5f8-kpmzr 9091:8888 &
```

```
[1] 9501
```

```
nohup: ignoring input and appending output to 'nohup.out'
```

```
[student@workstation ~]$ curl localhost:9091
```

```
Hello OpenShift!
```

```
[student@workstation ~]$
```

Maybe we can use a packet analyzer....

```
[student@workstation ~]$ sudo tcpdump -i any -s0 port 9091
```

```
[student@workstation ~]$ fg
```

```
nohup oc port-forward pod/hello-7764b7f5f8-kpmzr 9091:8888
```

```
^C[student@workstation ~]$
```

Logs and Events

Check the logs !

podman logs

cc logs

```
[student@workstation ~]$ oc logs pod/hello-7764b7f5f8-kpmzr
```

serving on 8888

serving on 8080

Servicing request.

Servicing request.

Servicing request.

```
$ oc get events
```

```
$ oc describe
```

```
[student@workstation ~]$ oc describe pod/hello-7764b7f5f8-kpmzr
```

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	Scheduled		<unknown>	default-scheduler Successfully assigned ablum-mytest/hello-7764b7f5f8-kpmzr to na45-mh9qn-worker-l886x
Normal	AddedInterface	10m	multus	Add eth0 [10.129.3.248/23]
Normal	Pulled	10m	kubelet, na45-mh9qn-worker-l886x	Container image "openshift/hello-openshift@sha256:aaea76ff622d2f8bcb32e538e7b3cd0ef6d291953f3e7c9f556c1ba5baf47e2e" already present on machine
Normal	Created	10m	kubelet, na45-mh9qn-worker-l886x	Created container hello
Normal	Started	10m	kubelet, na45-mh9qn-worker-l886x	Started container hello

Access running containers

```
$ podman exec
```

```
$ oc rsh
```

```
$ oc exec
```

```
[student@workstation ~]$ oc exec -h
```

Examples:

```
# Get output from running 'date' command from pod mypod, using the first container by default
oc exec mypod -- date
```

```
# Get output from running 'date' command in ruby-container from pod mypod
oc exec mypod -c ruby-container -- date
```

```
# Switch to raw terminal mode, sends stdin to 'bash' in ruby-container from pod mypod
# and sends stdout/stderr from 'bash' back to the client
oc exec mypod -c ruby-container -i -t -- bash -il
```

```
[student@workstation ~]$ oc exec -it hello-7764b7f5f8-kpmzr /bin/bash
```

```
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use
kubectl kubectl exec [POD] -- [COMMAND] instead.
```

```
ERRO[0000] exec failed: container_linux.go:348: starting container process caused "exec:
\"/bin/bash\": stat /bin/bash: no such file or directory"
```

```
exec failed: container_linux.go:348: starting container process caused "exec: \"bin/bash\": stat
/bin/bash: no such file or directory"
```

```
command terminated with exit code 1
```

```
[student@workstation ~]$ oc rsh hello-7764b7f5f8-kpmzr
```

```
ERRO[0000] exec failed: container_linux.go:348: starting container process caused "exec:
\"/bin/sh\": stat /bin/sh: no such file or directory"
```

```
exec failed: container_linux.go:348: starting container process caused "exec: \"bin/sh\": stat
/bin/sh: no such file or directory"
```

```
command terminated with exit code 1
```

What to do when the image used by the running container does not have a shell built in it.

Check out hello-openshift for example

<https://github.com/openshift/origin/blob/master/examples/hello-openshift/Dockerfile> :

Here, you will need to use alternative troubleshooting like:

- 1.) Build a new troubleshooting image that has additional tools (like an interactive shell)
- 2.) Study the application logs
- 3.) Use port-forward to redirect application traffic to a local debugging program
- 4.) Create a debug pod:

```
[student@workstation ~]$ oc debug deployment.apps/hello
```

```
Starting pod/hello-debug ...
```

```
Removing debug pod ...
```

```
error: container create failed: time="2020-10-02T11:16:20Z" level=error
```

```
msg="container_linux.go:348: starting container process caused \"exec: \\\"/bin/sh\\\": stat /bin/sh: no such file or directory\\\""
```

```
container_linux.go:348: starting container process caused "exec: \"/bin/sh\": stat /bin/sh: no such file or directory"
```

```
[student@workstation ~]$ oc debug --image=registry.access.redhat.com/ubi8:latest deployment.apps/hello
```

```
sh-4.4$ id
```

```
uid=1002450000(1002450000) gid=0(root) groups=0(root),1002450000
```

```
sh-4.4$ env | grep HELLO
```

```
HELLO_SERVICE_PORT=8080
```

```
HELLO_SERVICE_PORT_8888_TCP=8888
```

```
HELLO_PORT_8888_TCP=tcp://172.30.89.146:8888
```

```
HELLO_PORT_8080_TCP_PROTO=tcp
```

```
HELLO_SERVICE_HOST=172.30.89.146
```

```
HELLO_PORT_8888_TCP_PORT=8888
```

```
HELLO_PORT_8080_TCP_ADDR=172.30.89.146
```

```
HELLO_PORT=tcp://172.30.89.146:8080
```

```
HELLO_PORT_8888_TCP_PROTO=tcp
```

```
HELLO_SERVICE_PORT_8080_TCP=8080
```

```
HELLO_PORT_8080_TCP=tcp://172.30.89.146:8080
```

```
HELLO_PORT_8888_TCP_ADDR=172.30.89.146
```

```
HELLO_PORT_8080_TCP_PORT=8080
```

```
sh-4.4$ curl $HELLO_SERVICE_HOST:$HELLO_SERVICE_PORT
```

```
Hello OpenShift!
```

```
sh-4.4$ exit
```

```
exit
```

```
Removing debug pod ...
```

Copying files in and out of a container

```
# podman cp  
# oc rsync -h
```

Watch out though, there are some caveats:

<https://docs.openshift.com/container-platform/4.5/nodes/containers/nodes-containers-copying-files.html>

```
[student@workstation ~]$ mkdir data  
[student@workstation ~]$ echo hello > data/testfile
```

```
[student@workstation ~]$ oc rsync data hello-7764b7f5f8-kpmzr:/  
WARNING: cannot use rsync: rsync not available in container  
WARNING: cannot use tar: tar not available in container  
error: No available strategies to copy.
```